

Introduction

mongoDB

สถาบัน ไอเเวมซี



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

What is MongoDB ?

Humongous



What is MongoDB ?

Non-relational database

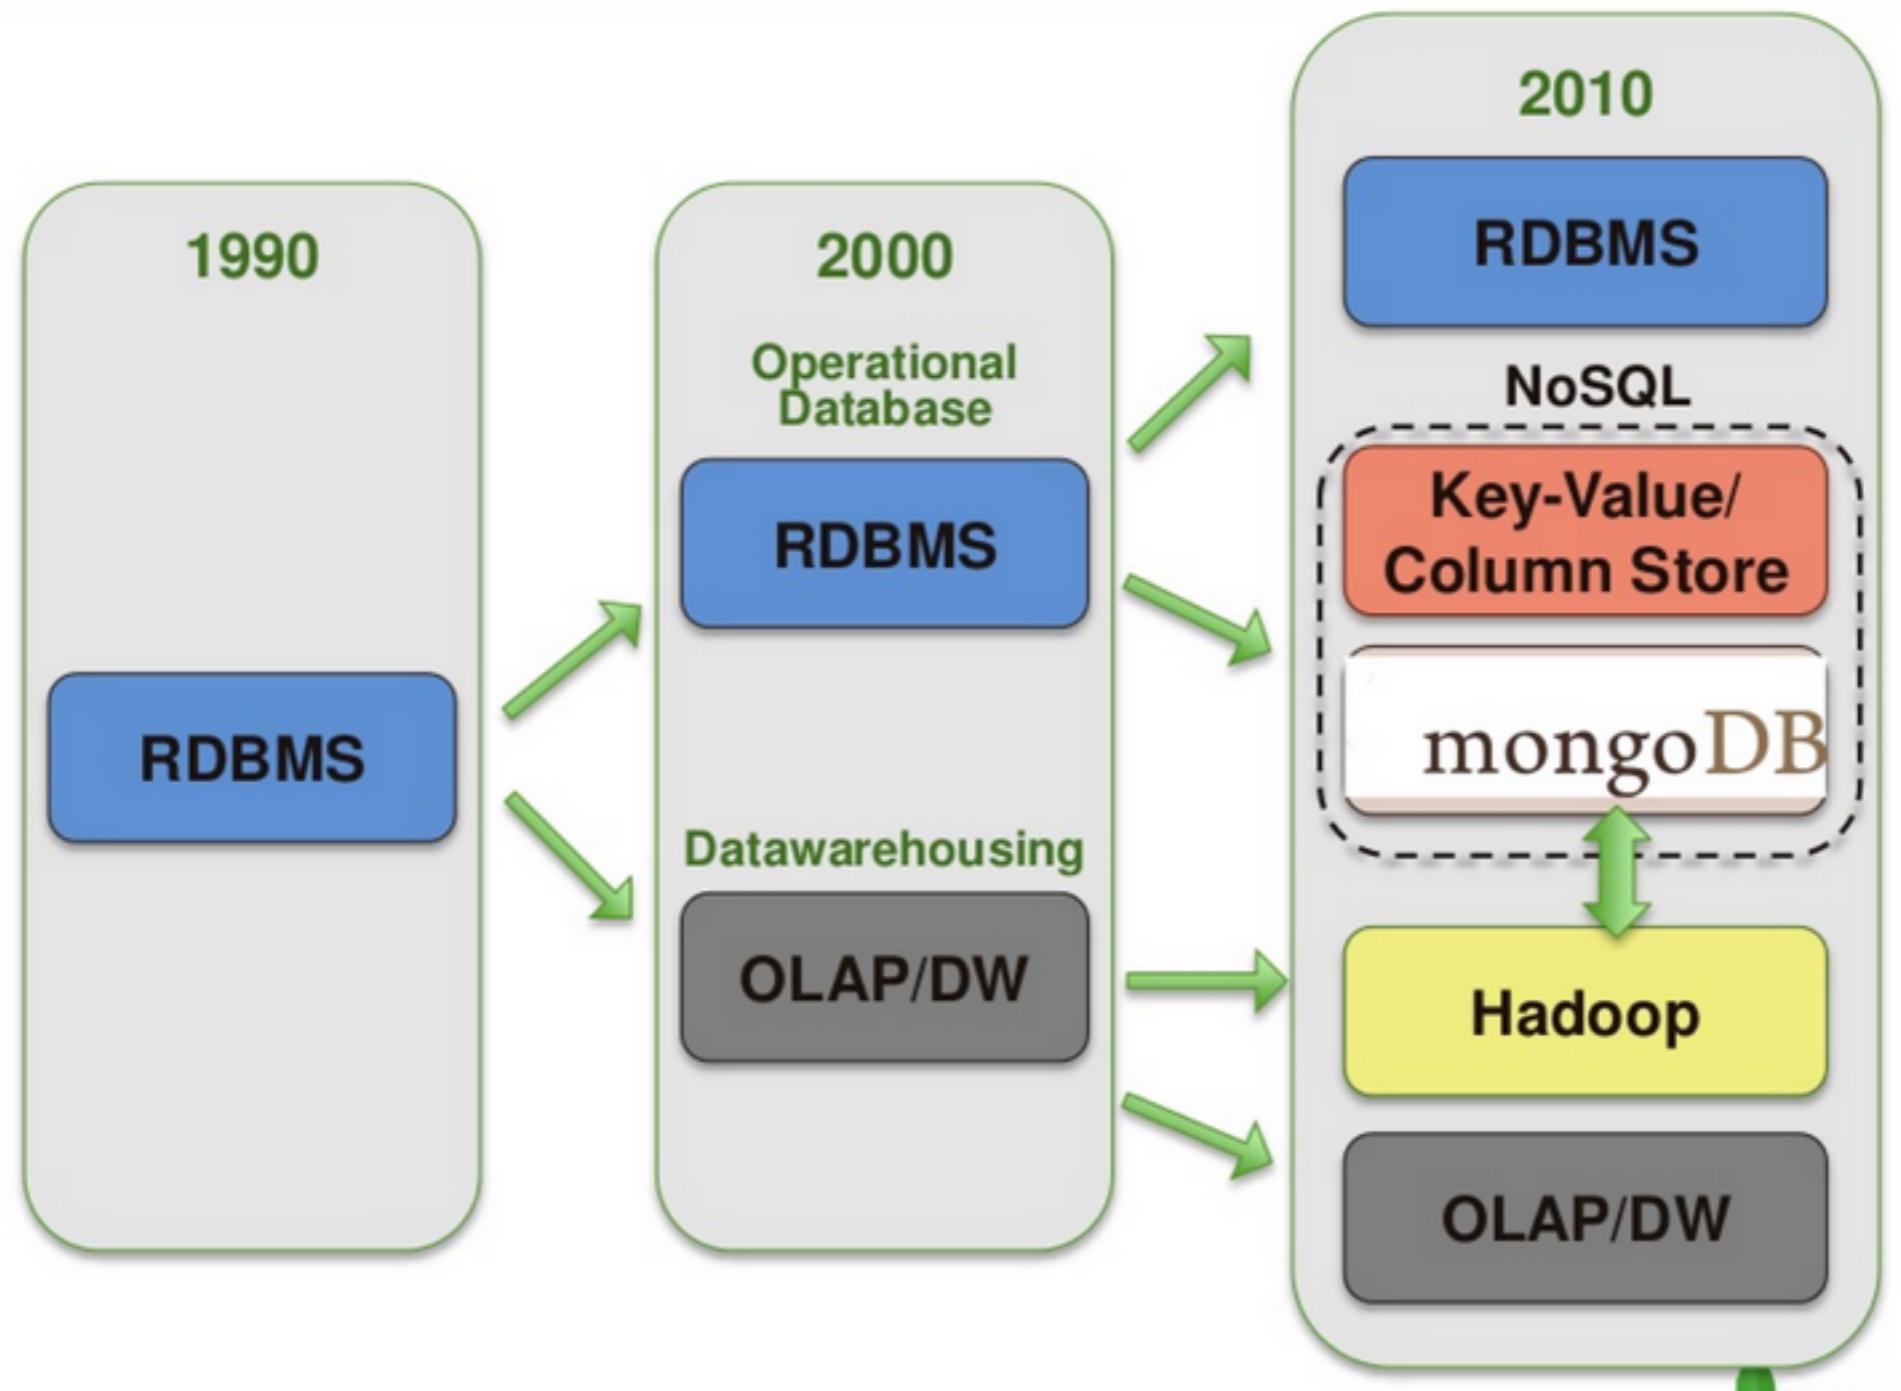
Store JSON document

Schema less

Not support Joins



Evolution of database



What is NoSQL ?

Not Only SQL

Non-relational database

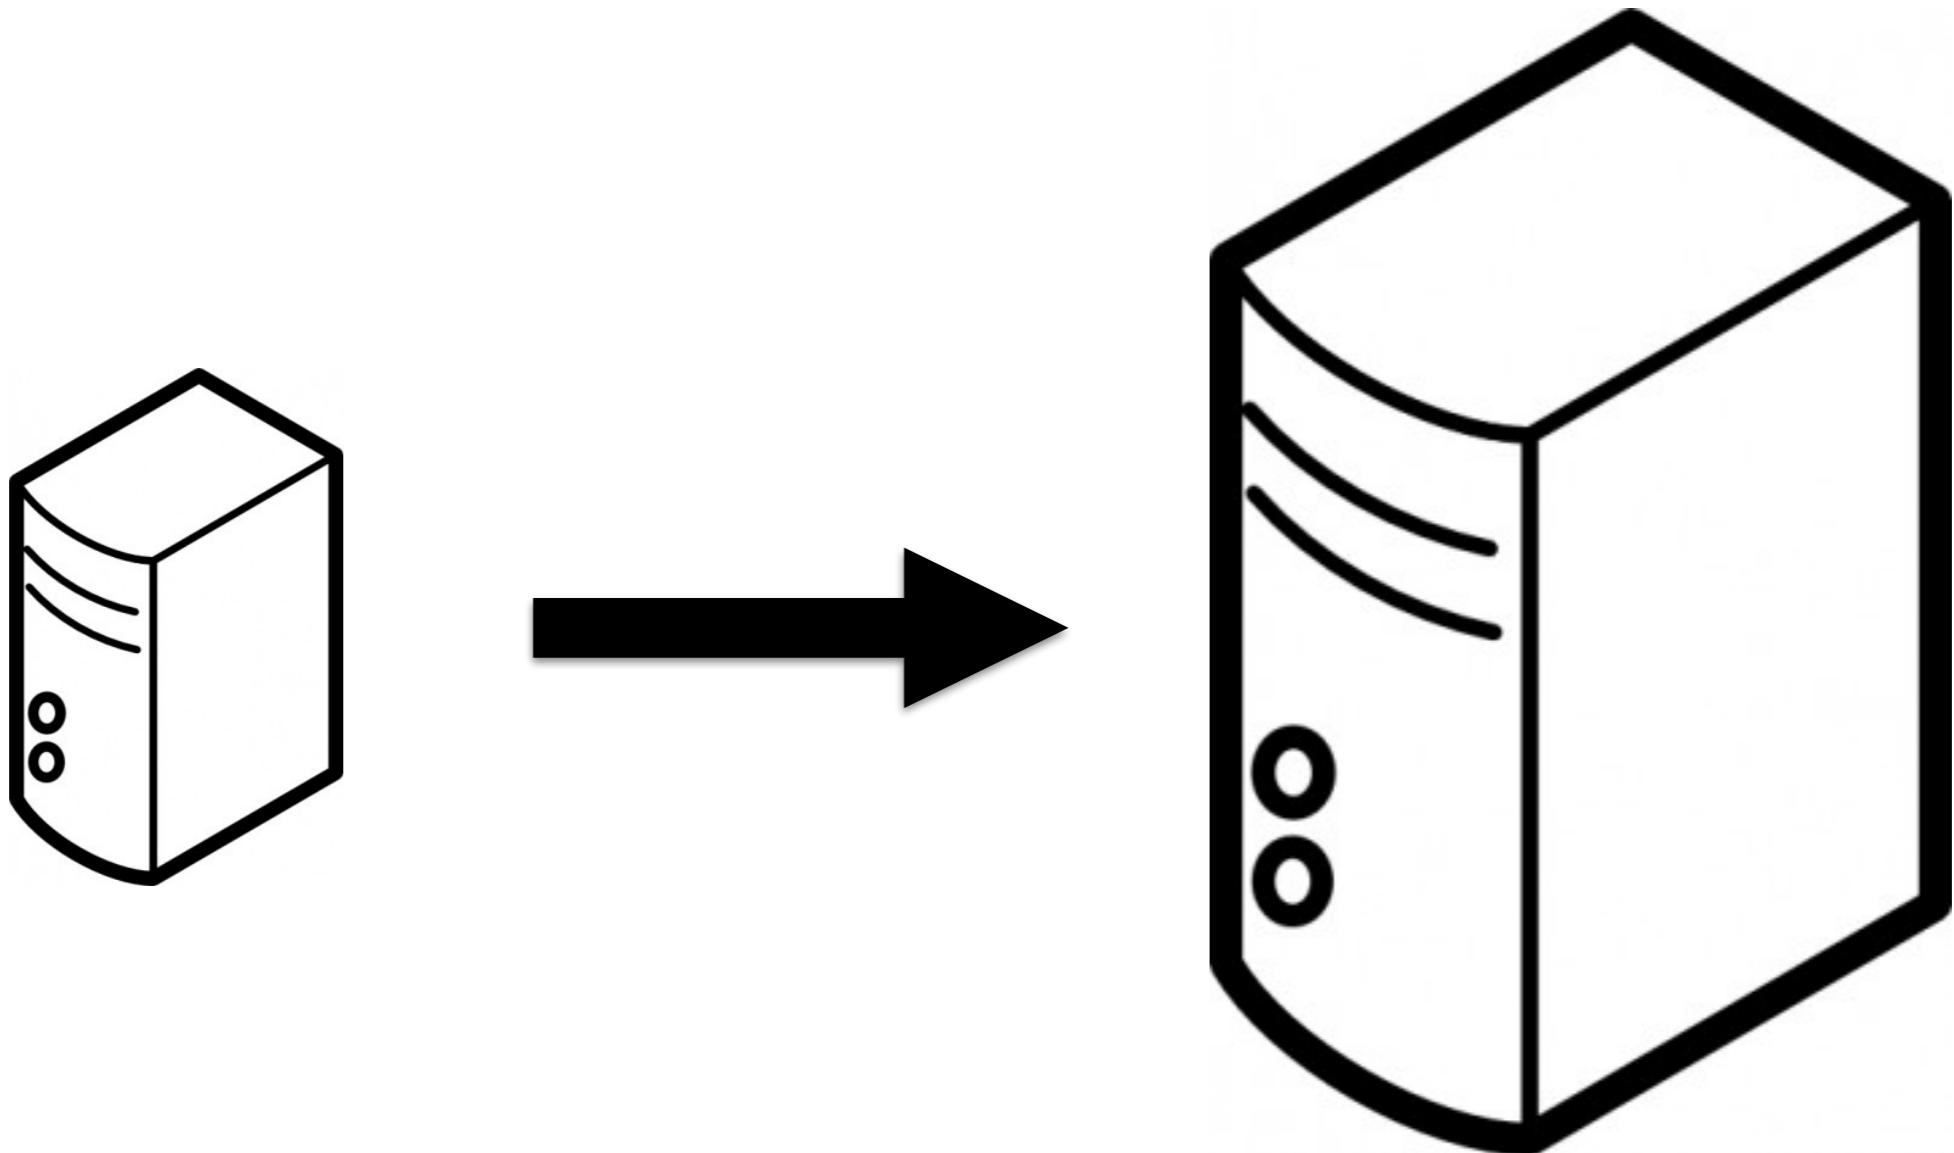
No joins

No complex transactions

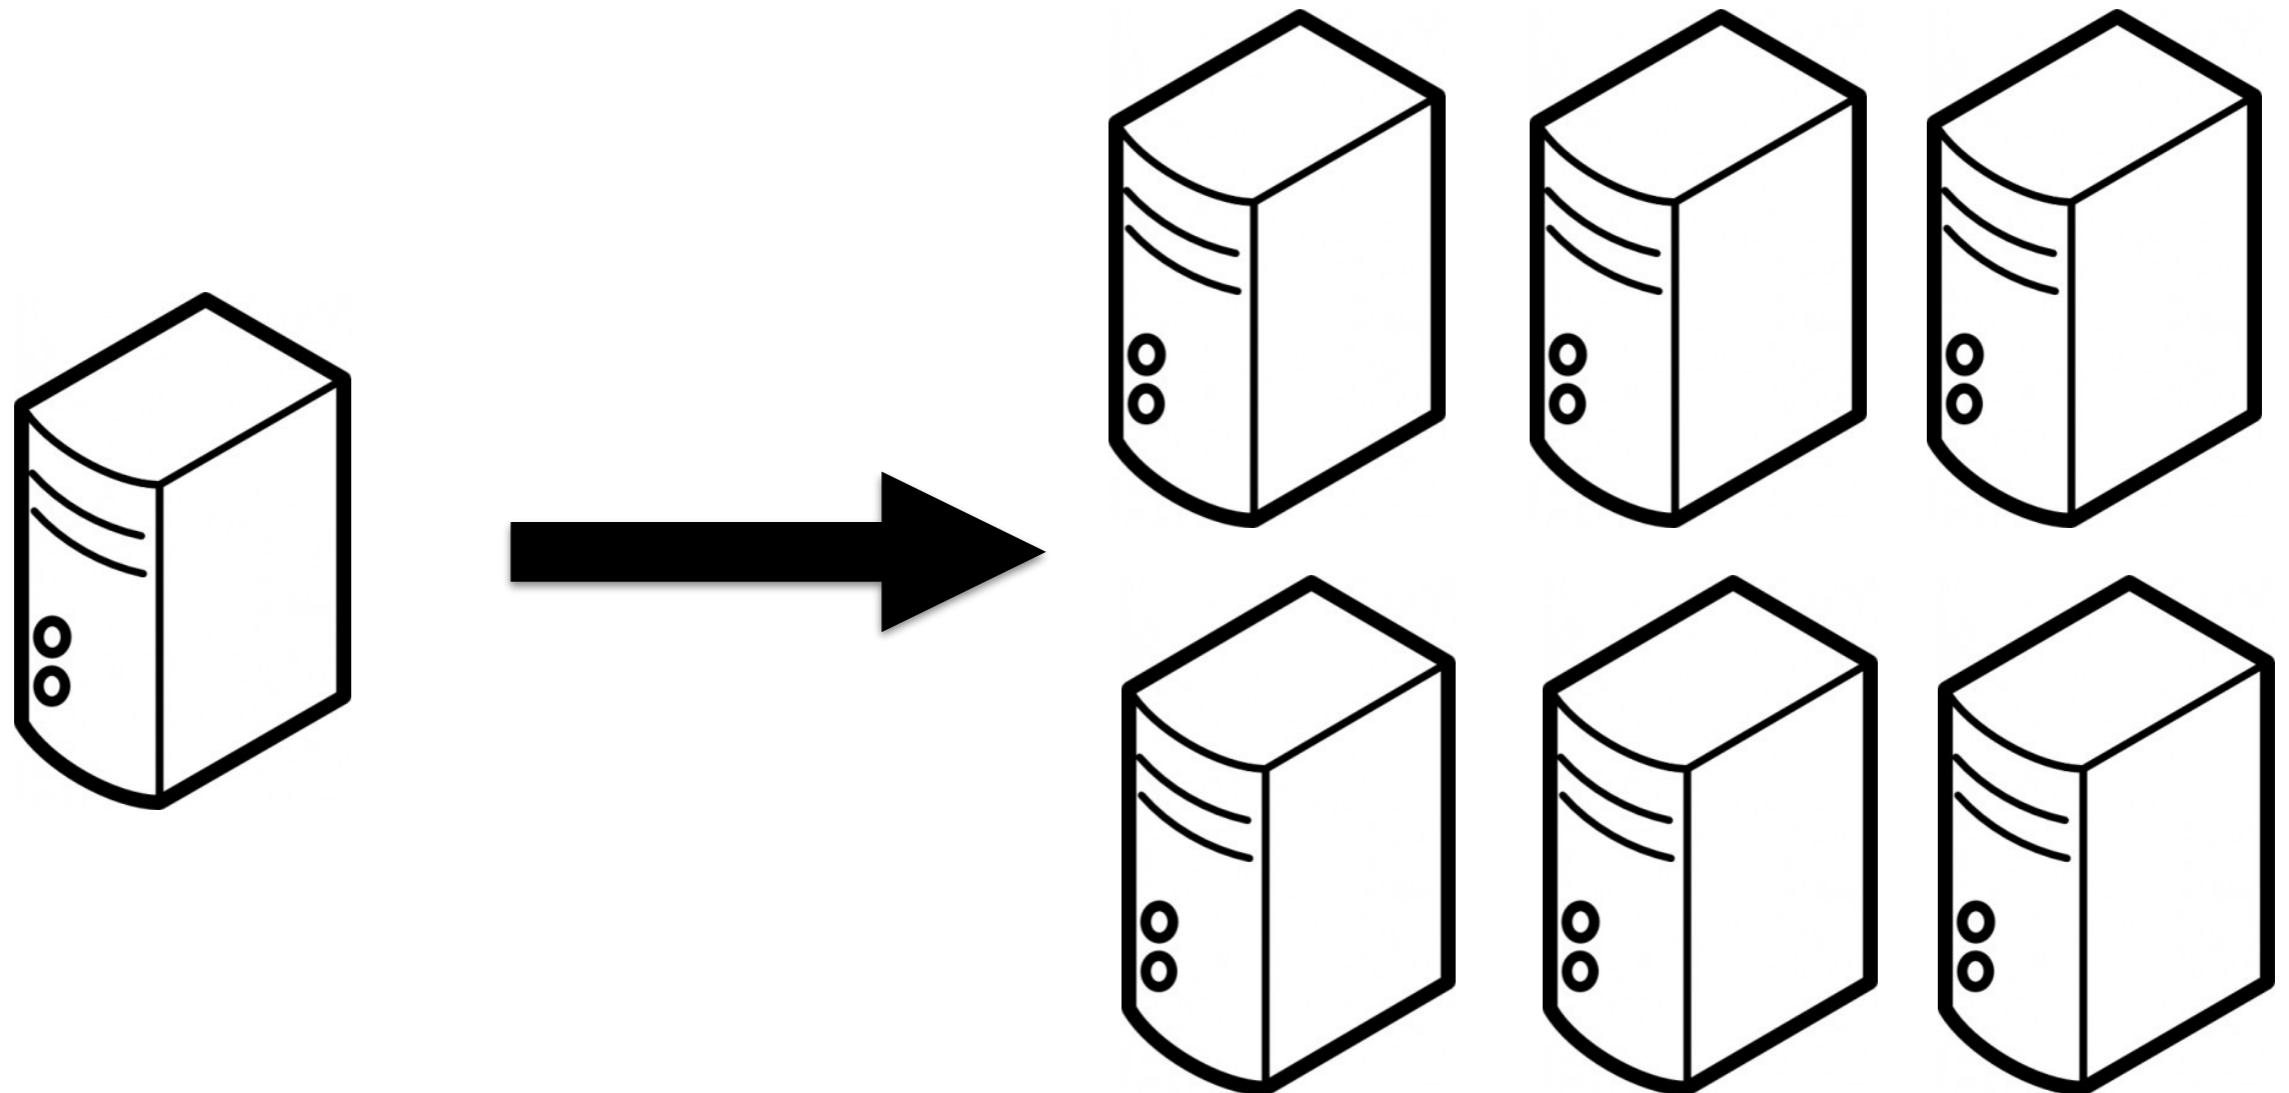
Focus on horizontal scalable

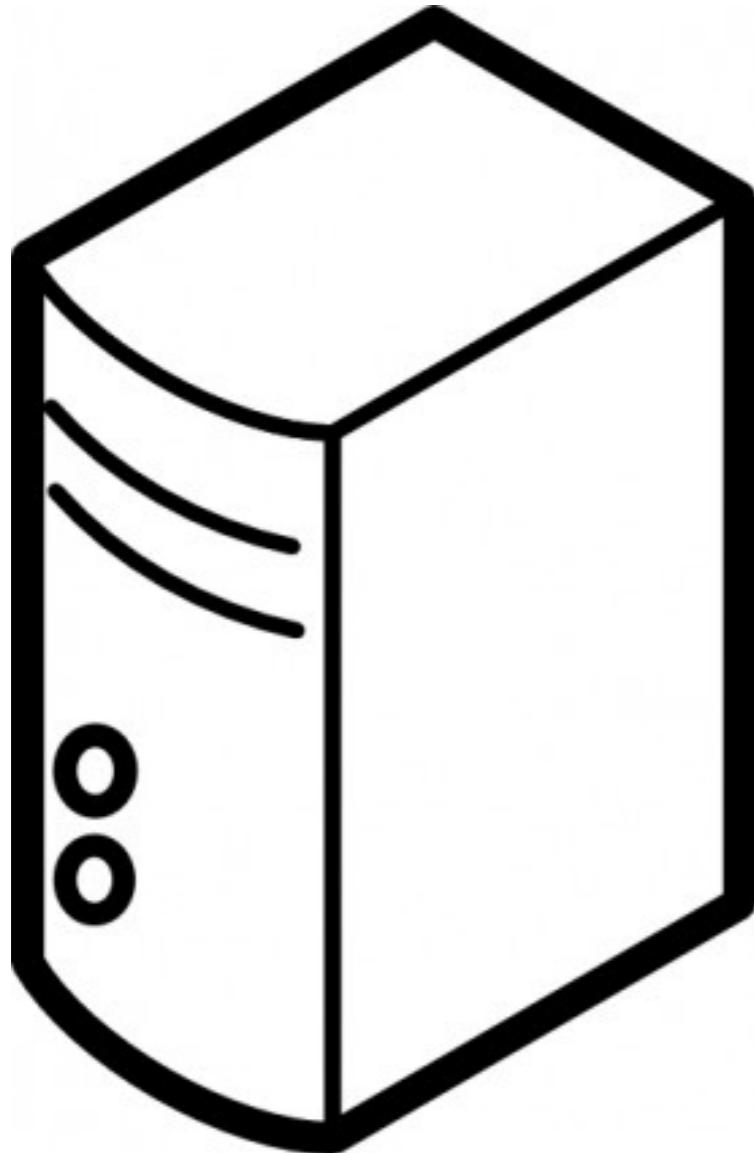


Vertical scaling

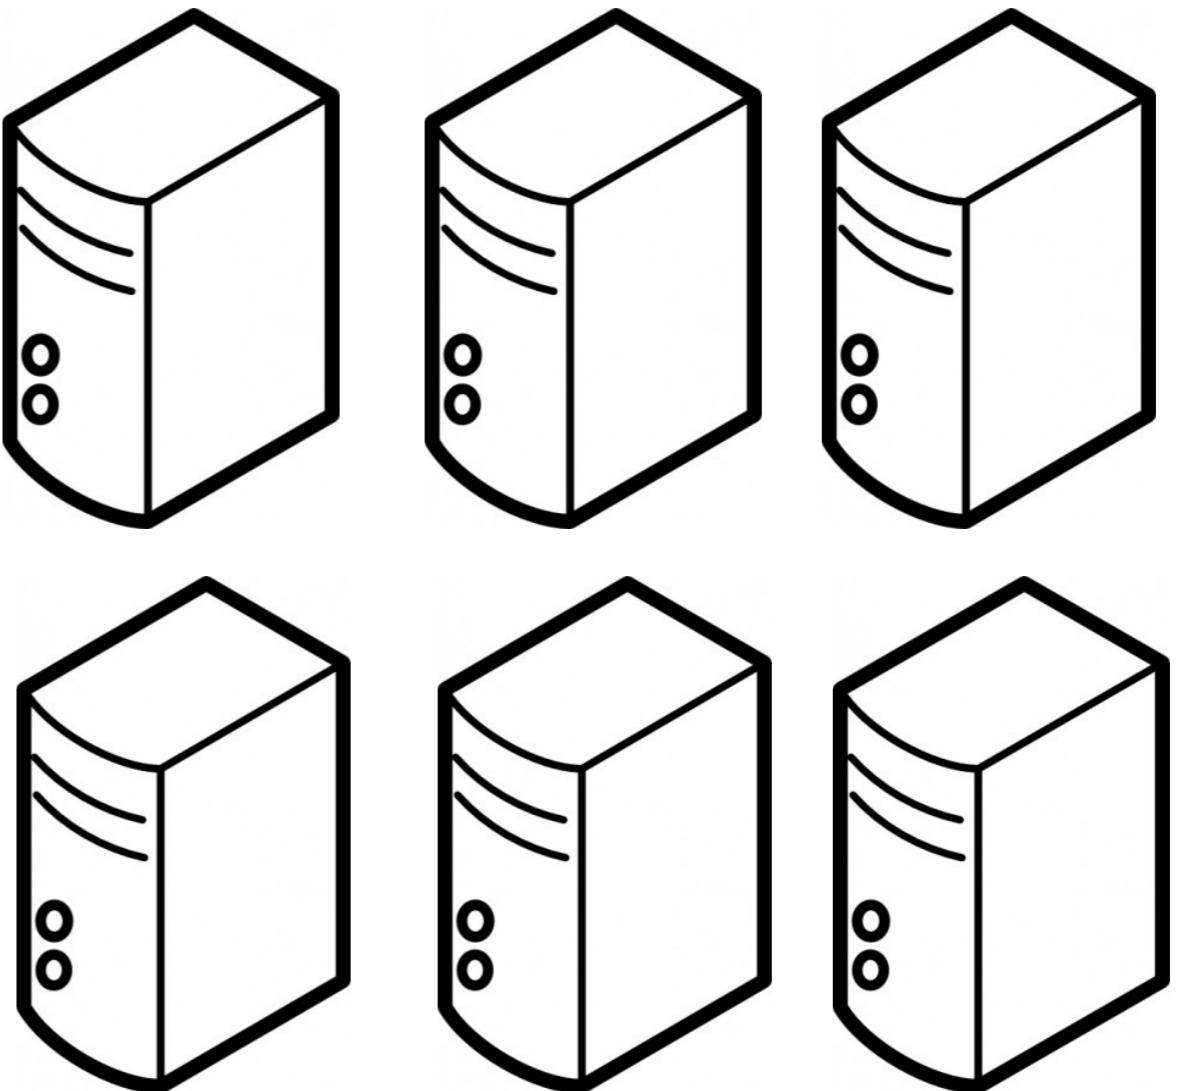


Horizontal scaling





VS



Characteristics of NoSQL

Non-relational

Open-source

Cluster-friendly

Schema less

For new web app

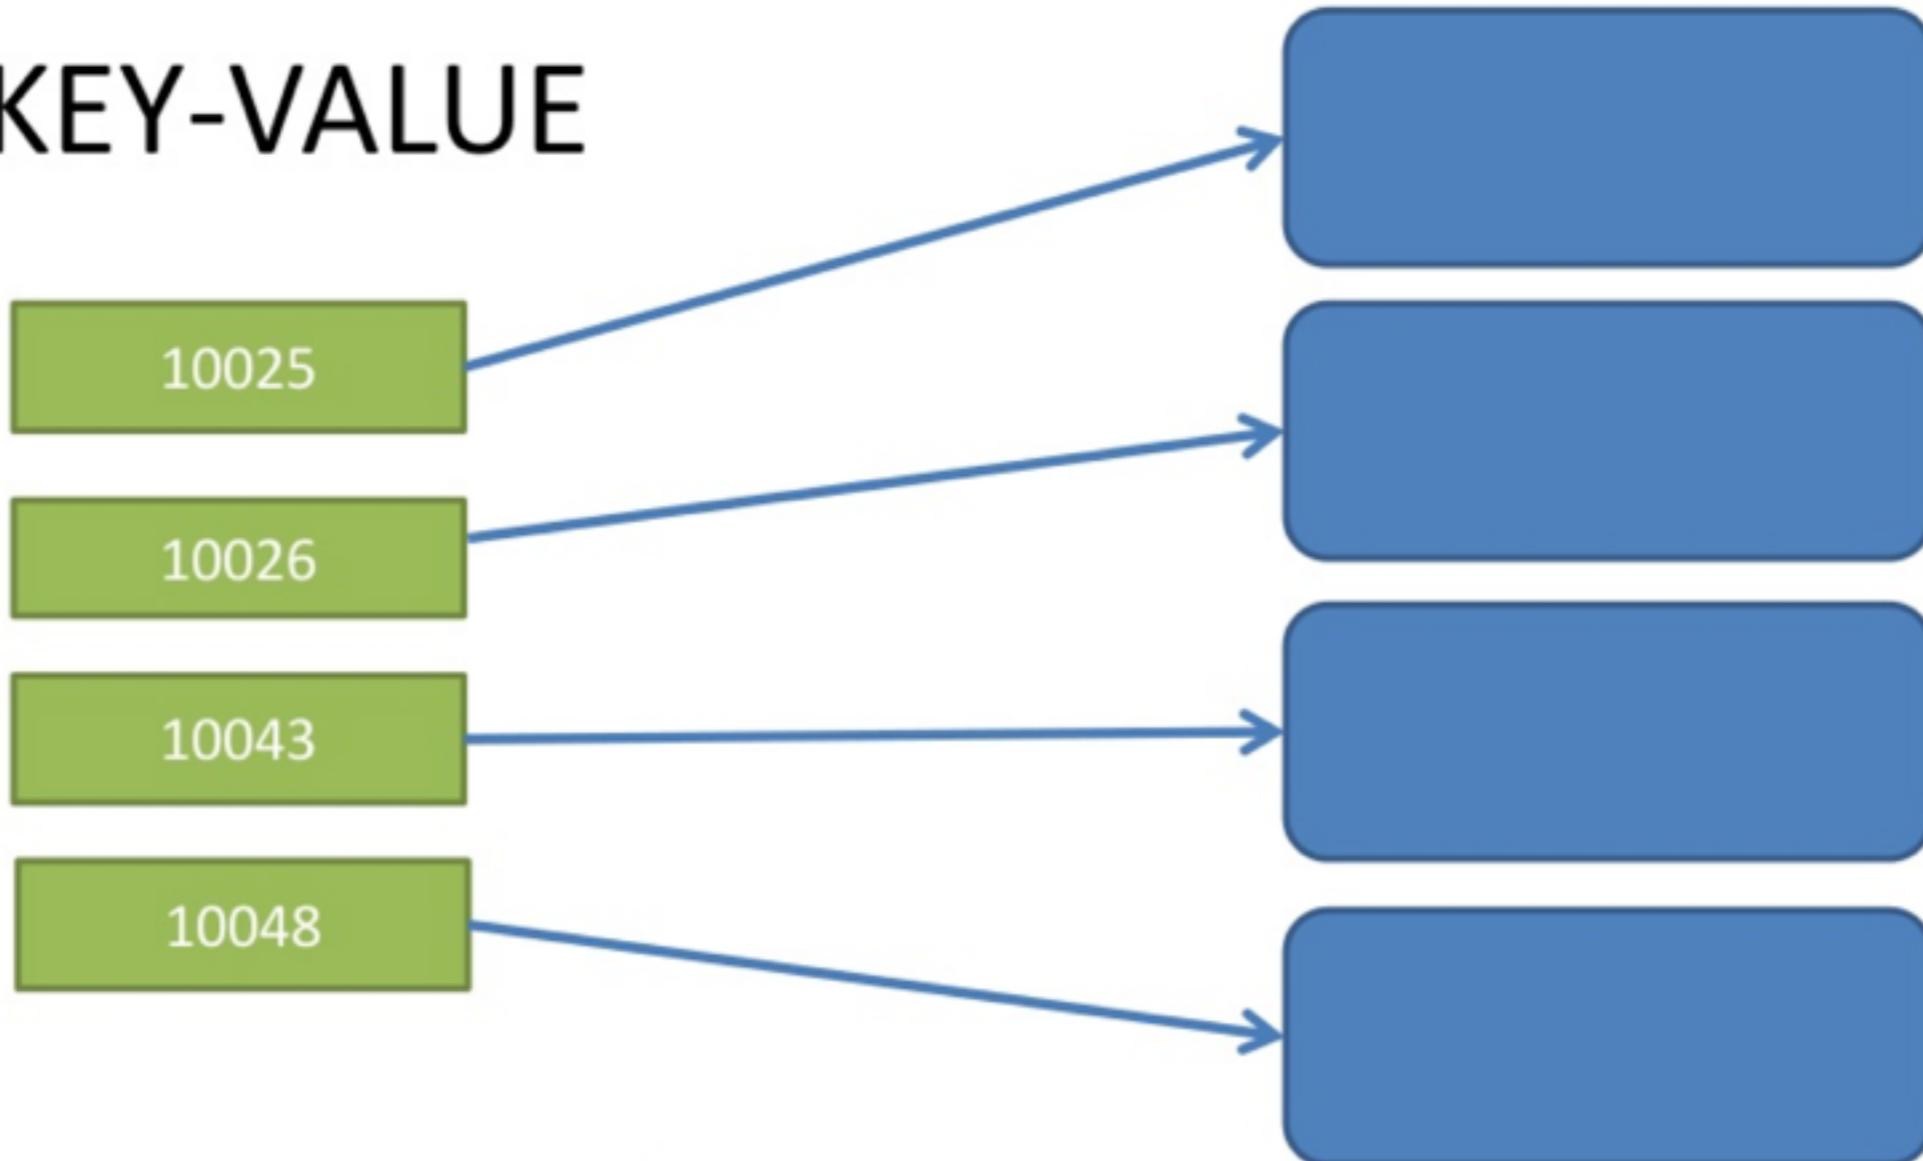


Data model of NoSQL

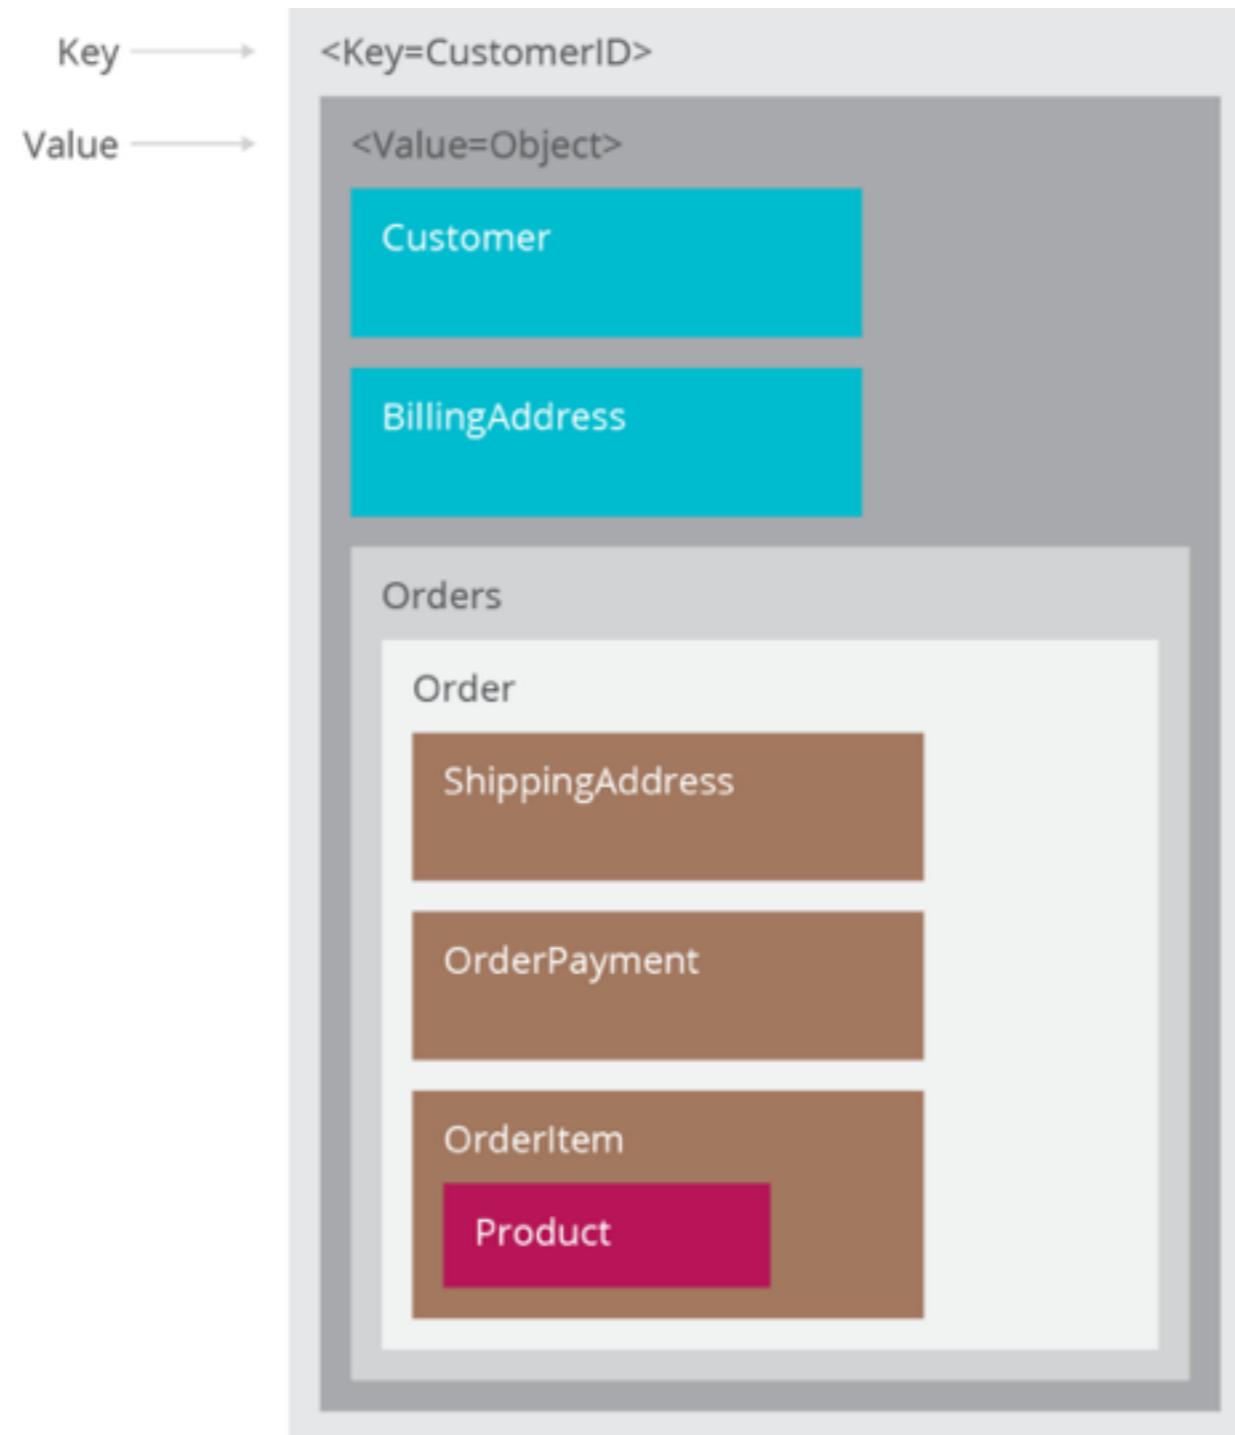


Key-value

KEY-VALUE



Document



Column-family

Column Family

Row

Row KeyX

Column1

name1:value1

Column2

name2:value2

ColumnN

nameN:valueN

Row

Row KeyY

Column1

name1:value1

Column9

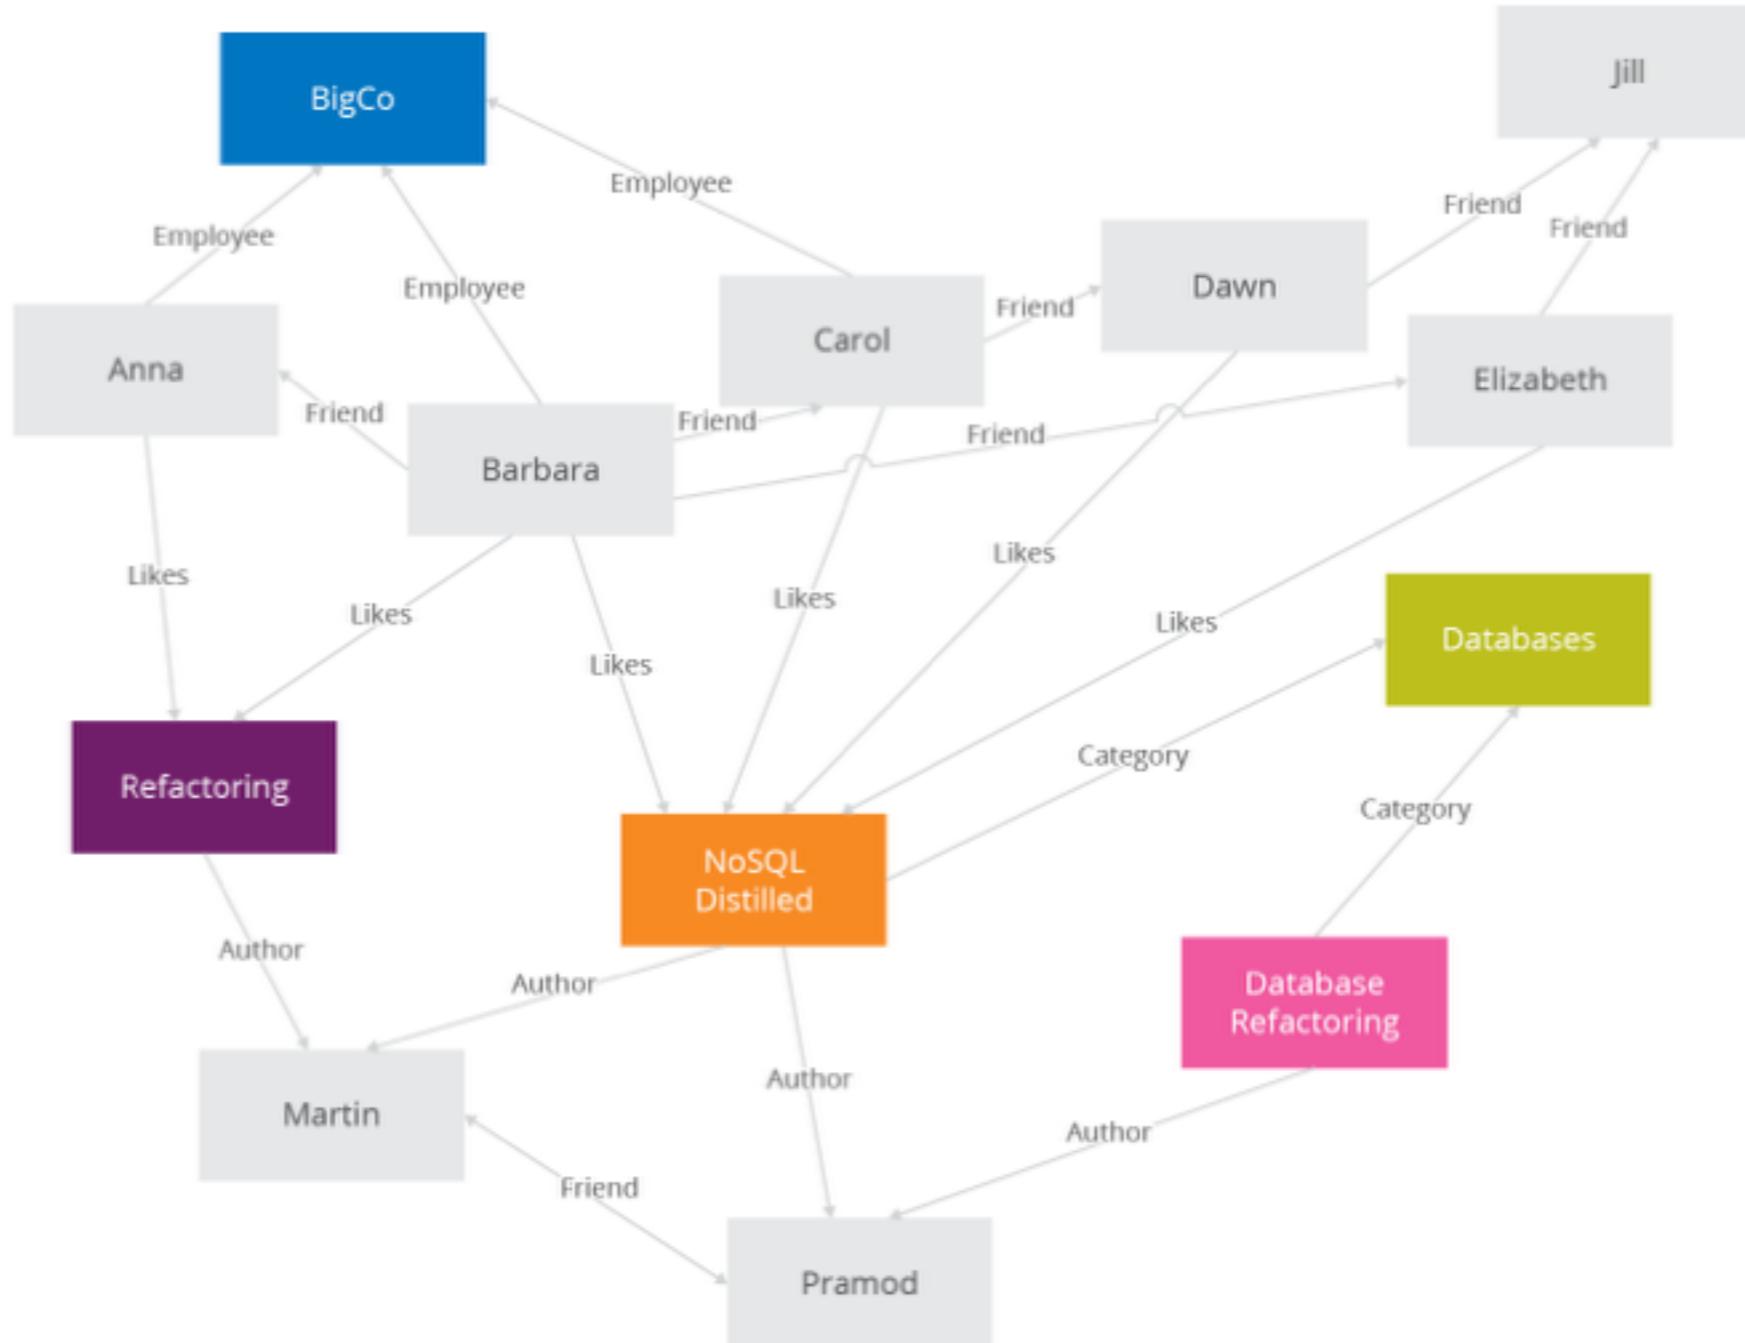
name9:value9

ColumnN

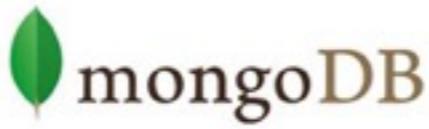
nameN:valueN



Graph



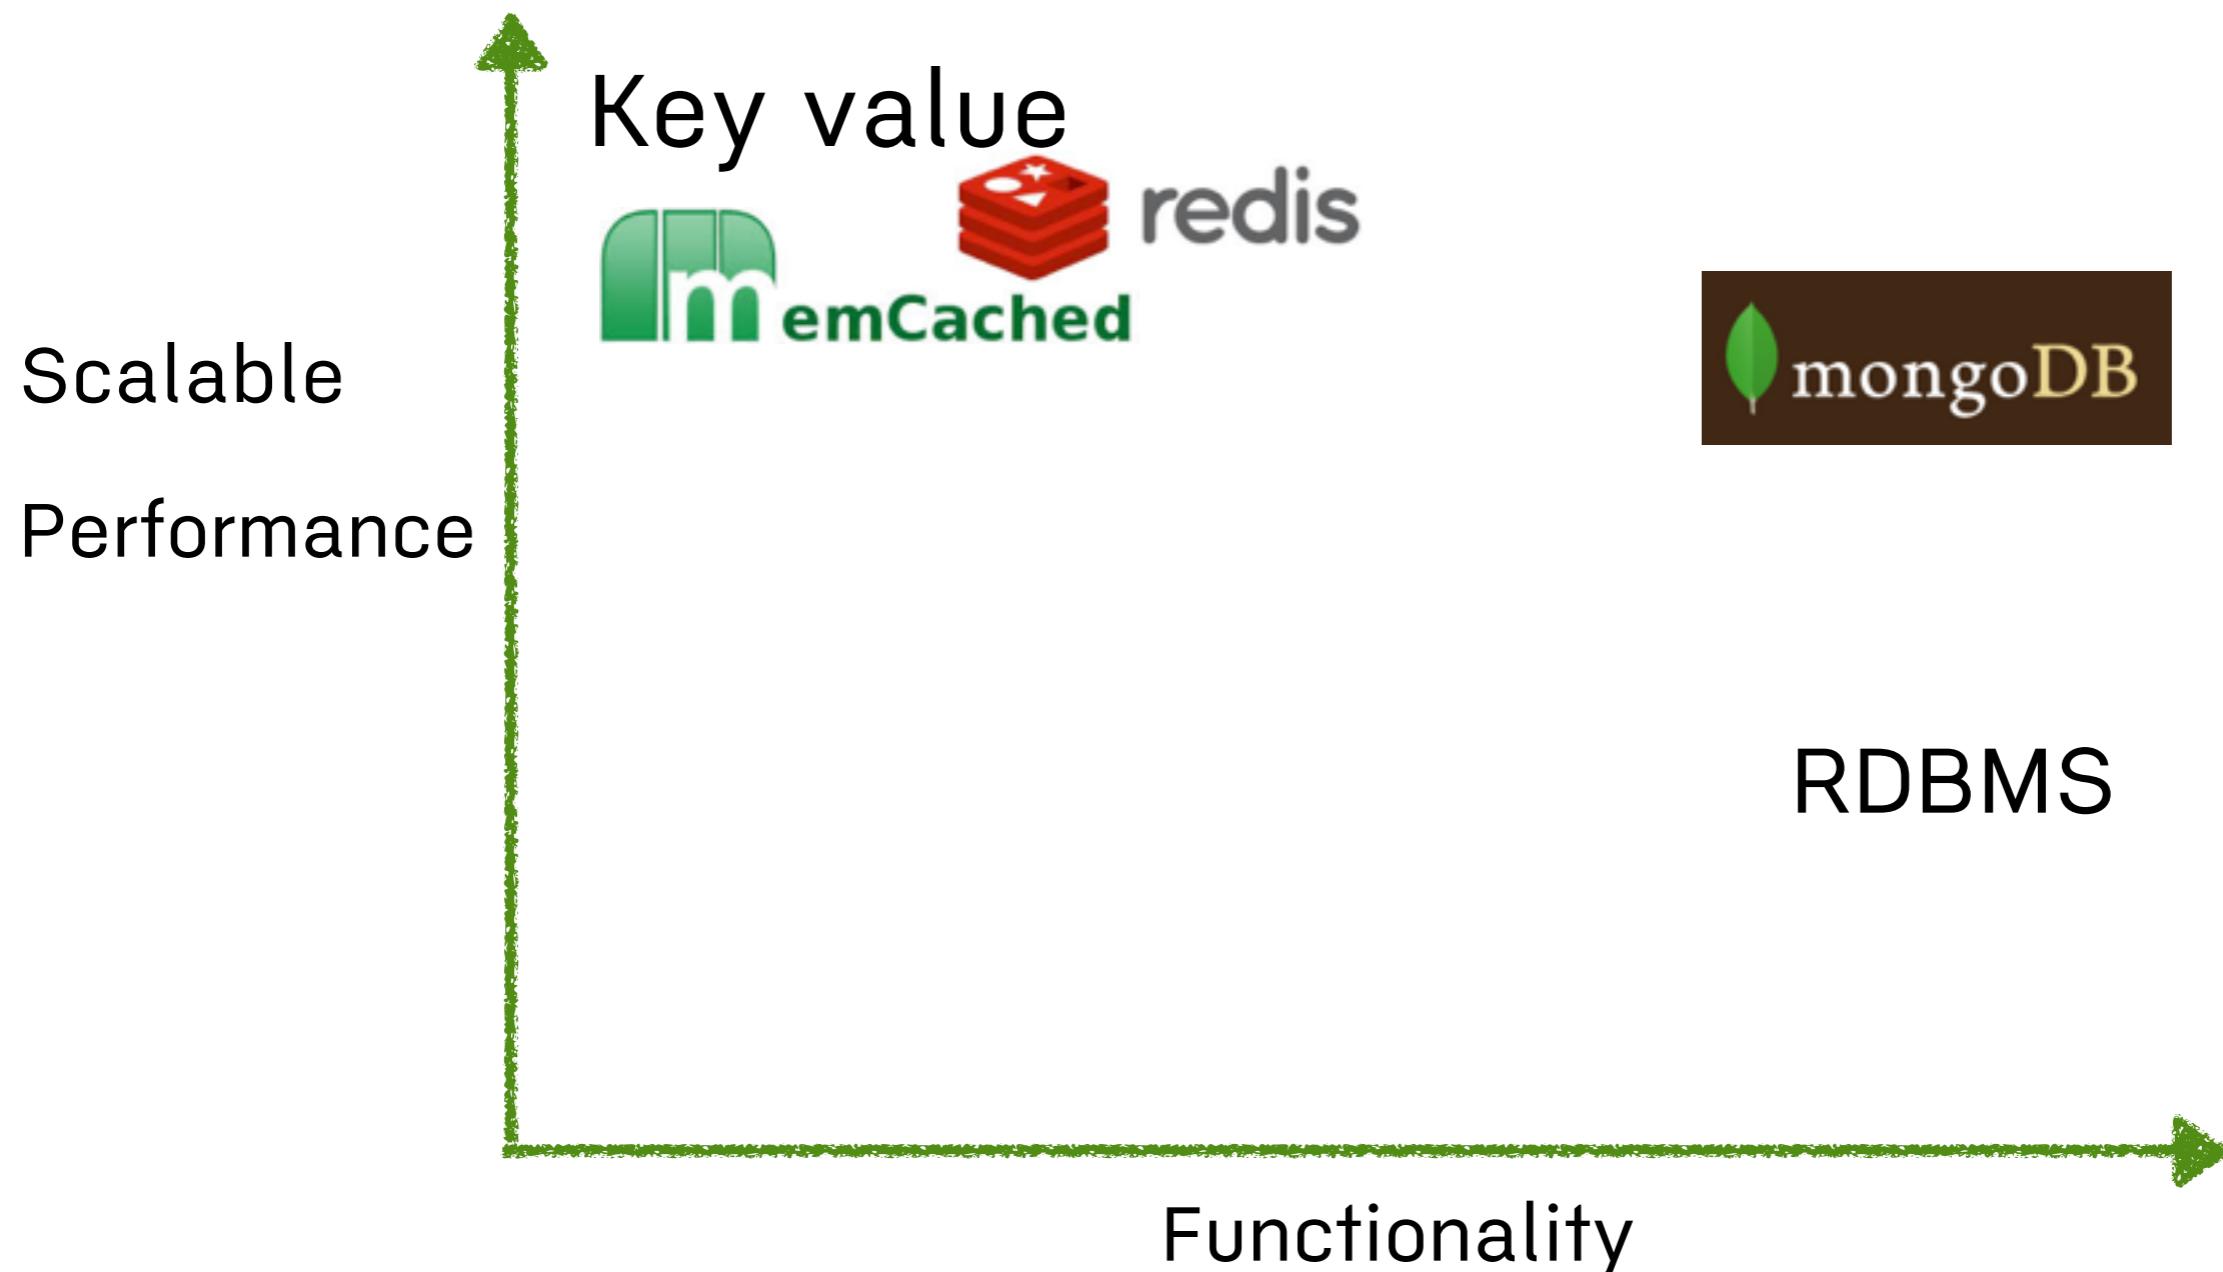
Data model of NoSQL

Document Database	Graph Databases
   	 
Wide Column Stores	Key-Value Databases
   	  HYPERTABLE INC  Cassandra  Apache HBASE Amazon SimpleDB

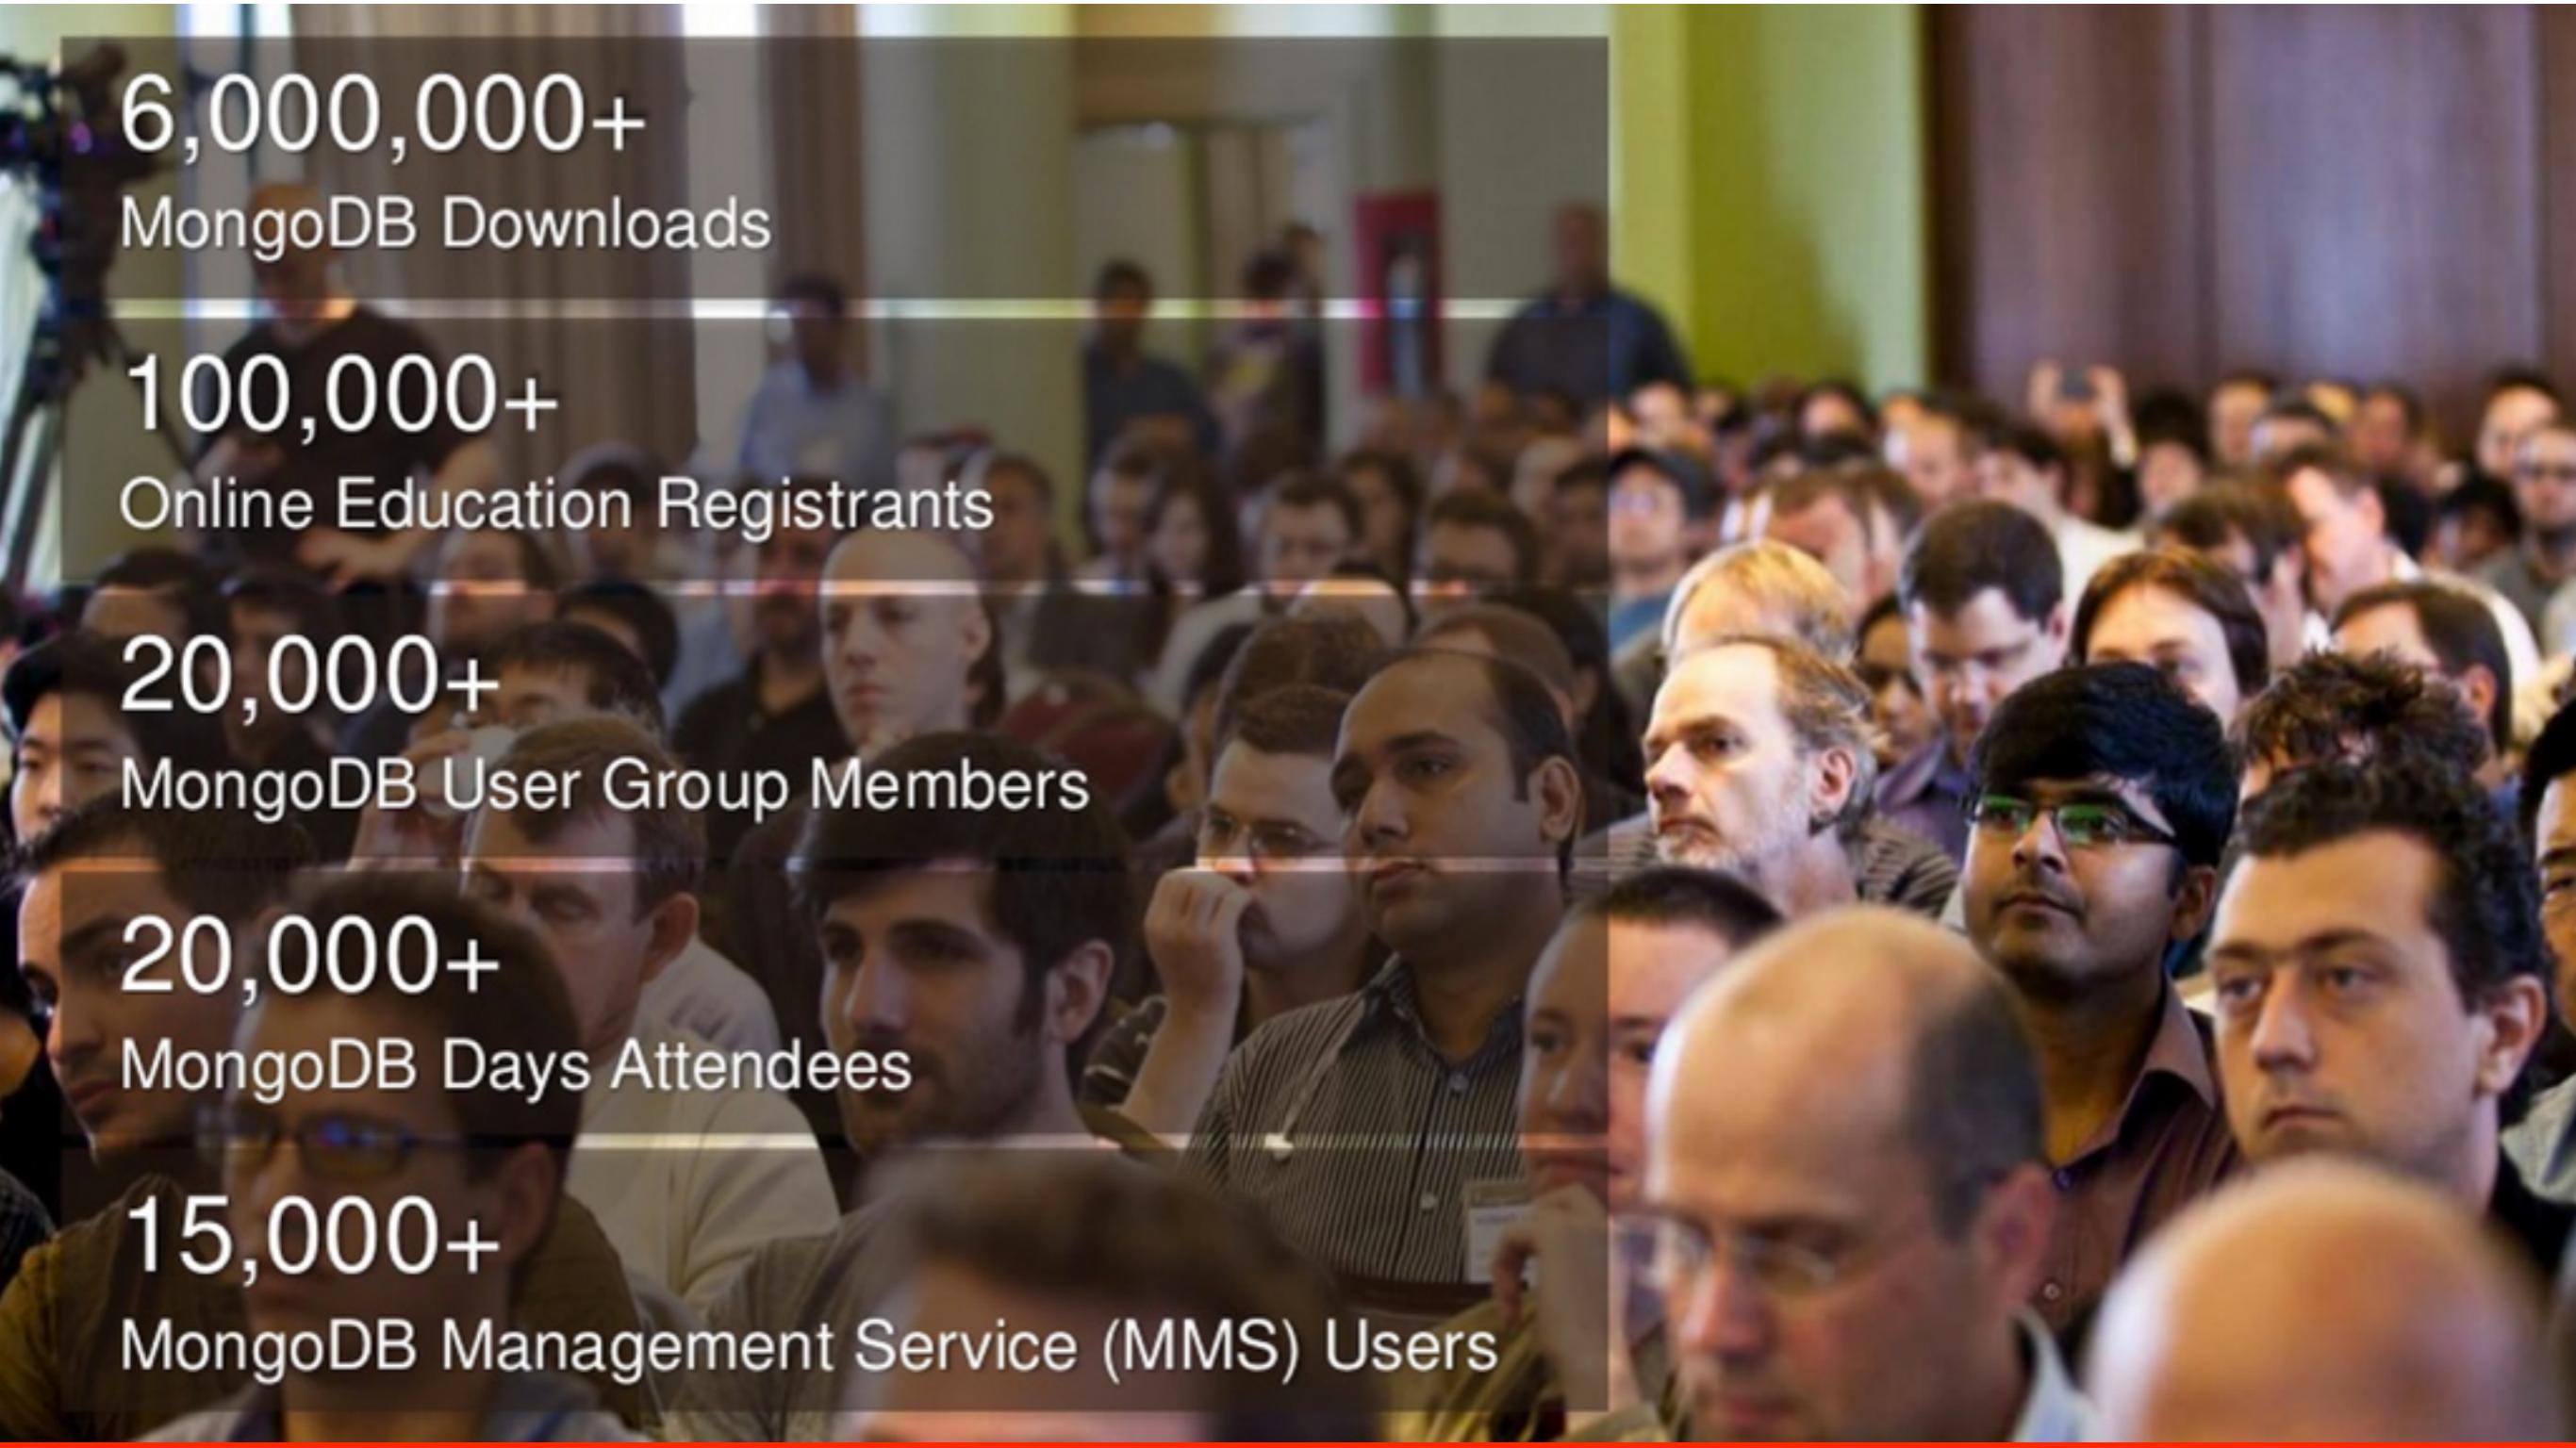
@cloudbit <http://www.aryannava.com>



NoSQL



Community of MongoDB



6,000,000+
MongoDB Downloads

100,000+
Online Education Registrants

20,000+
MongoDB User Group Members

20,000+
MongoDB Days Attendees

15,000+
MongoDB Management Service (MMS) Users



303 systems in ranking, April 2016

Rank			DBMS	Database Model	Score		
Apr 2016	Mar 2016	Apr 2015			Apr 2016	Mar 2016	Apr 2015
1.	1.	1.	Oracle	Relational DBMS	1467.53	-4.48	+21.40
2.	2.	2.	MySQL +	Relational DBMS	1370.11	+22.39	+85.53
3.	3.	3.	Microsoft SQL Server	Relational DBMS	1135.05	-1.45	-14.07
4.	4.	4.	MongoDB +	Document store	312.44	+7.11	+33.85
5.	5.	5.	PostgreSQL	Relational DBMS	303.73	+4.10	+35.41
6.	6.	6.	DB2	Relational DBMS	184.08	-3.85	-13.56
7.	7.	7.	Microsoft Access	Relational DBMS	131.97	-3.06	-10.22
8.	8.	8.	Cassandra +	Wide column store	129.67	-0.66	+24.78
9.	9.	↑ 10.	Redis +	Key-value store	111.24	+5.02	+16.69
10.	10.	↓ 9.	SQLite	Relational DBMS	107.96	+2.19	+5.67
11.	11.	↑ 14.	Elasticsearch +	Search engine	82.58	+2.41	+17.92
12.	12.	↓ 11.	SAP Adaptive Server	Relational DBMS	73.32	-3.33	-13.37
13.	13.	13.	Teradata	Relational DBMS	72.26	-1.81	+2.00
14.	14.	↓ 12.	Solr	Search engine	66.02	-3.35	-15.98
15.	15.	15.	HBase	Wide column store	51.49	-0.92	-9.65
16.	16.	↑ 17.	Hive	Relational DBMS	49.08	-1.43	+6.33
17.	17.	↓ 16.	FileMaker	Relational DBMS	46.10	-1.83	-5.72

http://db-engines.com/en/ranking_trend



One size fit all !!



what is JSON ?

JavaScript Object Notation
Key-value



JSON document

```
{  
    "name": "Somkiat",  
    "age": 30,  
    "province": "Bangkok"  
    "country": "Thailand"  
}
```



JSON document

```
{  
    "name": "MongoDB book,  
    "page": 300,  
    "author": ["author 1", "author 2"]  
}
```



Schema less

Two documents don't have the same schema !!



MongoDB Book

```
{  
  "name": "MongoDB book,  
  "page": 300,  
  "author": ["author 1", "author 2"]  
}
```



NoSQL Book

```
{  
  "name": "NoSQL book,  
  "page": 200,  
  "author": ["author 1", "author 2"],  
  "discount}
```



Installation

<https://www.mongodb.com/download-center?jmp=nav#community>



Supported platforms

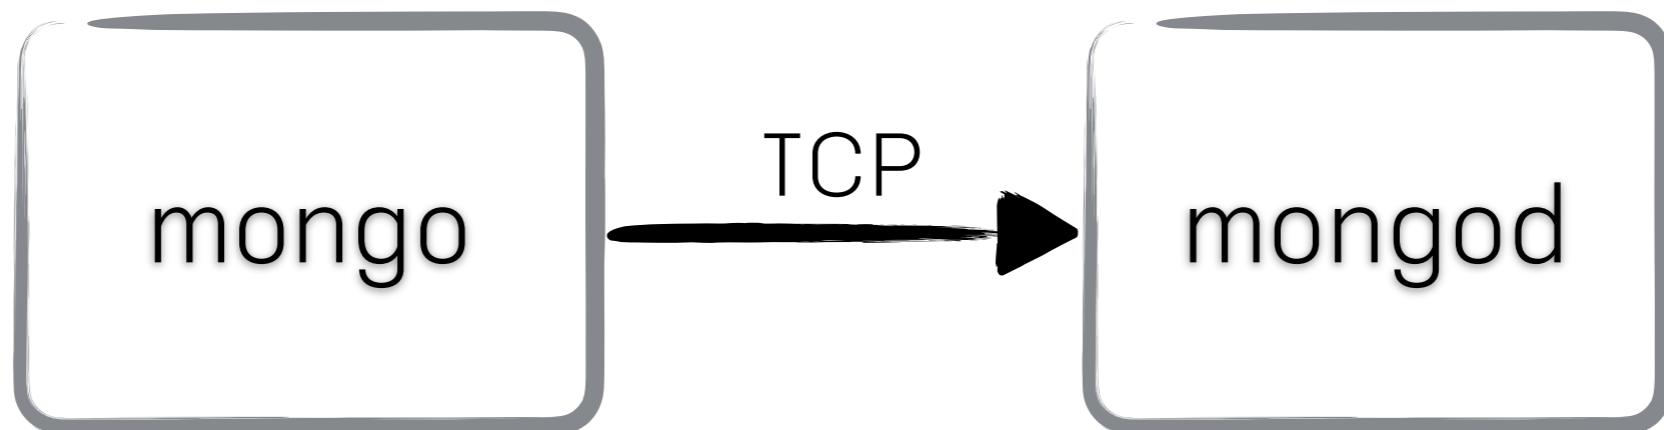
Platform	3.2	3.0	2.6	2.4	2.2
Amazon Linux	✓	✓	✓	✓	✓
Debian 7	✓	✓	✓	✓	✓
Fedora 8+			✓	✓	✓
RHEL/CentOS 6.2+	✓	✓	✓	✓	✓
RHEL/CentOS 7.0+	✓	✓	✓		
SLES 11	✓	✓	✓	✓	✓
SLES 12	✓				
Solaris 64-bit	✓	✓	✓	✓	✓
Ubuntu 12.04	✓	✓	✓	✓	✓
Ubuntu 14.04	✓	✓	✓		
Microsoft Azure	✓	✓	✓	✓	✓
Windows Vista/Server 2008R2/2012+	✓	✓	✓	✓	✓
OSX 10.7+	✓	✓	✓	✓	



Let's start

mongod = MongoDB database

mongo = MongoDB shell



Start MongoDB

```
$mongod --dbpath d:\mongodb\data
```

```
[initandlisten]
[initandlisten] ** WARNING: soft rlimits too low. Number of files
[HostnameCanonicalizationWorker] Starting hostname canonicalization worker
[initandlisten] Initializing full-time diagnostic data capture...
[initandlisten] waiting for connections on port 27017
```



Mongo shell

\$mongo

```
MongoDB shell version: 3.2.6
connecting to: test
Server has startup warnings:
2016-05-11T21:59:35.256+0700 I CONTROL  [initandlisten]
2016-05-11T21:59:35.256+0700 I CONTROL  [initandlisten]
** WARNING: soft rlimits too low. Number of files is 256
, should be at least 1000
> █
```



RDBMS vs MongoDB concept

RDBMS	MongoDB
database	database
table	collection
row	document/JSON
column	field
index	index
table joins	embedded documents and linking
primary key	primary key
aggregation	aggregation pipeline



Create database

>use store



Drop database

```
>db.dropDatabase()
```



Insert new document

```
>db.book.insert(  
  {  
    "name": "Book 01",  
    "page": 250  
  }  
)
```



Count all documents

```
>db.book.count()
```



Query for all documents

```
>db.book.find()
```



Query by a field

```
>db.book.find( {"name": "Book 01"} )
```



Count by query

```
>db.book.find( {"name": "Book 01"} ).count()
```



Query with operator

```
>db.book.find( {"page":{ $gt: 200} } )
```

```
>db.book.find( {"page":{ $lt: 250} } )
```



Query with limit

```
>db.book.find().limit(2)
```



Query with skip

```
>db.book.find().limit(1).skip(1)
```



Sort query results

```
>db.book.find().sort({"page": 1})
```

```
>db.book.find().sort({"page": -1})
```



Update a document

```
>db.book.update(  
  { "name": "Book 01" },  
  {  
    $set: { "author": "Somkiat" },  
    $currentDate: { "lastModified": true }  
  }  
)
```



Update a document

```
>db.book.update(  
  { "name": "Book 01"},  
  {  
    $set: { "author": "Somkiat" },  
    $currentDate: { "lastModified": true }  
  }  
)
```



Update a document

```
>db.book.update(  
  { "name": "Book 01" },  
  {  
    $set: { "author": "Somkiat" },  
    $currentDate: { "lastModified": true }  
  }  
)
```



Update multiple documents

```
>db.book.update(  
  {"page": { $gt: 200 }},  
  {  
    $set: { "author": "New" },  
    $currentDate: { "lastModified": true }  
  },  
  { multi: true}  
)
```



Replace a document

```
>db.book.update(  
  { "name": "Book 01" },  
  {  
    "name": "Update book 01",  
    "page": 500  
  }  
)
```



Remove all documents

```
>db.book.remove( { } )
```



Remove some documents

```
>db.book.remove({"name": "Book 02"})
```



Remove only one document

```
>db.book.remove( { }, { justOne: true } )
```



Drop a collection

```
>db.book.drop()
```



Workshop

POST

COMMENTS

Create a post

Add a comment to post

01_blog.txt



Create database

>use blog



Create a new post

```
>post =  
{   _id: "1",  
    topic: "Hello my blog",  
    text: "this is my first blog",  
    tag: [ "intro", "mongodb" ],  
    author: "somkiat",  
    date: new Date()  
}
```

```
>db.posts.insert(post)
```



Add a comment to post

```
>comment =  
{   text: "Cool post!",  
    author: "somkiat",  
    date: new Date()  
}  
  
>db.posts.update(  
  { _id: "1"},  
  {  
    $push: { comments : comment }  
  }  
)
```



Query data

>db.posts.find()

>db.posts.find().pretty()



Workshop with query

POST

COMMENTS

Post by author

Last 10 posts

Posts from 1 May 2016

Posts with a tag

Paging



Post by author

```
>db.posts.find( {author:"somkiat"} )
```



Last 10 posts

```
>db.posts.find().sort( {date: -1} ).limit(10)
```



Posts from 1 May 2016

```
>may_1_2016 = new Date(2016, 5, 1)
```

```
>db.posts.find(  
  {  
    date: { $gte: may_1_2016}  
  }  
)
```



Posts with tag

```
>db.posts.find( {tag:"mongodb"} )
```



Paging

```
>page = 0
```

```
>page_size = 10
```

```
>db.posts.find().limit(page_size).skip( page *  
page_size )
```



Analyze Query Performance



บริษัท สยามช่างนาฏกิจ จำกัด และเพื่อนพ้องน้องพี่

Posts with tag

```
>db.posts.find( {tag:"mongodb"} ).explain()
```

```
"queryPlanner" : {  
    "plannerVersion" : 1,  
    "namespace" : "blog.posts",  
    "indexFilterSet" : false,  
    "parsedQuery" : {  
        "tag" : {  
            "$eq" : "mongodb"  
        }  
    },  
    "winningPlan" : {  
        "stage" : "COLLSCAN",  
        "filter" : {  
            "tag" : {  
                "$eq" : "mongodb"  
            }  
        },  
        "direction" : "forward"  
    },  
    "rejectedPlans" : [ ]  
},
```



Create new index

```
>db.posts.createIndex({ tag: 1 })
```



Posts with tag

```
>db.posts.find( {tag:"mongodb"} ).explain()
```

```
"winningPlan" : {  
    "stage" : "FETCH",  
    "inputStage" : {  
        "stage" : "IXSCAN",  
        "keyPattern" : {  
            "tag" : 1  
        },  
        "indexName" : "tag_1",  
        "isMultiKey" : true,  
        "isUnique" : false,  
        "isSparse" : false,  
        "isPartial" : false,  
        "indexVersion" : 1,  
        "direction" : "forward",  
        "indexBounds" : {  
            "tag" : [  
                "[\"mongodb\", \"mongodb\"]"  
            ]  
        }  
    }  
},
```



Drop index

```
>db.posts.dropIndex({ tag: 1 })
```



Aggregation

03_aggregation.txt



```
SELECT by_user, count(*)  
FROM mycol  
GROUP BY by_user
```



Count number of book by user

```
>db.book.aggregate(  
[  
  {  
    $group :  
      { _id : "$by_user",  
        num_tutorial : {$sum : 1}  
      }  
  }  
]  
)
```

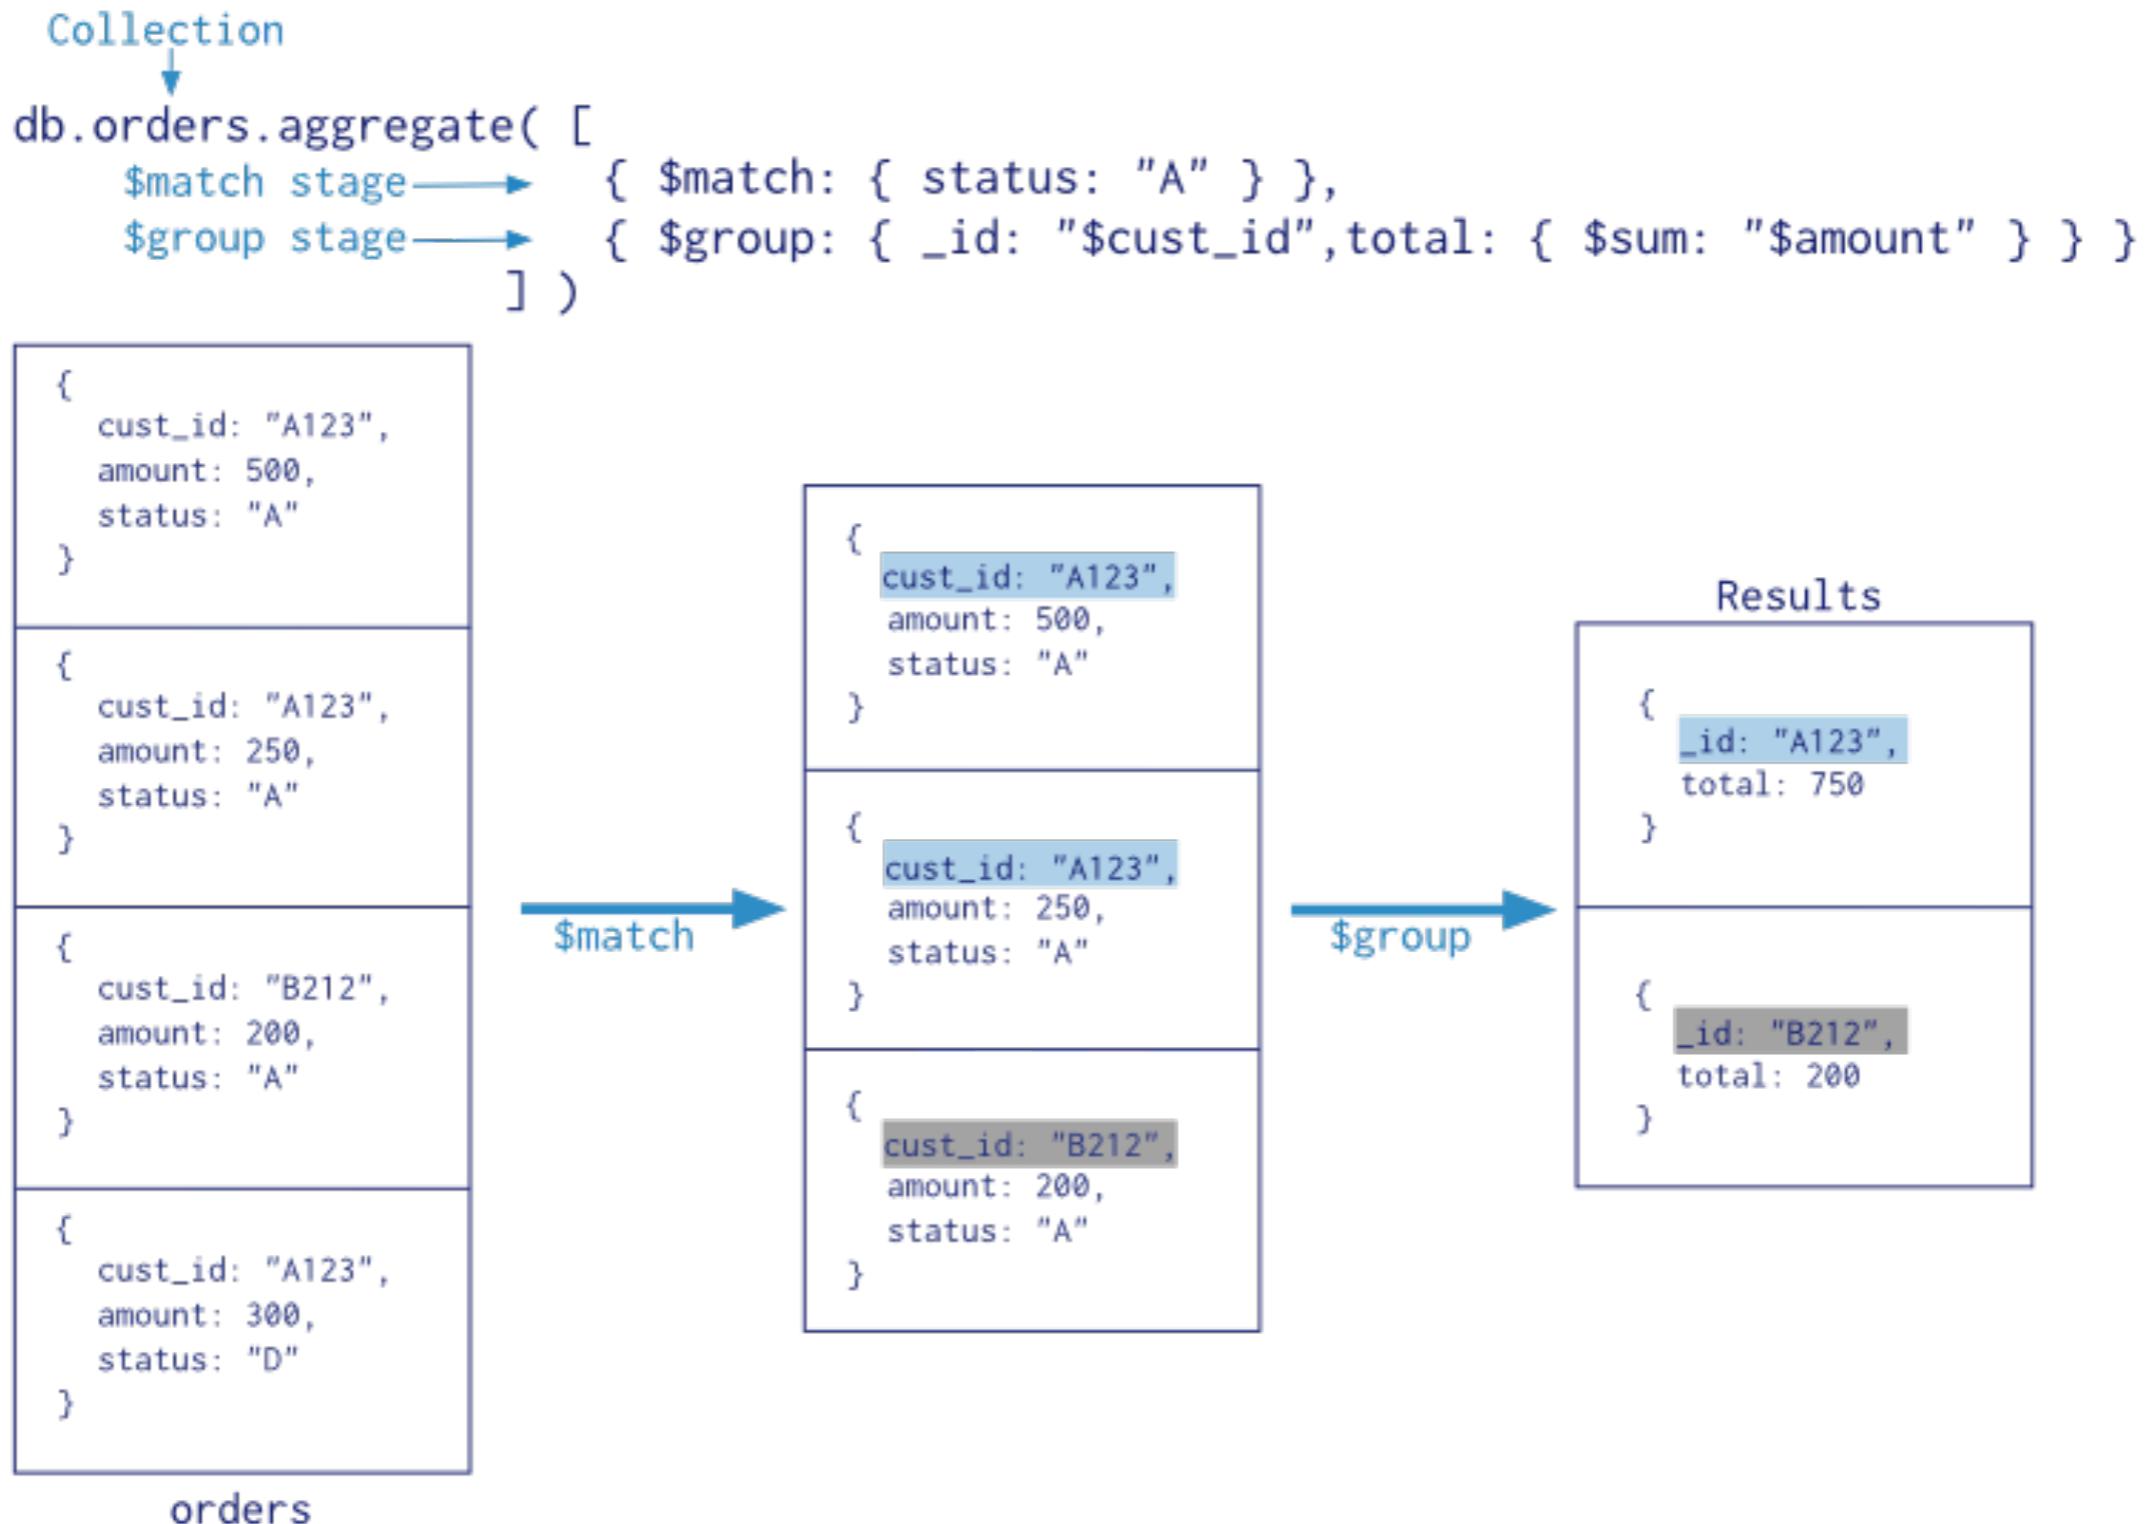


SQL vs Aggregation

SQL	Aggregation
SELECT	\$project
WHERE	\$match
GROUP BY	\$group
HAVING	\$match
ORDER BY	\$sort
LIMIT	\$limit
SUM()	\$sum
COUNT()	\$sum
join	\$lookup



Aggregation pipeline



Workshop

Working with order information

04_order.txt



Data Modeling



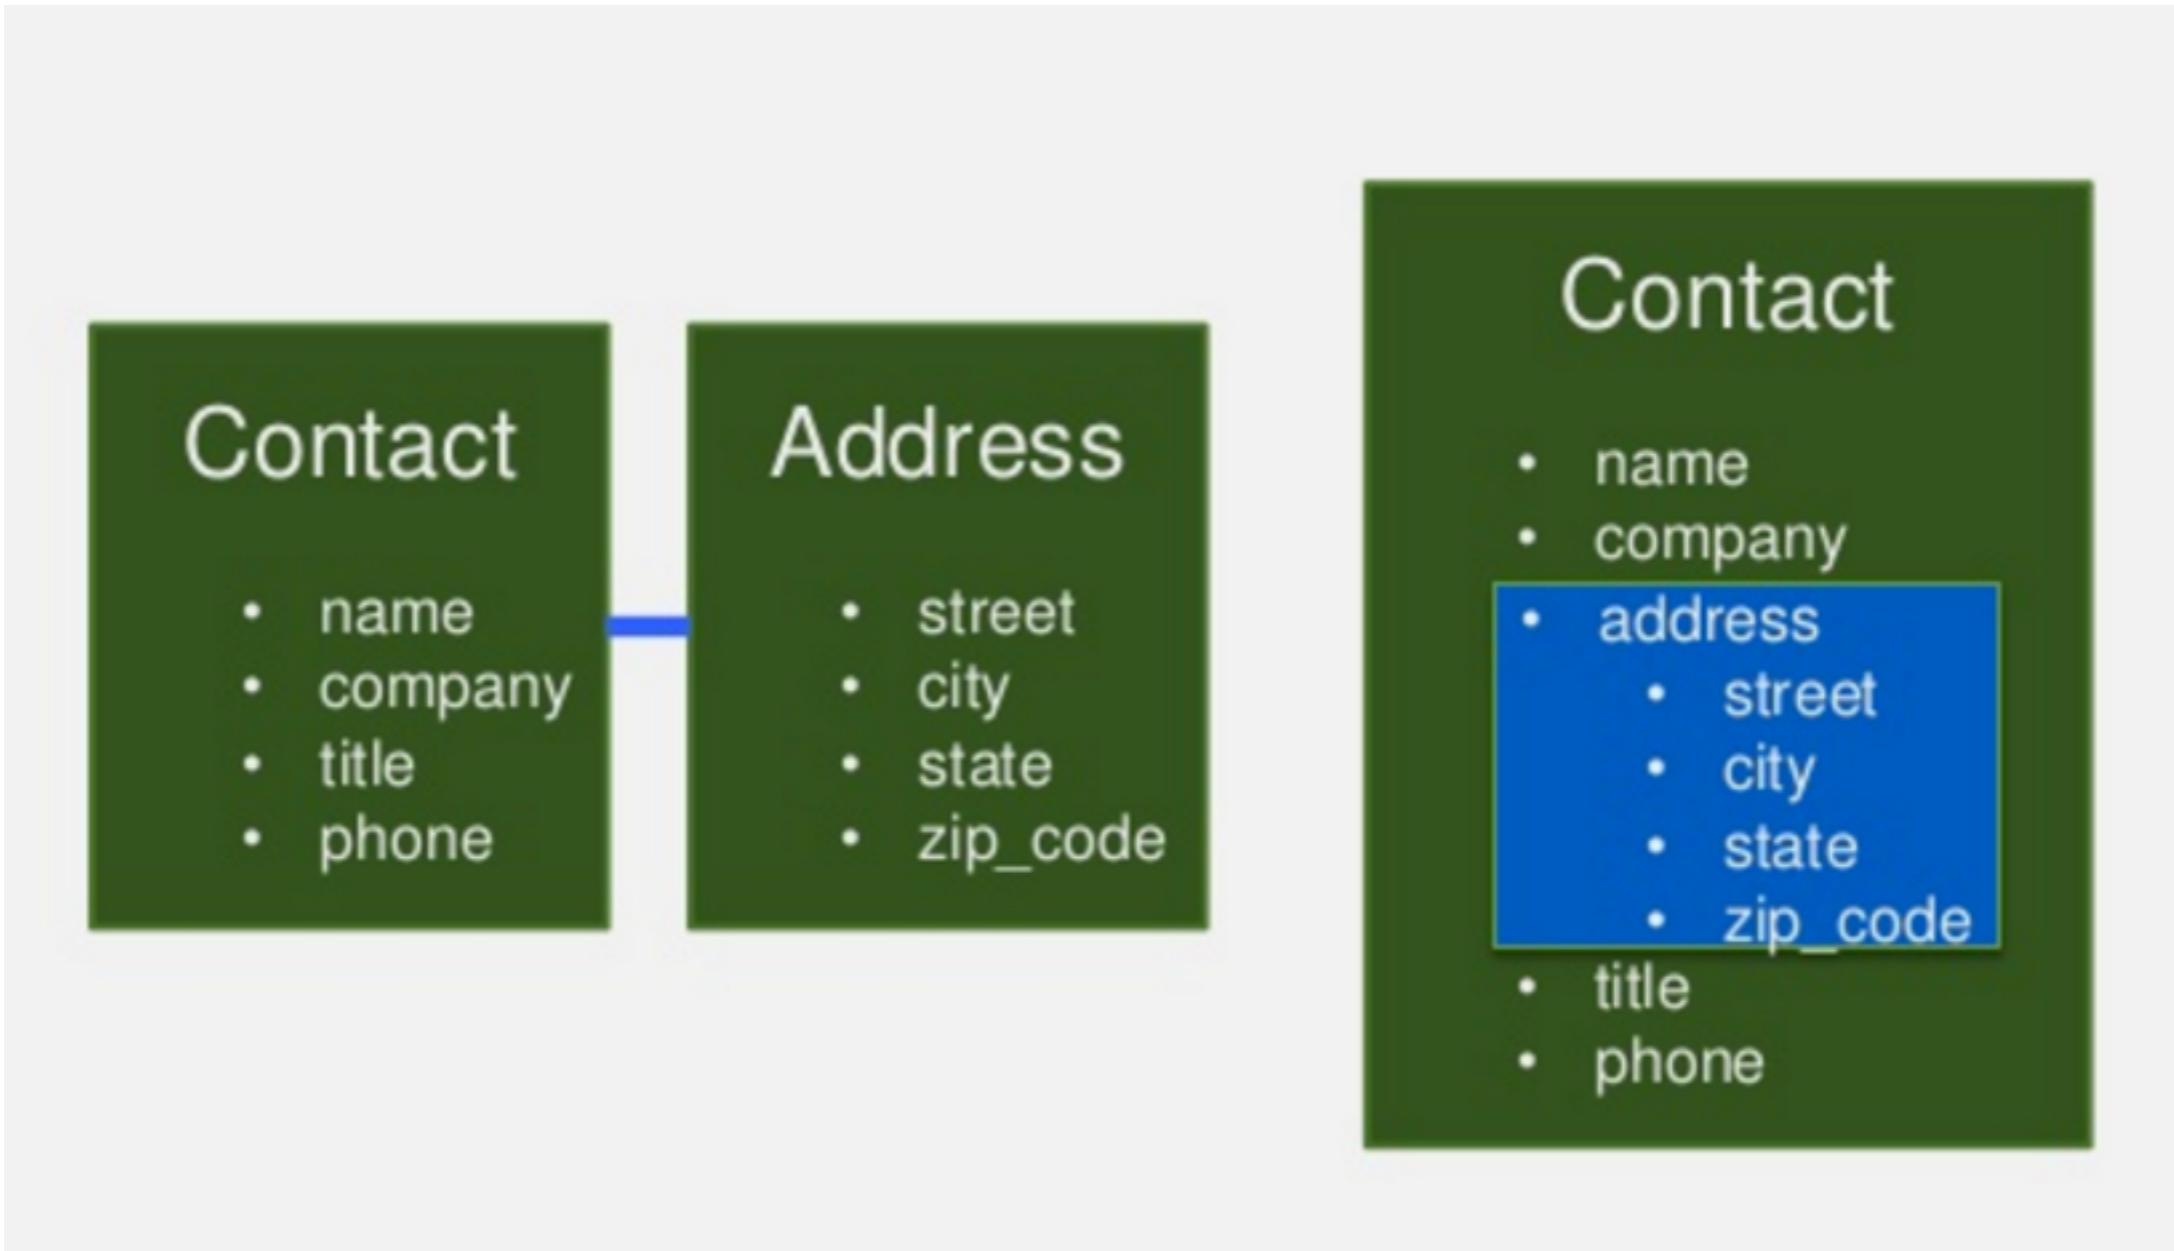
Embedding

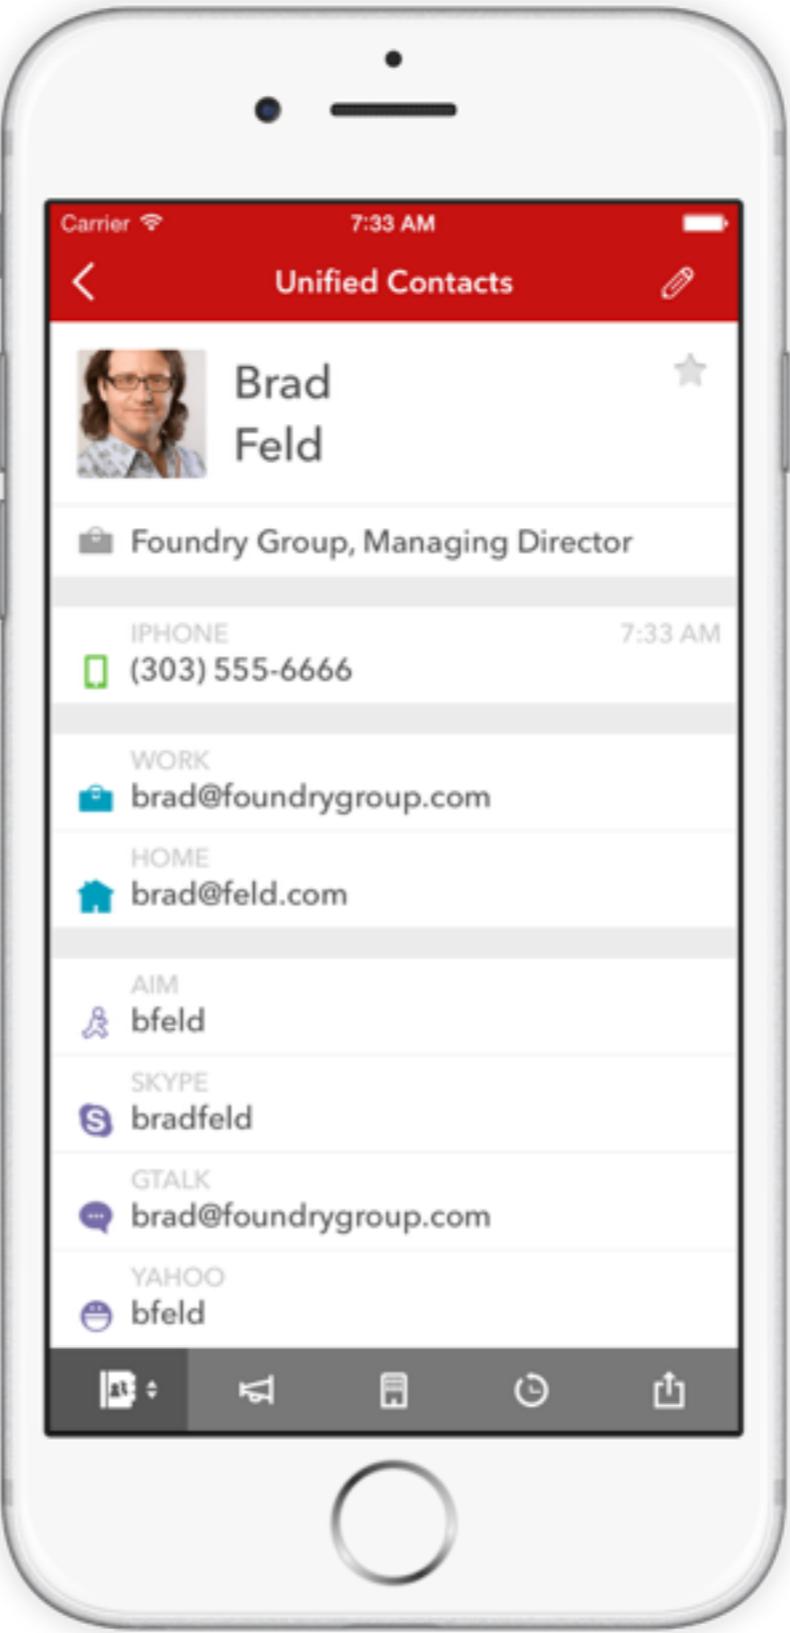
Contact

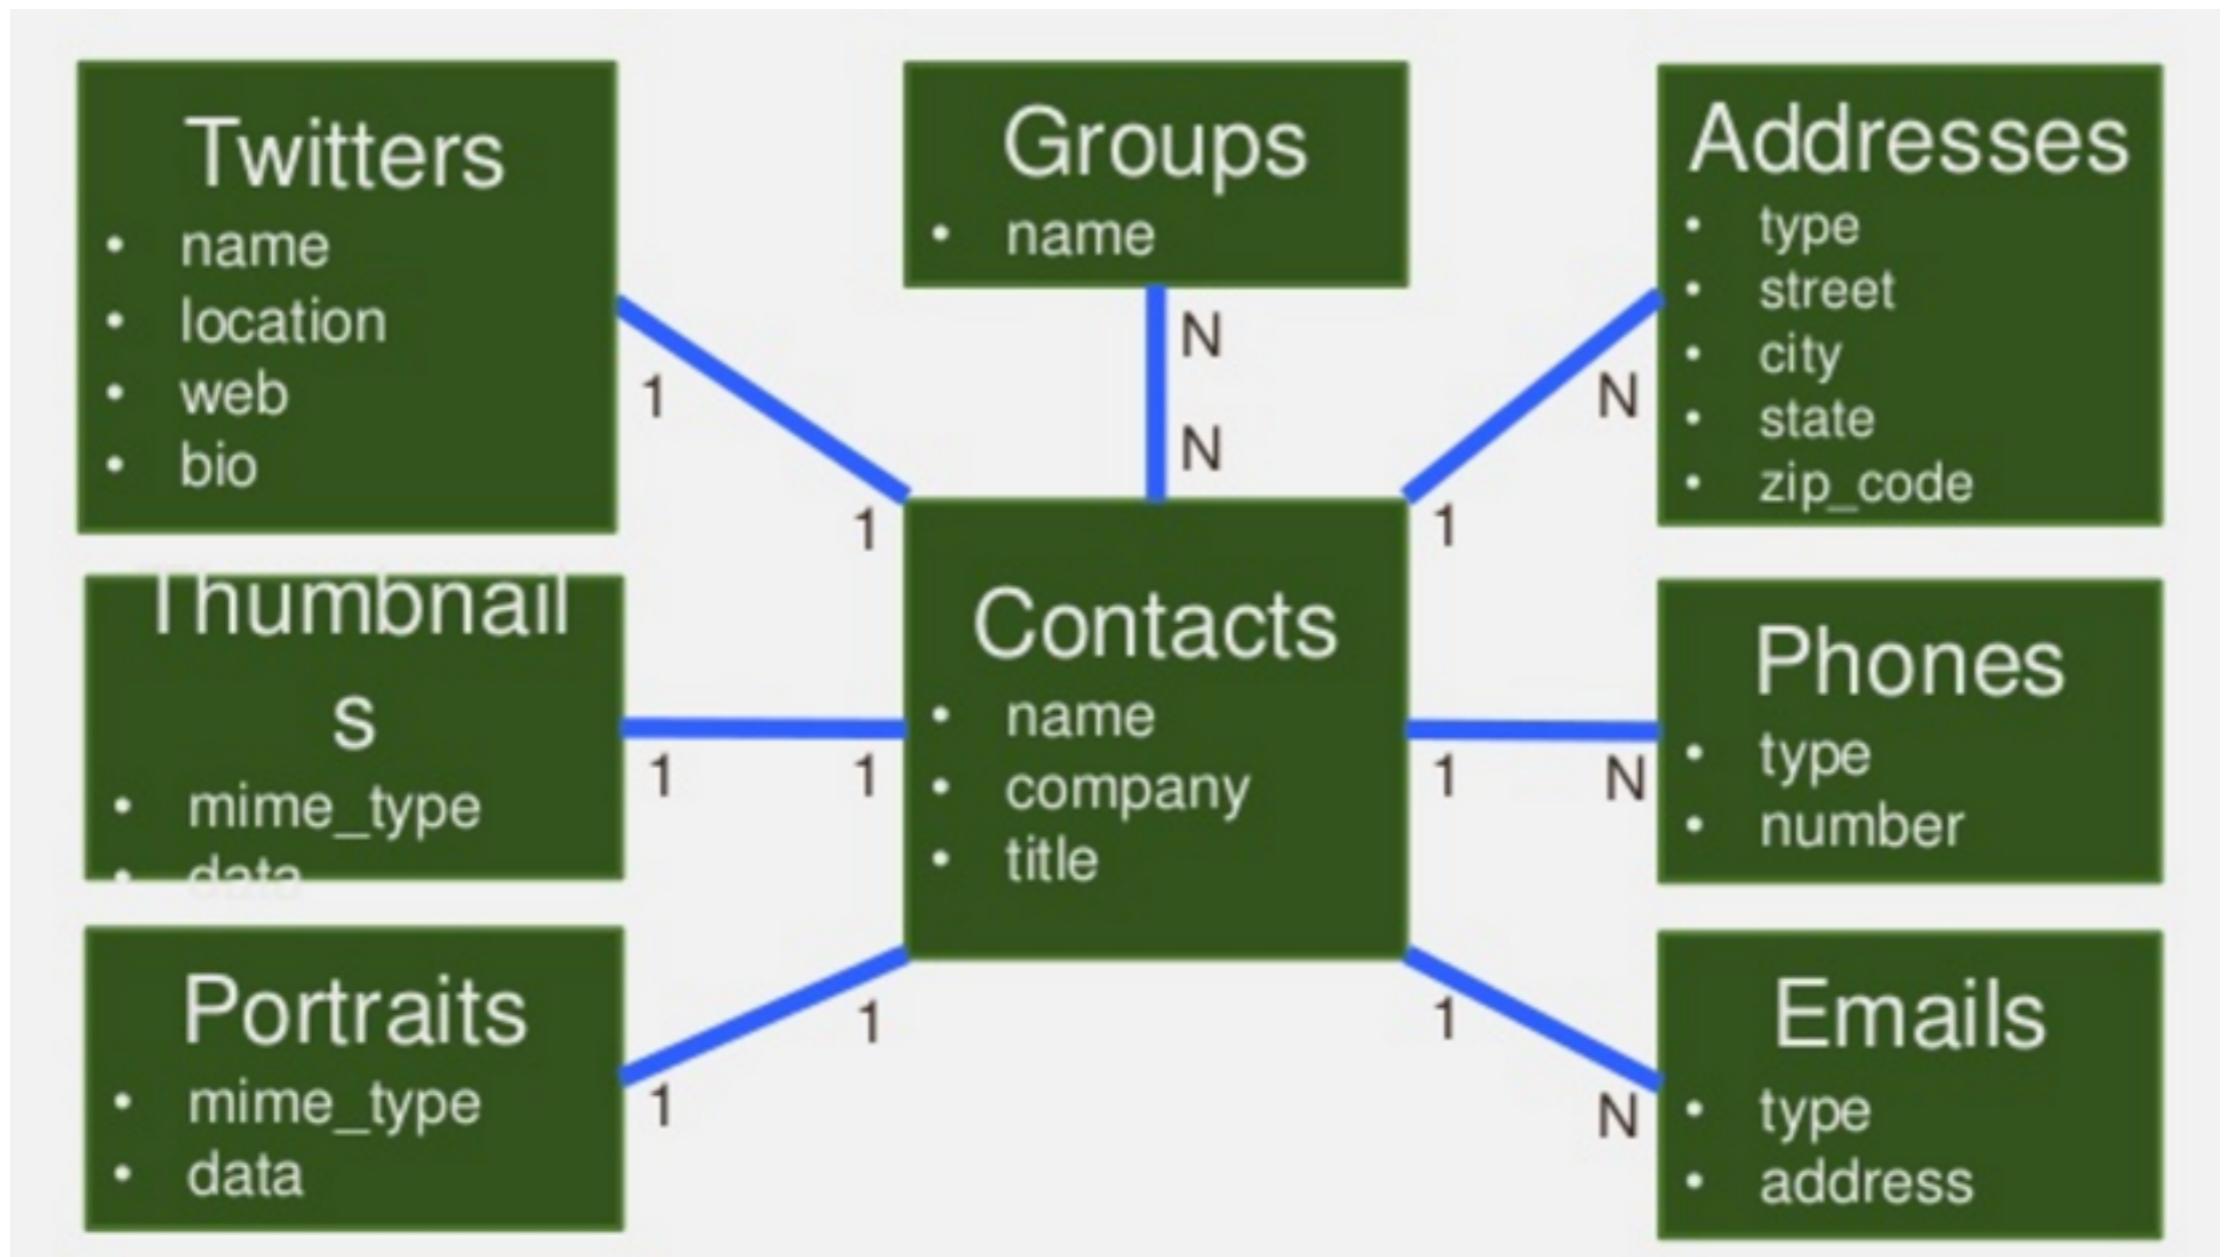
- name
- company
- address
 - street
 - city
 - State
 - zip_code
- title
- phone



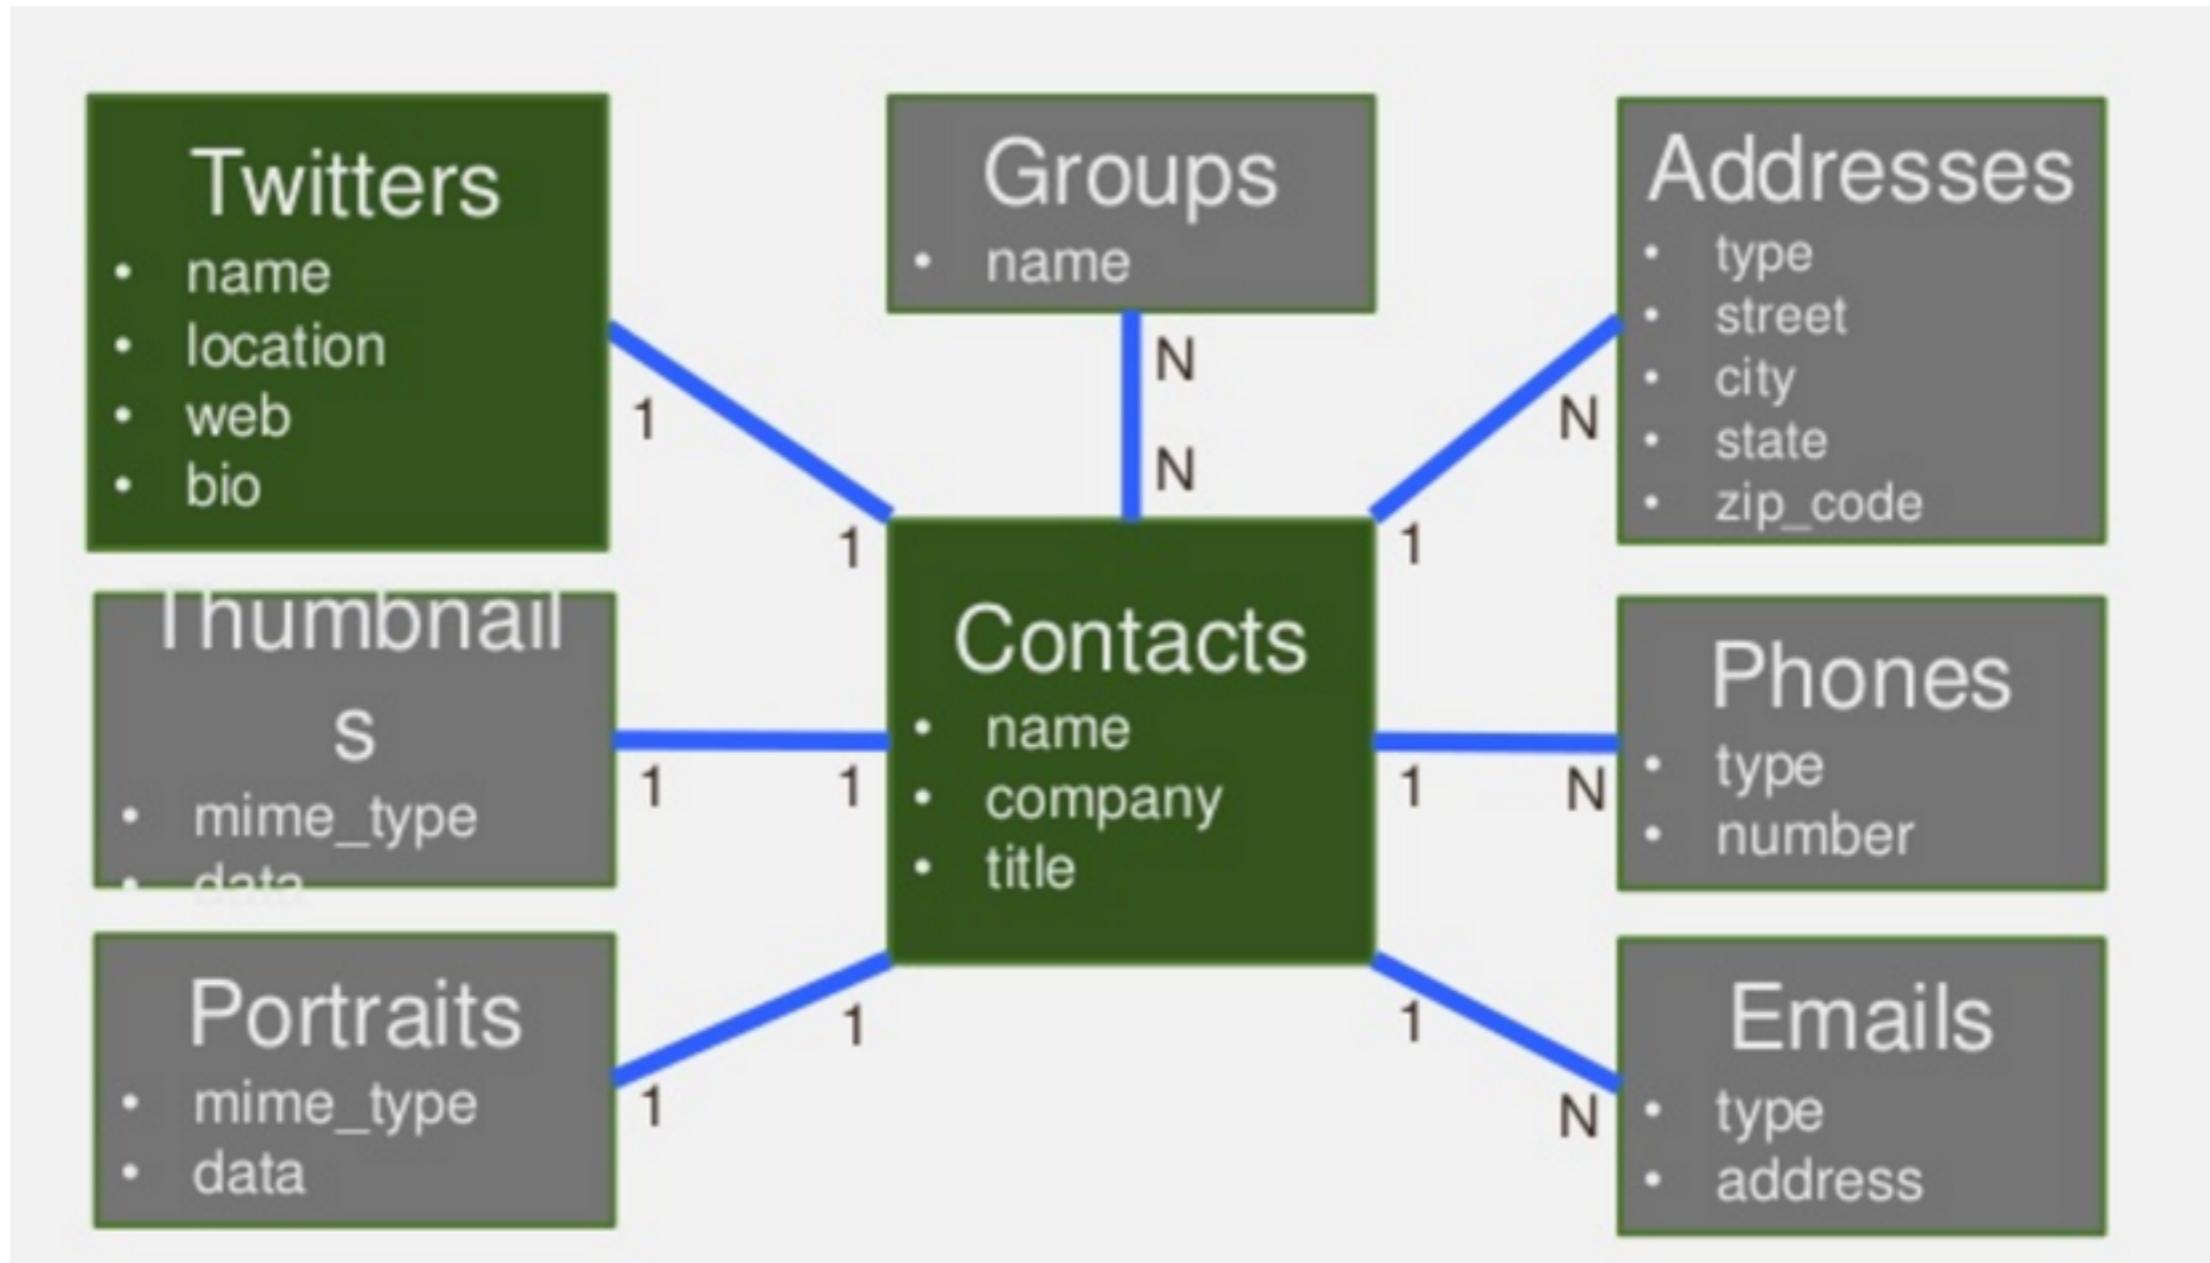
แตกต่างกันอย่างไร ? ทำไม ?







One-to-One



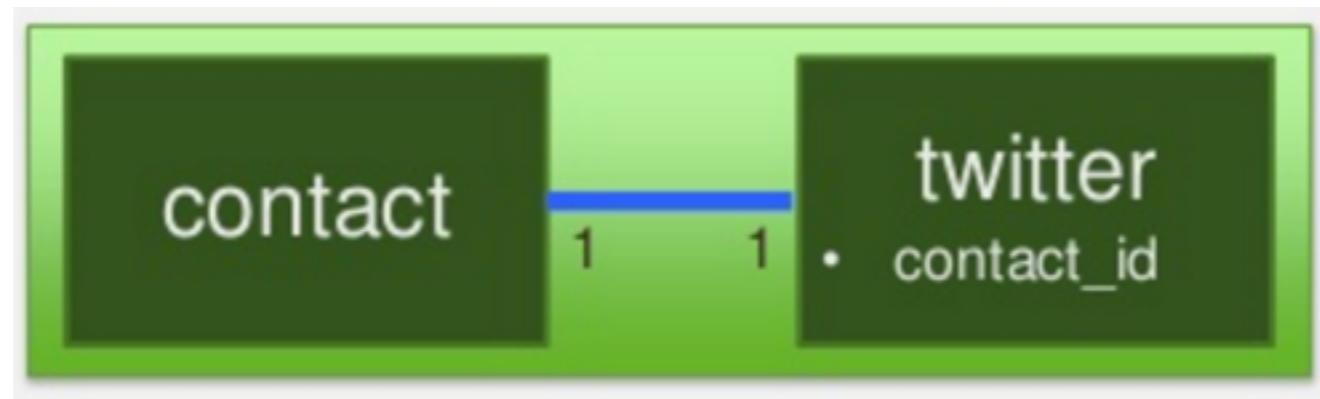
ออกแบบได้อย่างไรบ้าง ?



One-to-One :: 1



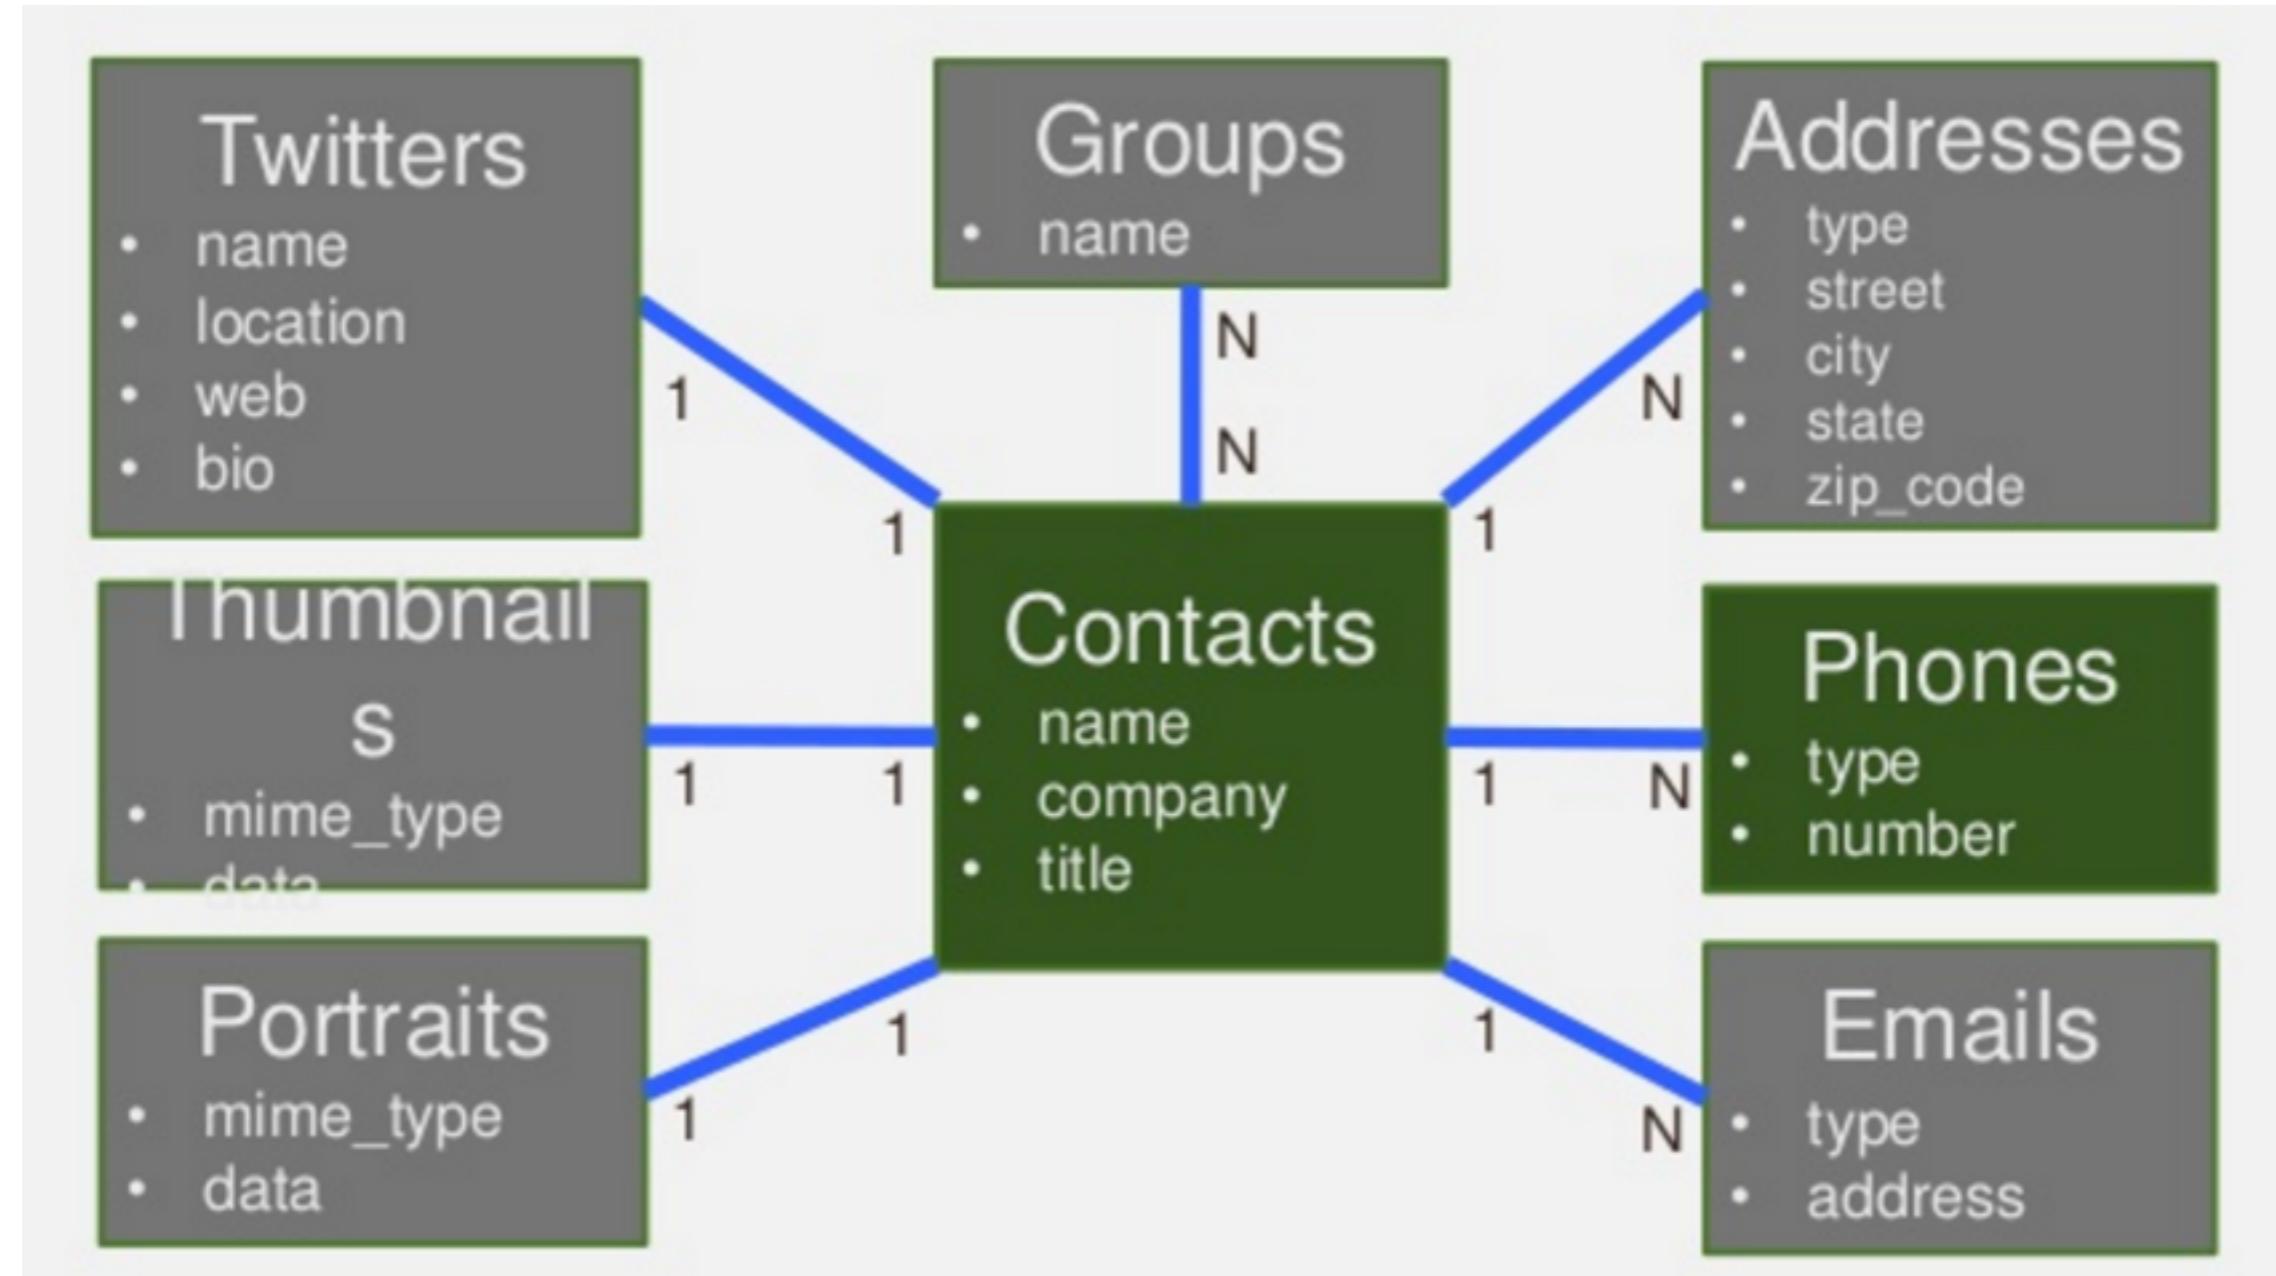
One-to-One :: 2



One-to-One :: 3



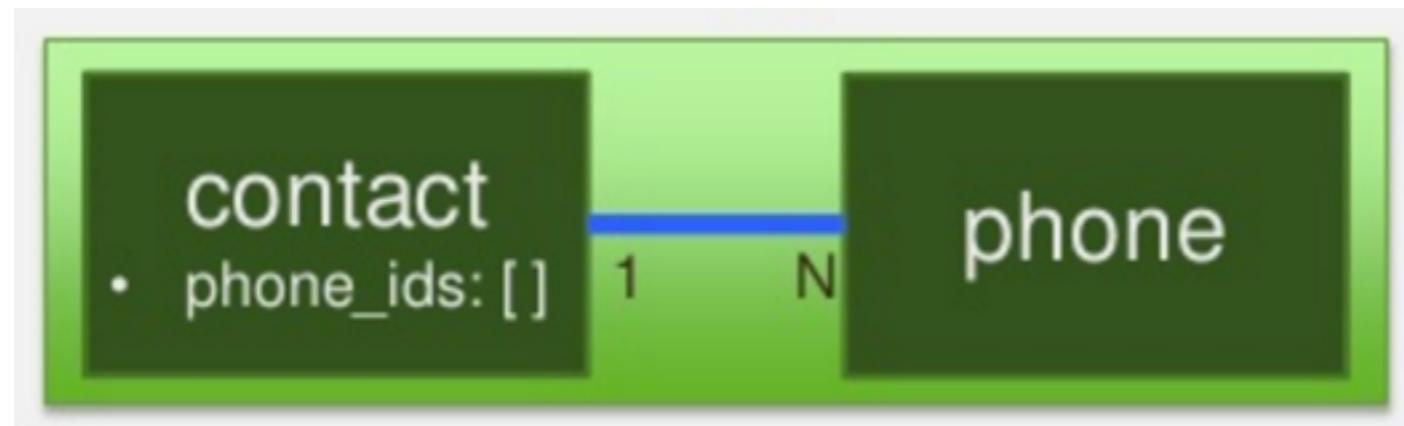
One-to-Many



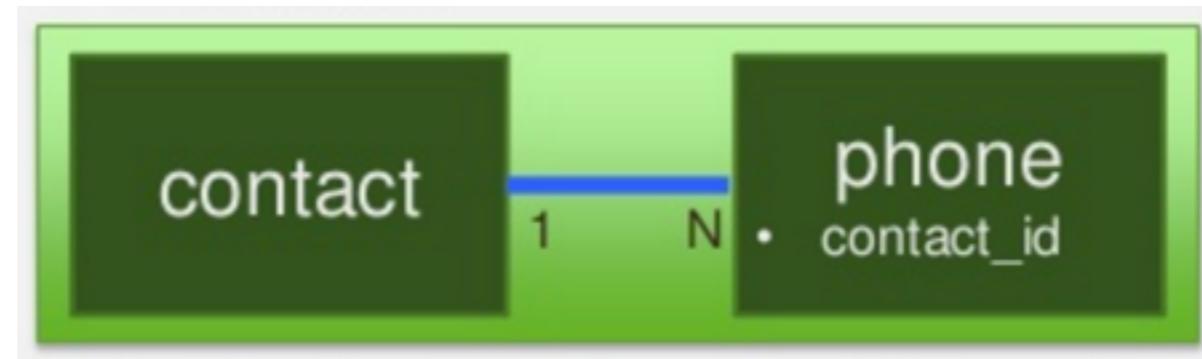
ออกแบบได้อย่างไรบ้าง ?



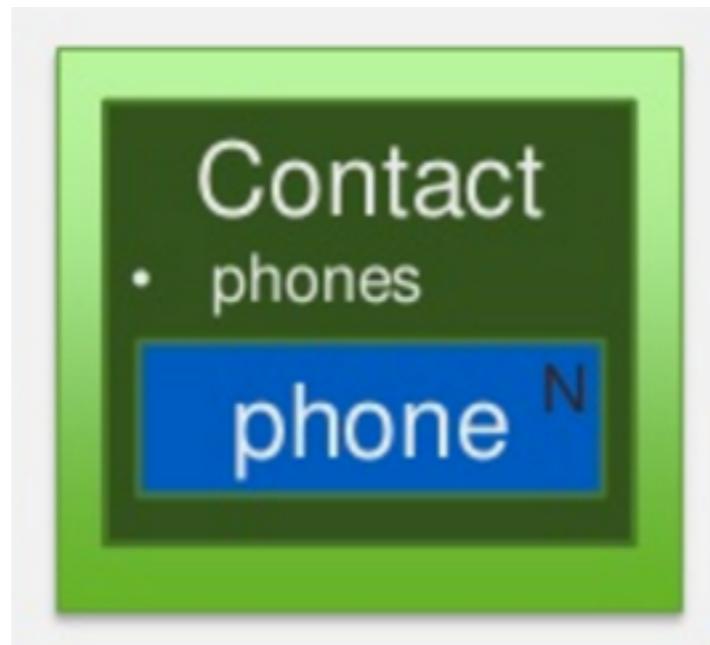
One-to-Many :: 1



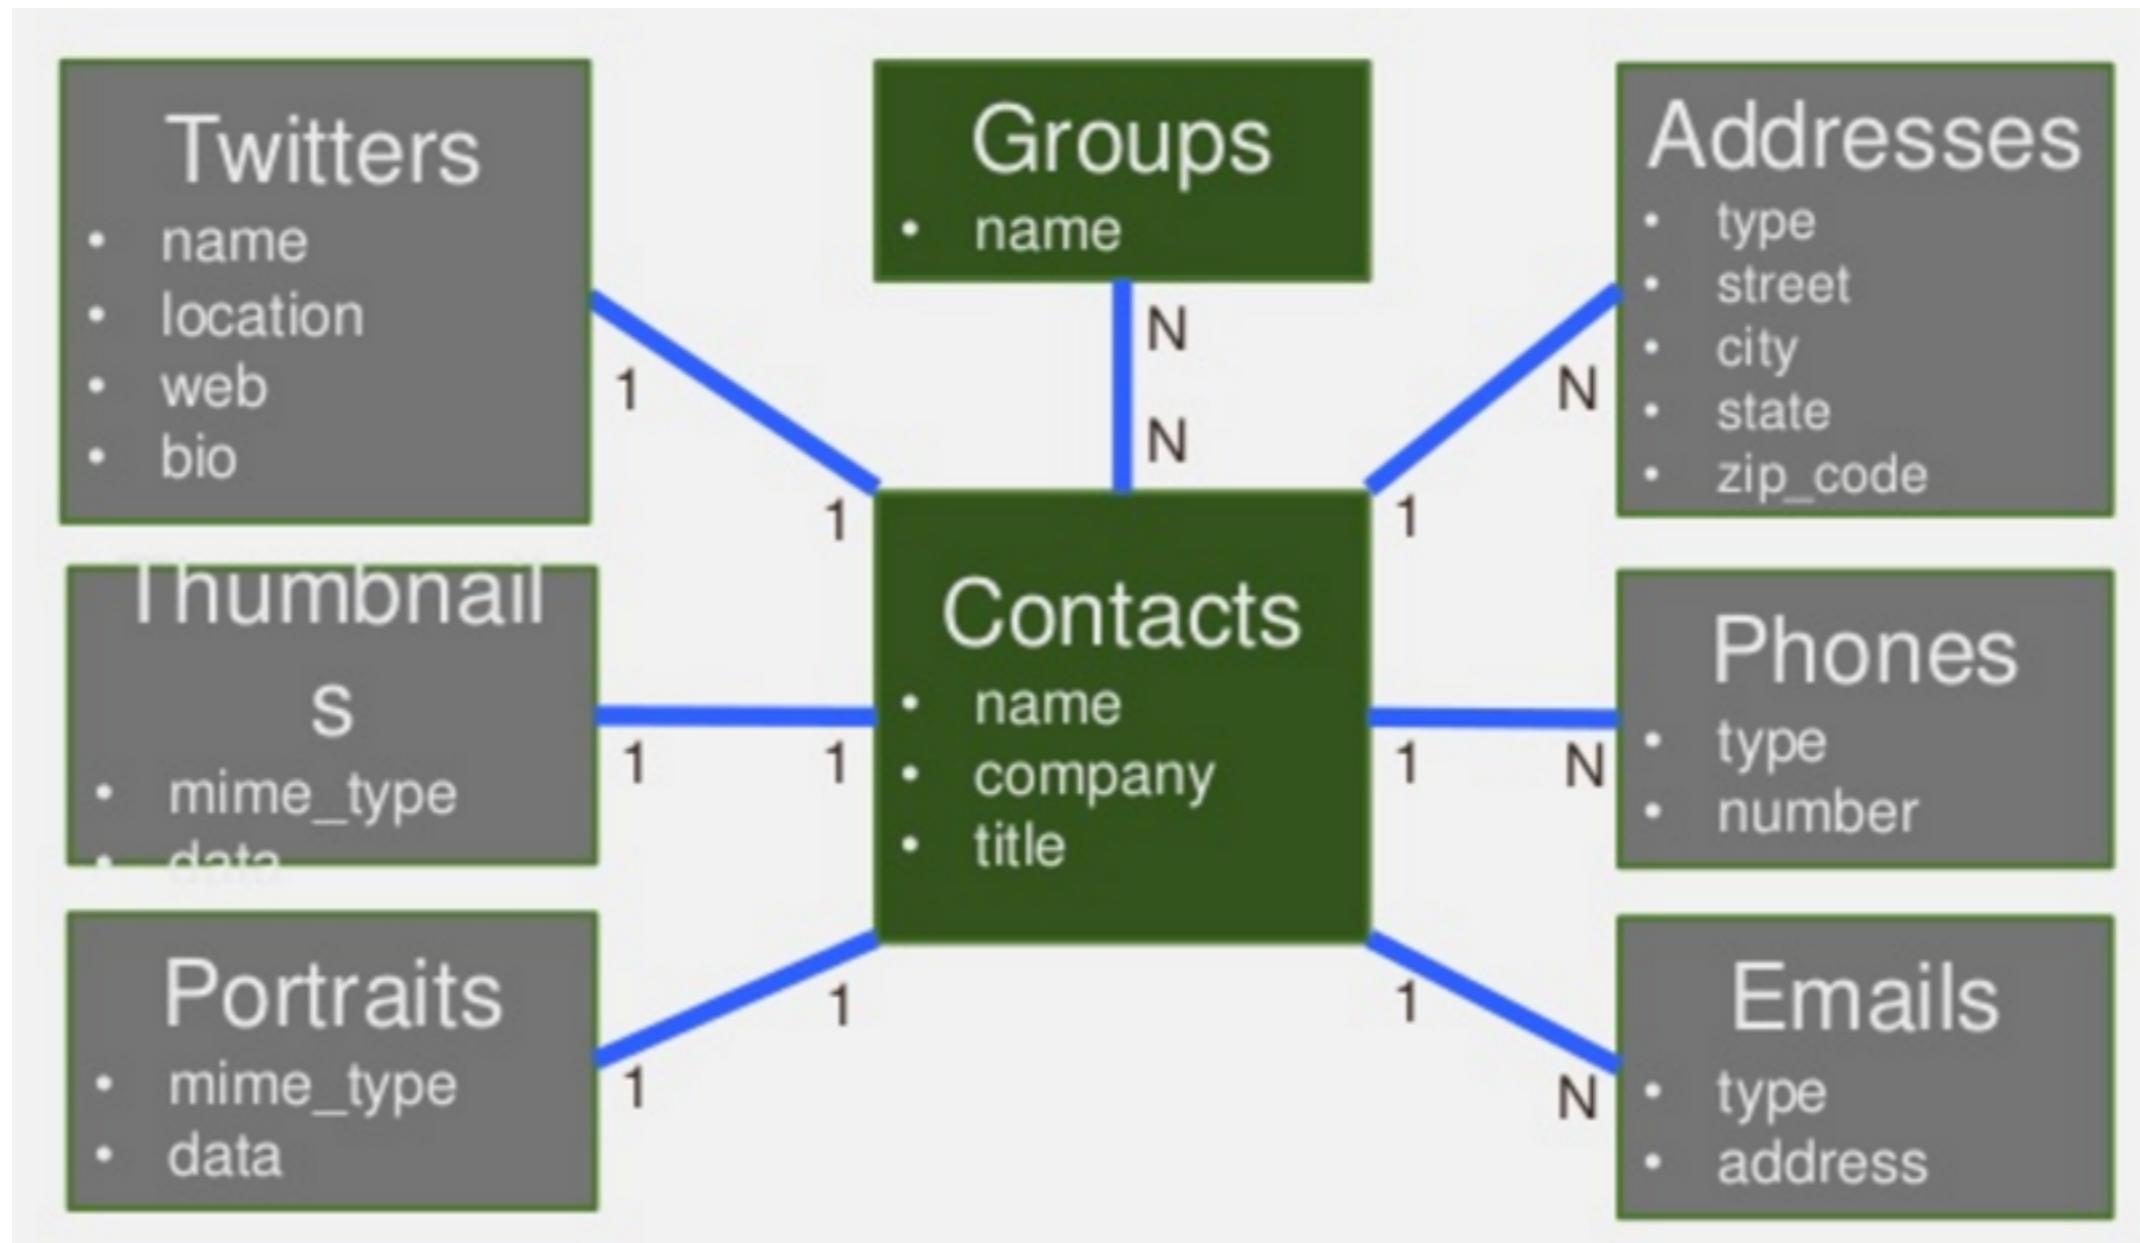
One-to-Many :: 2



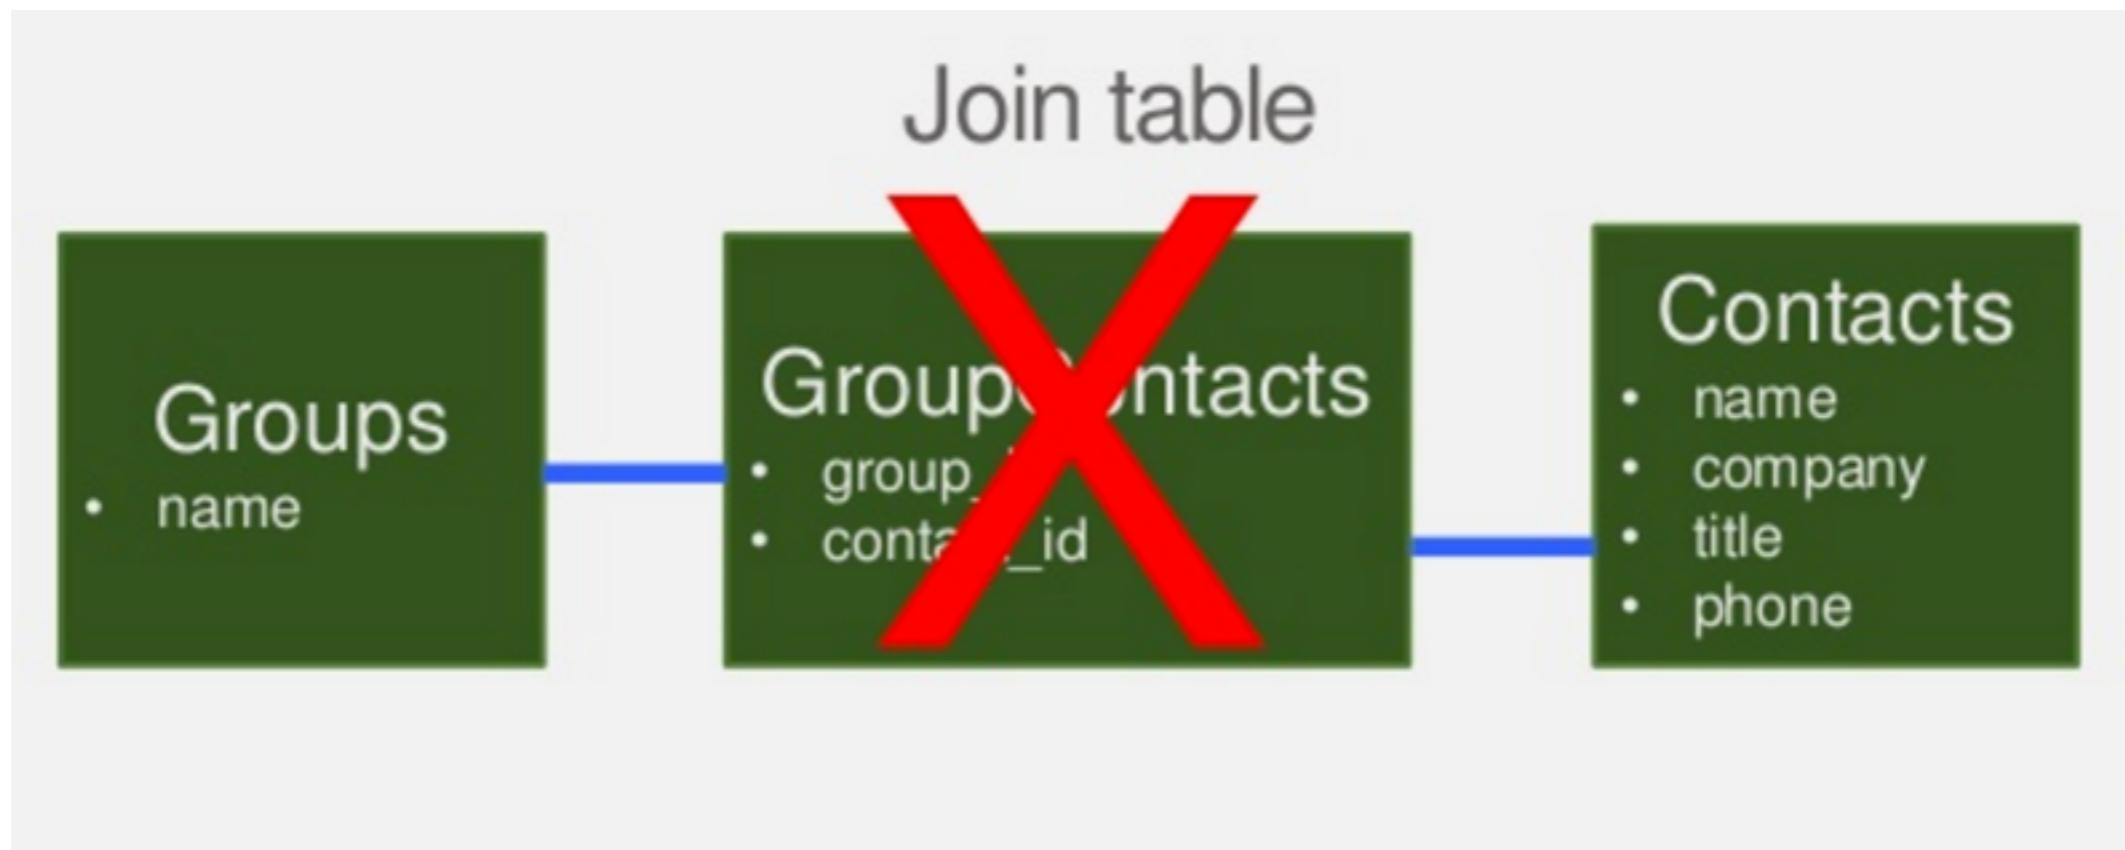
One-to-Many :: 3



Many-to-Many



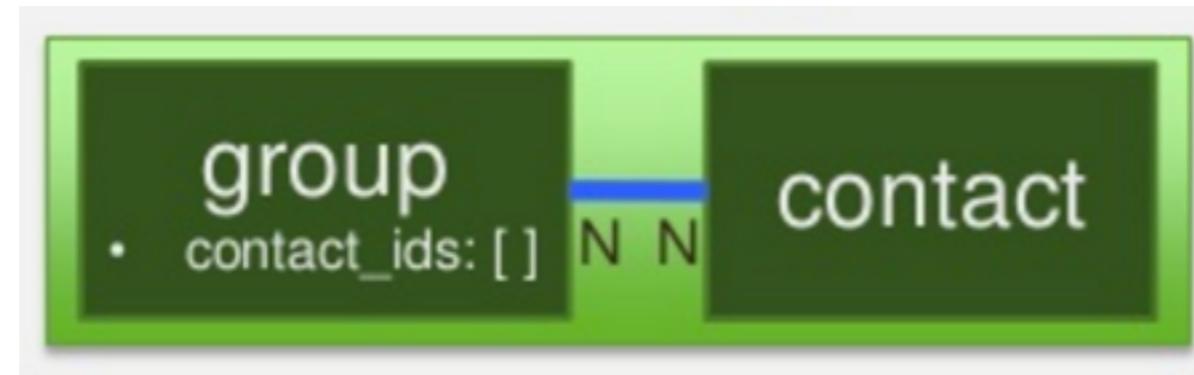
Many-to-Many



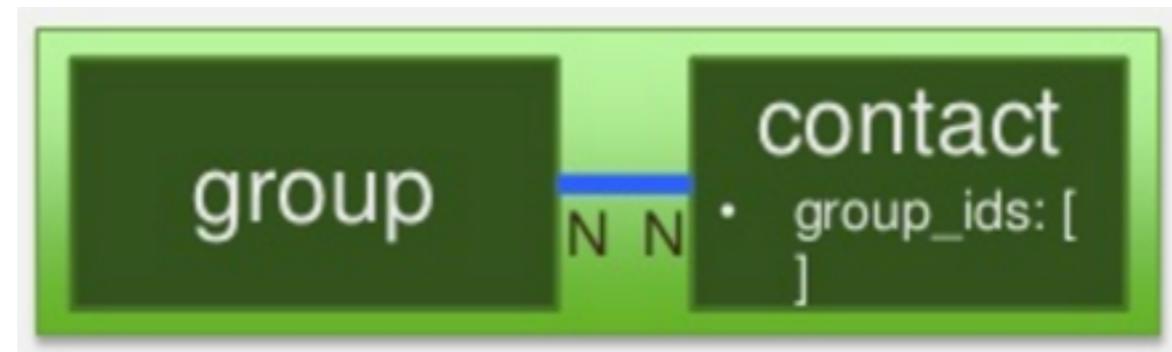
ออกแบบได้อย่างไรบ้าง ?



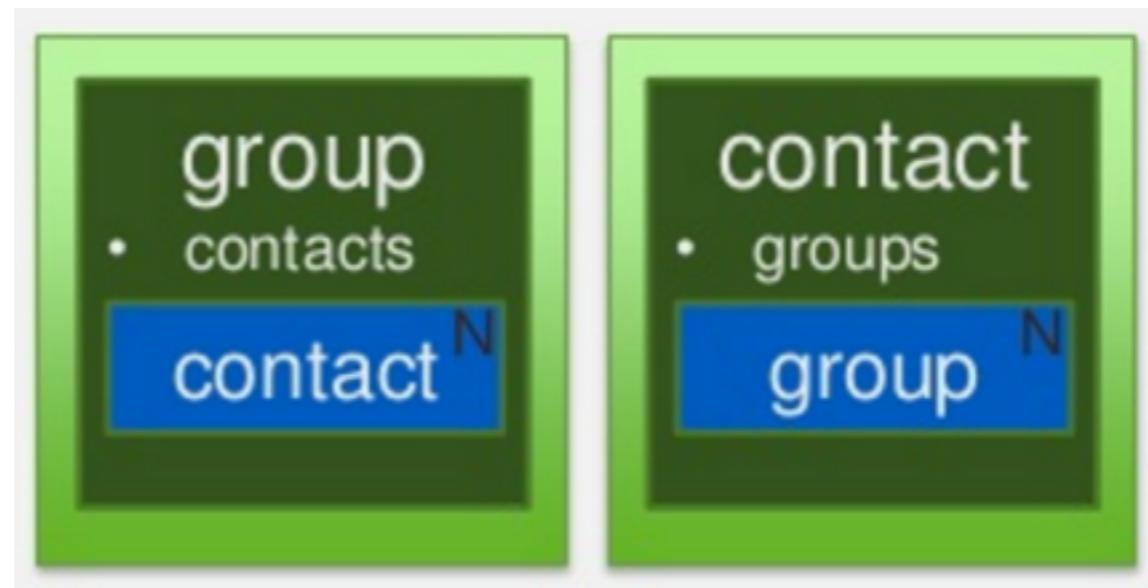
Many-to-Many :: 1



Many-to-Many :: 2

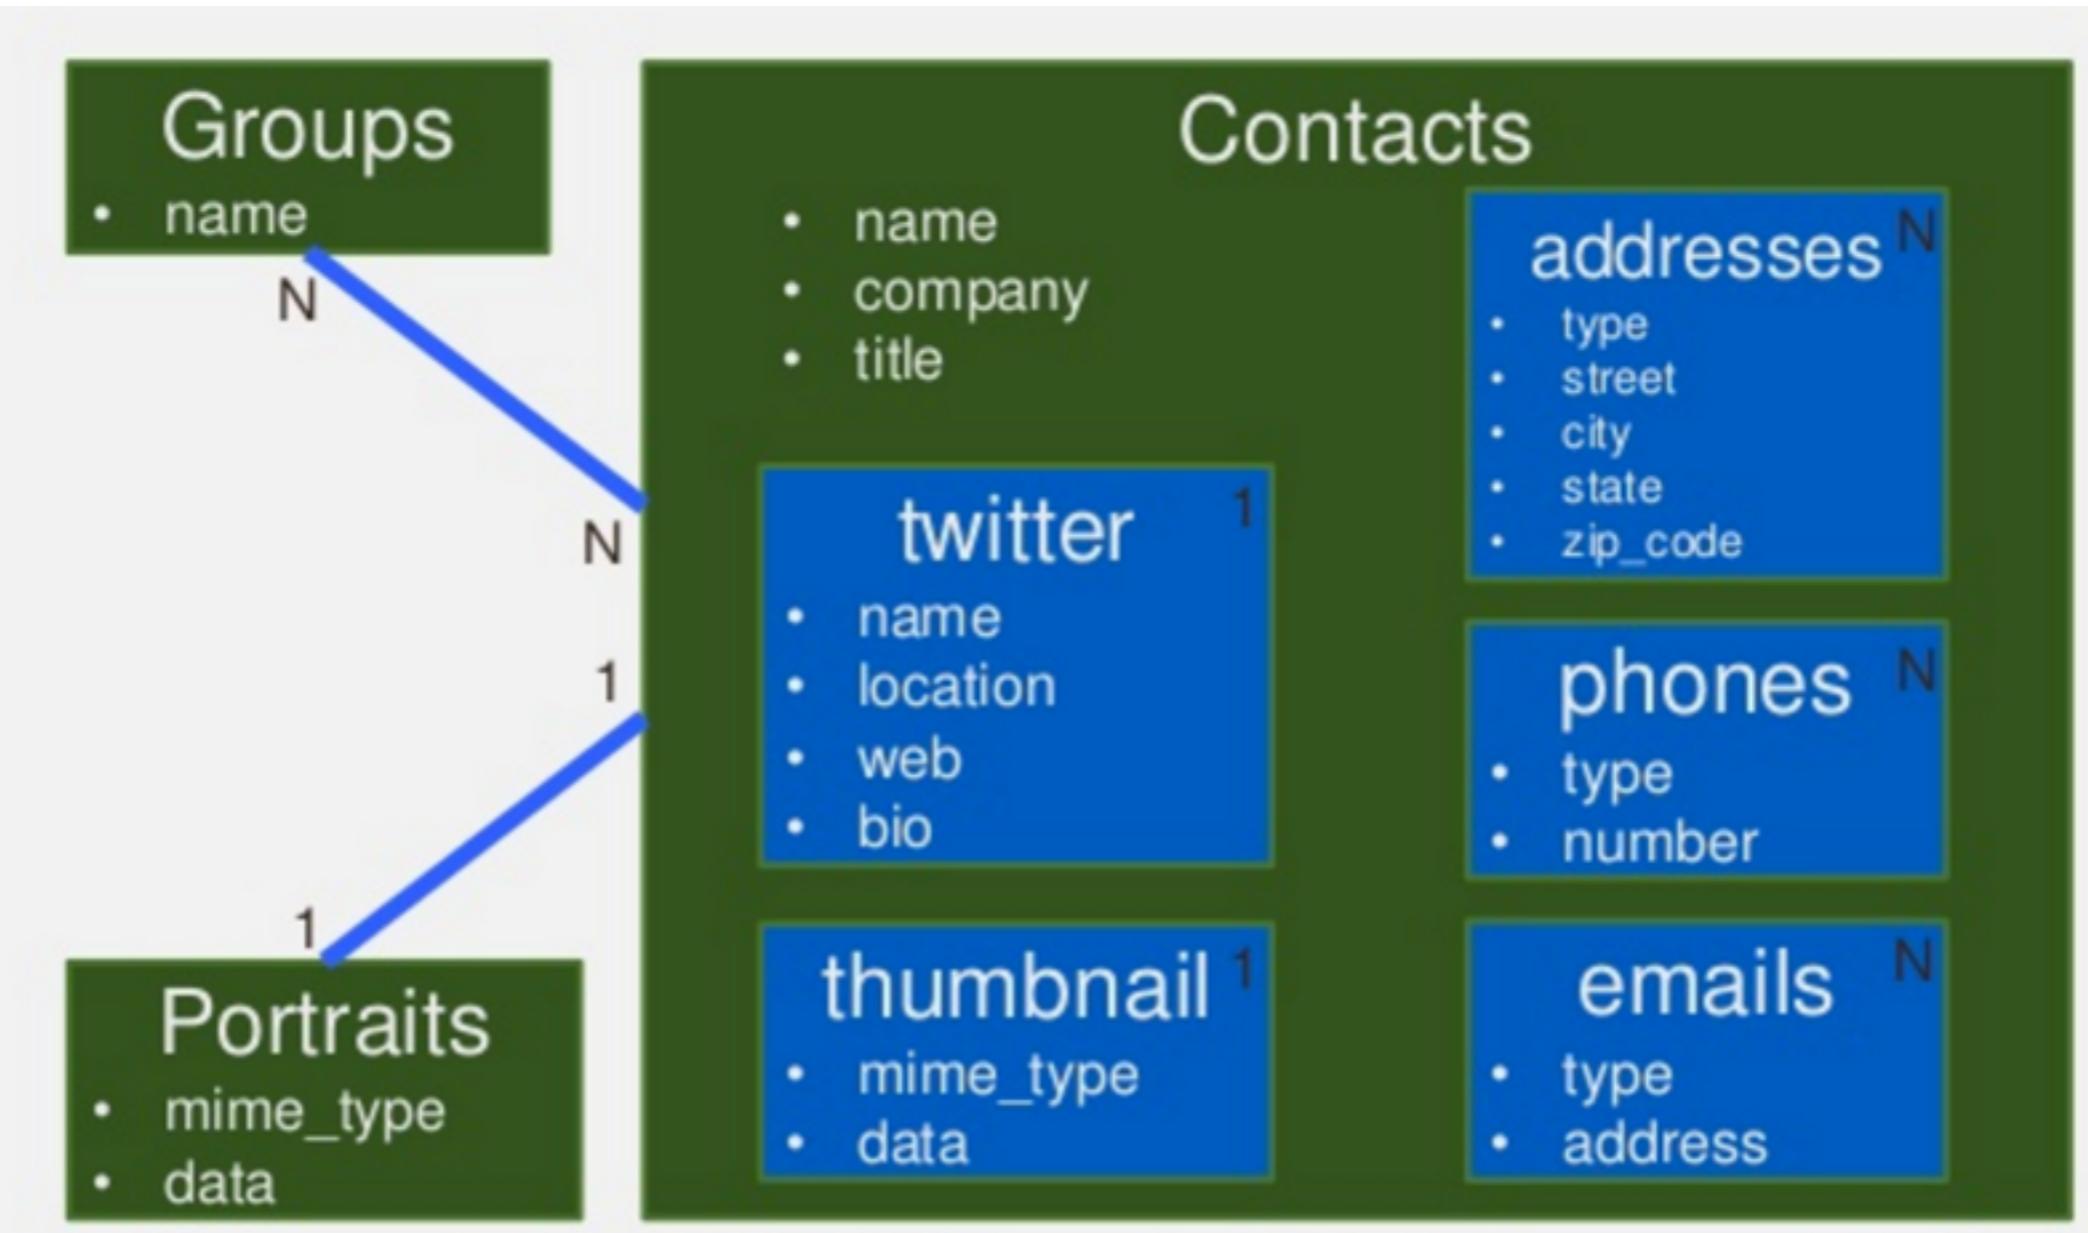


Many-to-Many :: 3



ສຸດທ້າຍແລ້ວ





Workshop

Working with contact information

05_contact.txt



ระบบตัวอย่าง





Learn, Share & Fun



หน้าแรก

เลือกห้อง ▼

แท็ก

กิจกรรม

N

อื่นๆ ▼

Smart Search



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

ออกแบบด้วยแนวคิดของ RDBMS ?





เทคนิคในการออกแบบ NoSQL



1. Denormalization

Query data volume หรือ I/O per query

vs.

Total data volume



1. Denormalization

Processing complexity

vs.

Total data volume



2. Aggregation

Minimization of one-to-many relationship
Reduction of joins



2. Aggregation

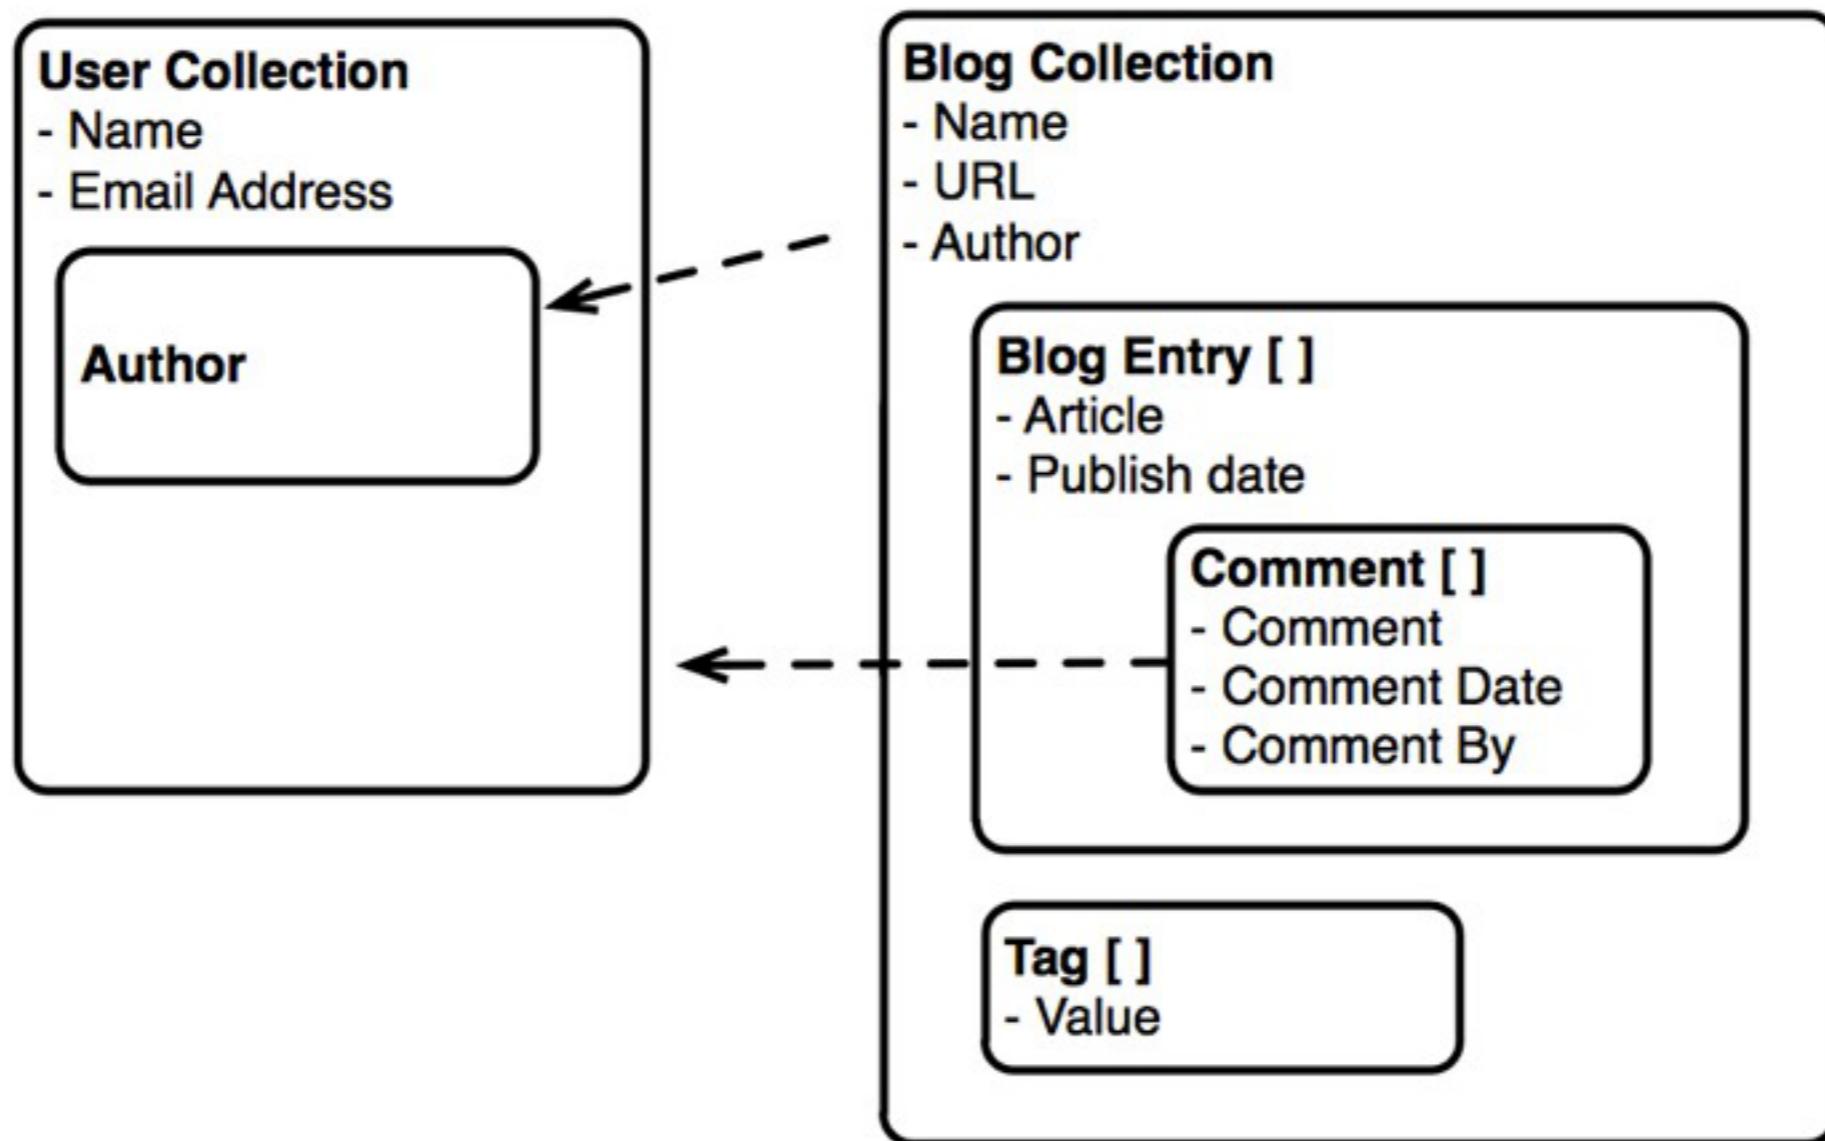
Minimization of one-to-many relationship

Reduction of joins

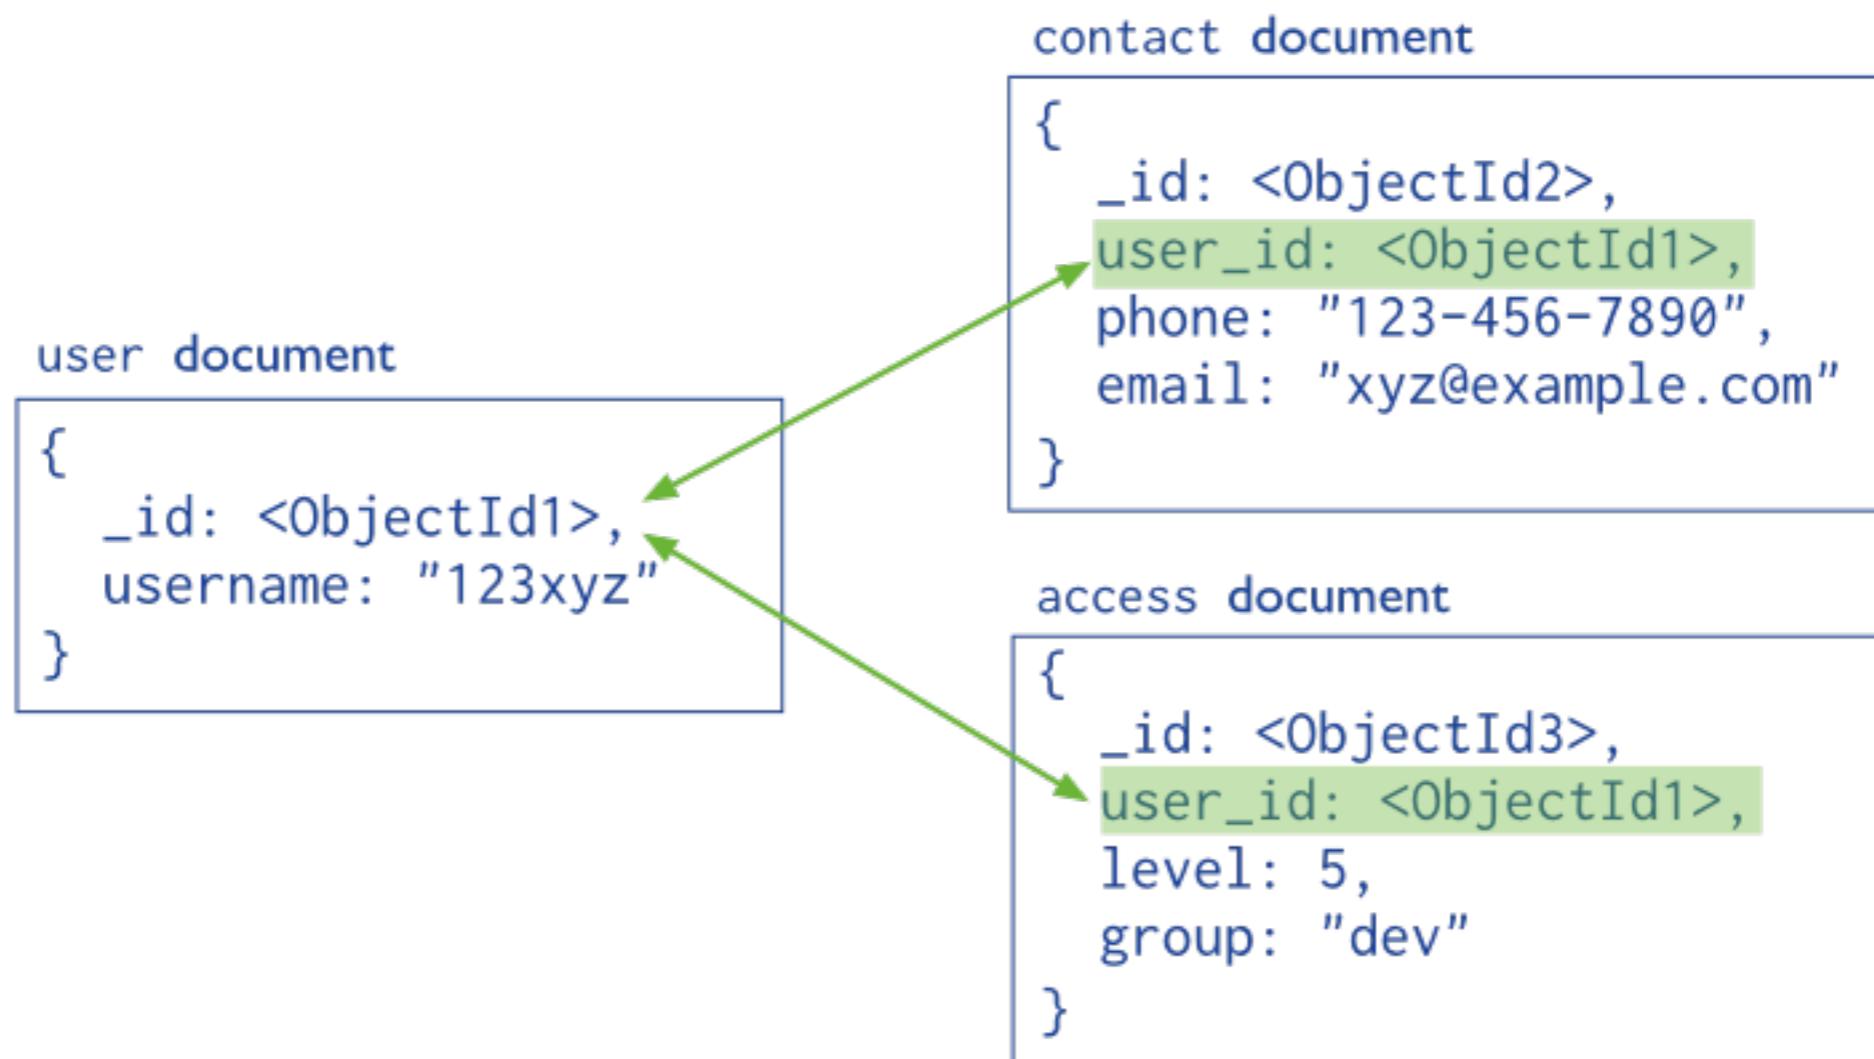
Nested or Embedded entity



Embedded document



Normalized data model



Embedded document

```
{  
  _id: <ObjectId1>,  
  username: "123xyz",  
  contact: {  
    phone: "123-456-7890",  
    email: "xyz@example.com"  
  },  
  access: {  
    level: 5,  
    group: "dev"  
  }  
}
```



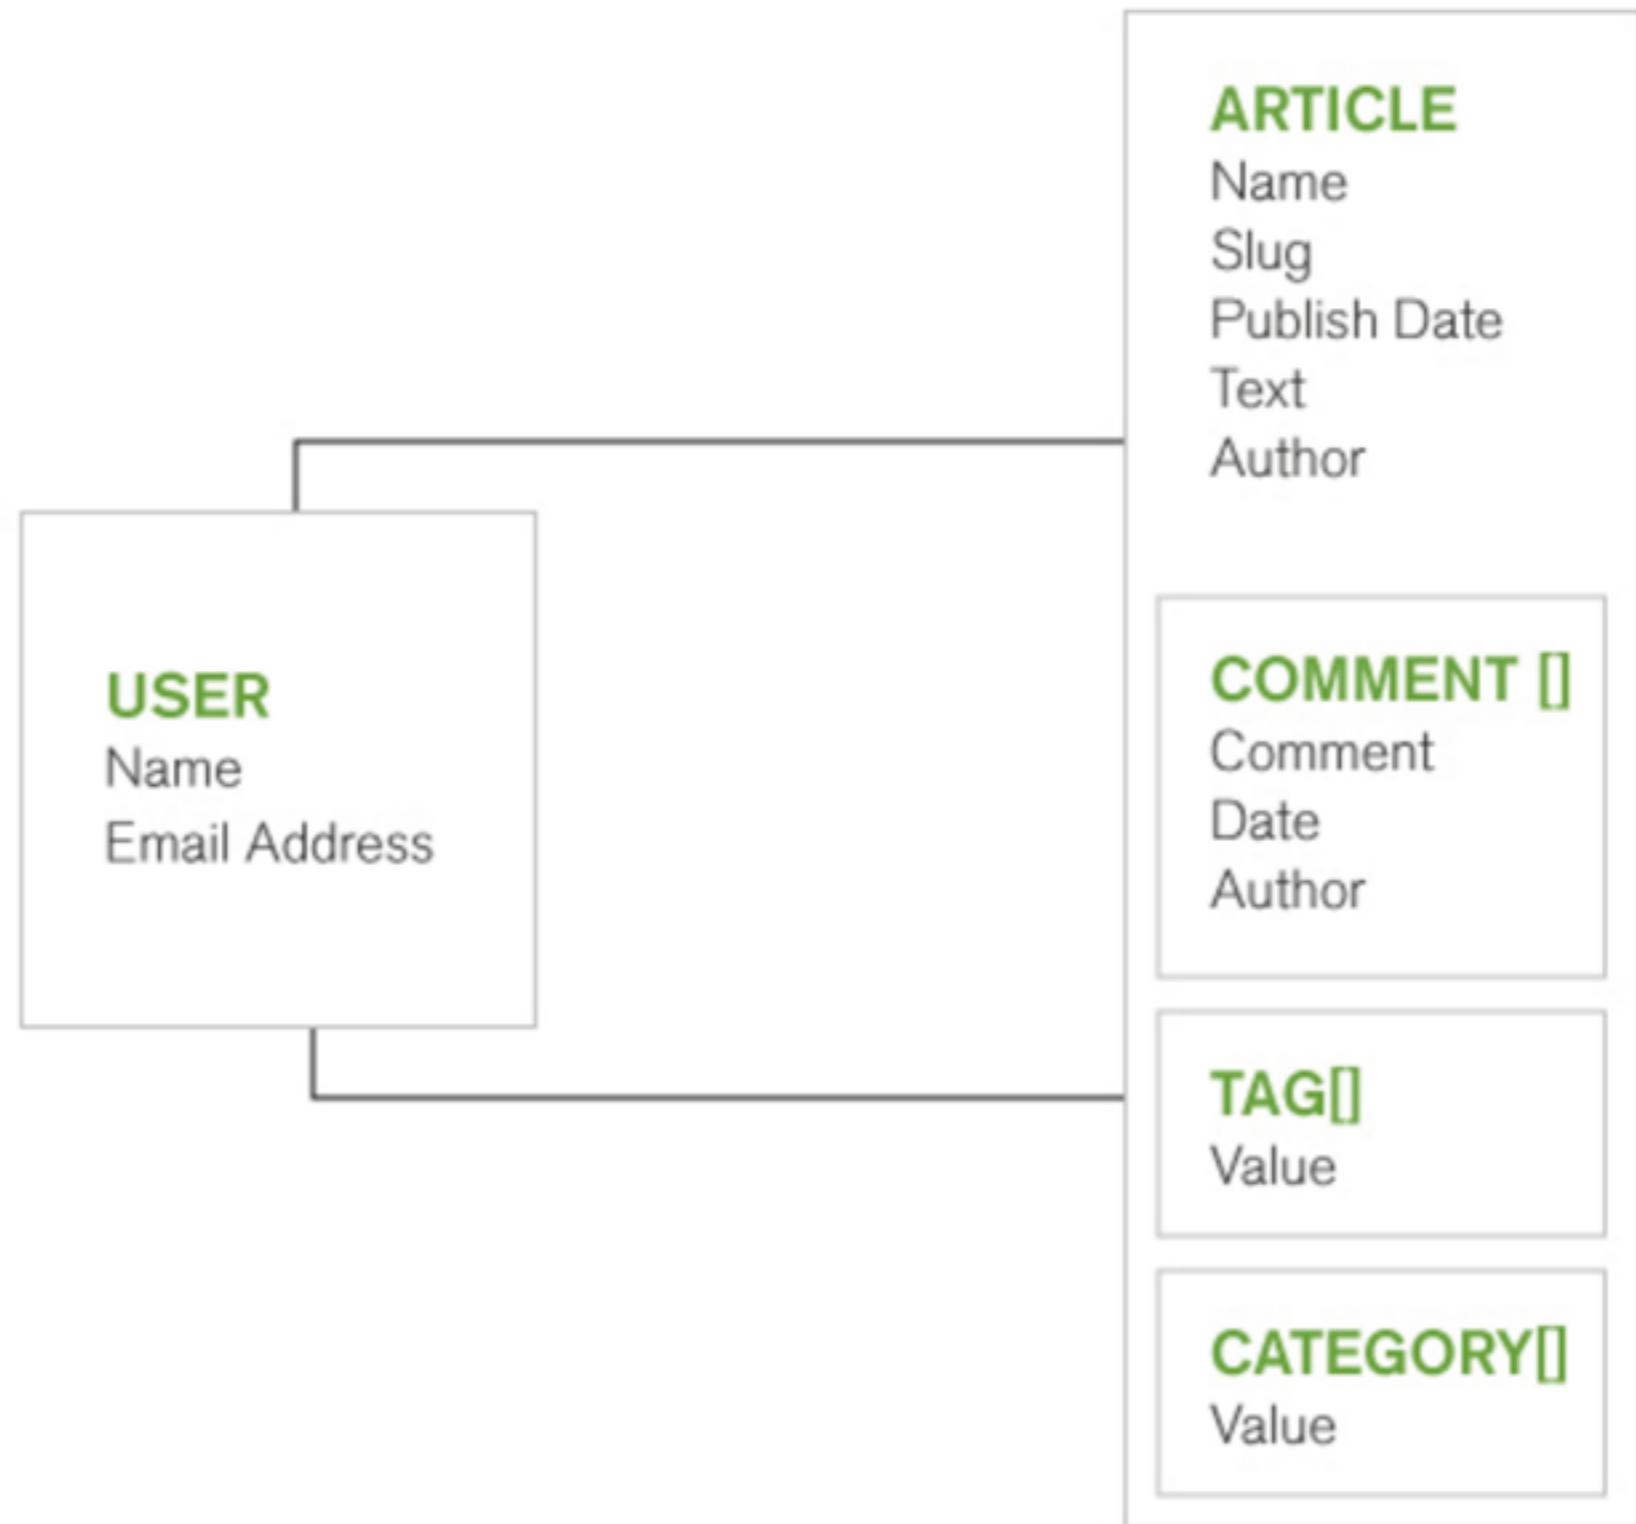
Embedded sub-document

Embedded sub-document



ออกแบบด้วยแนวคิดของ NoSQL





ปัญหาที่อาจจะเกิดขึ้น ?



บริษัท สยามชนาญกิจ จำกัด และเพื่อนพ้องน้องพี่



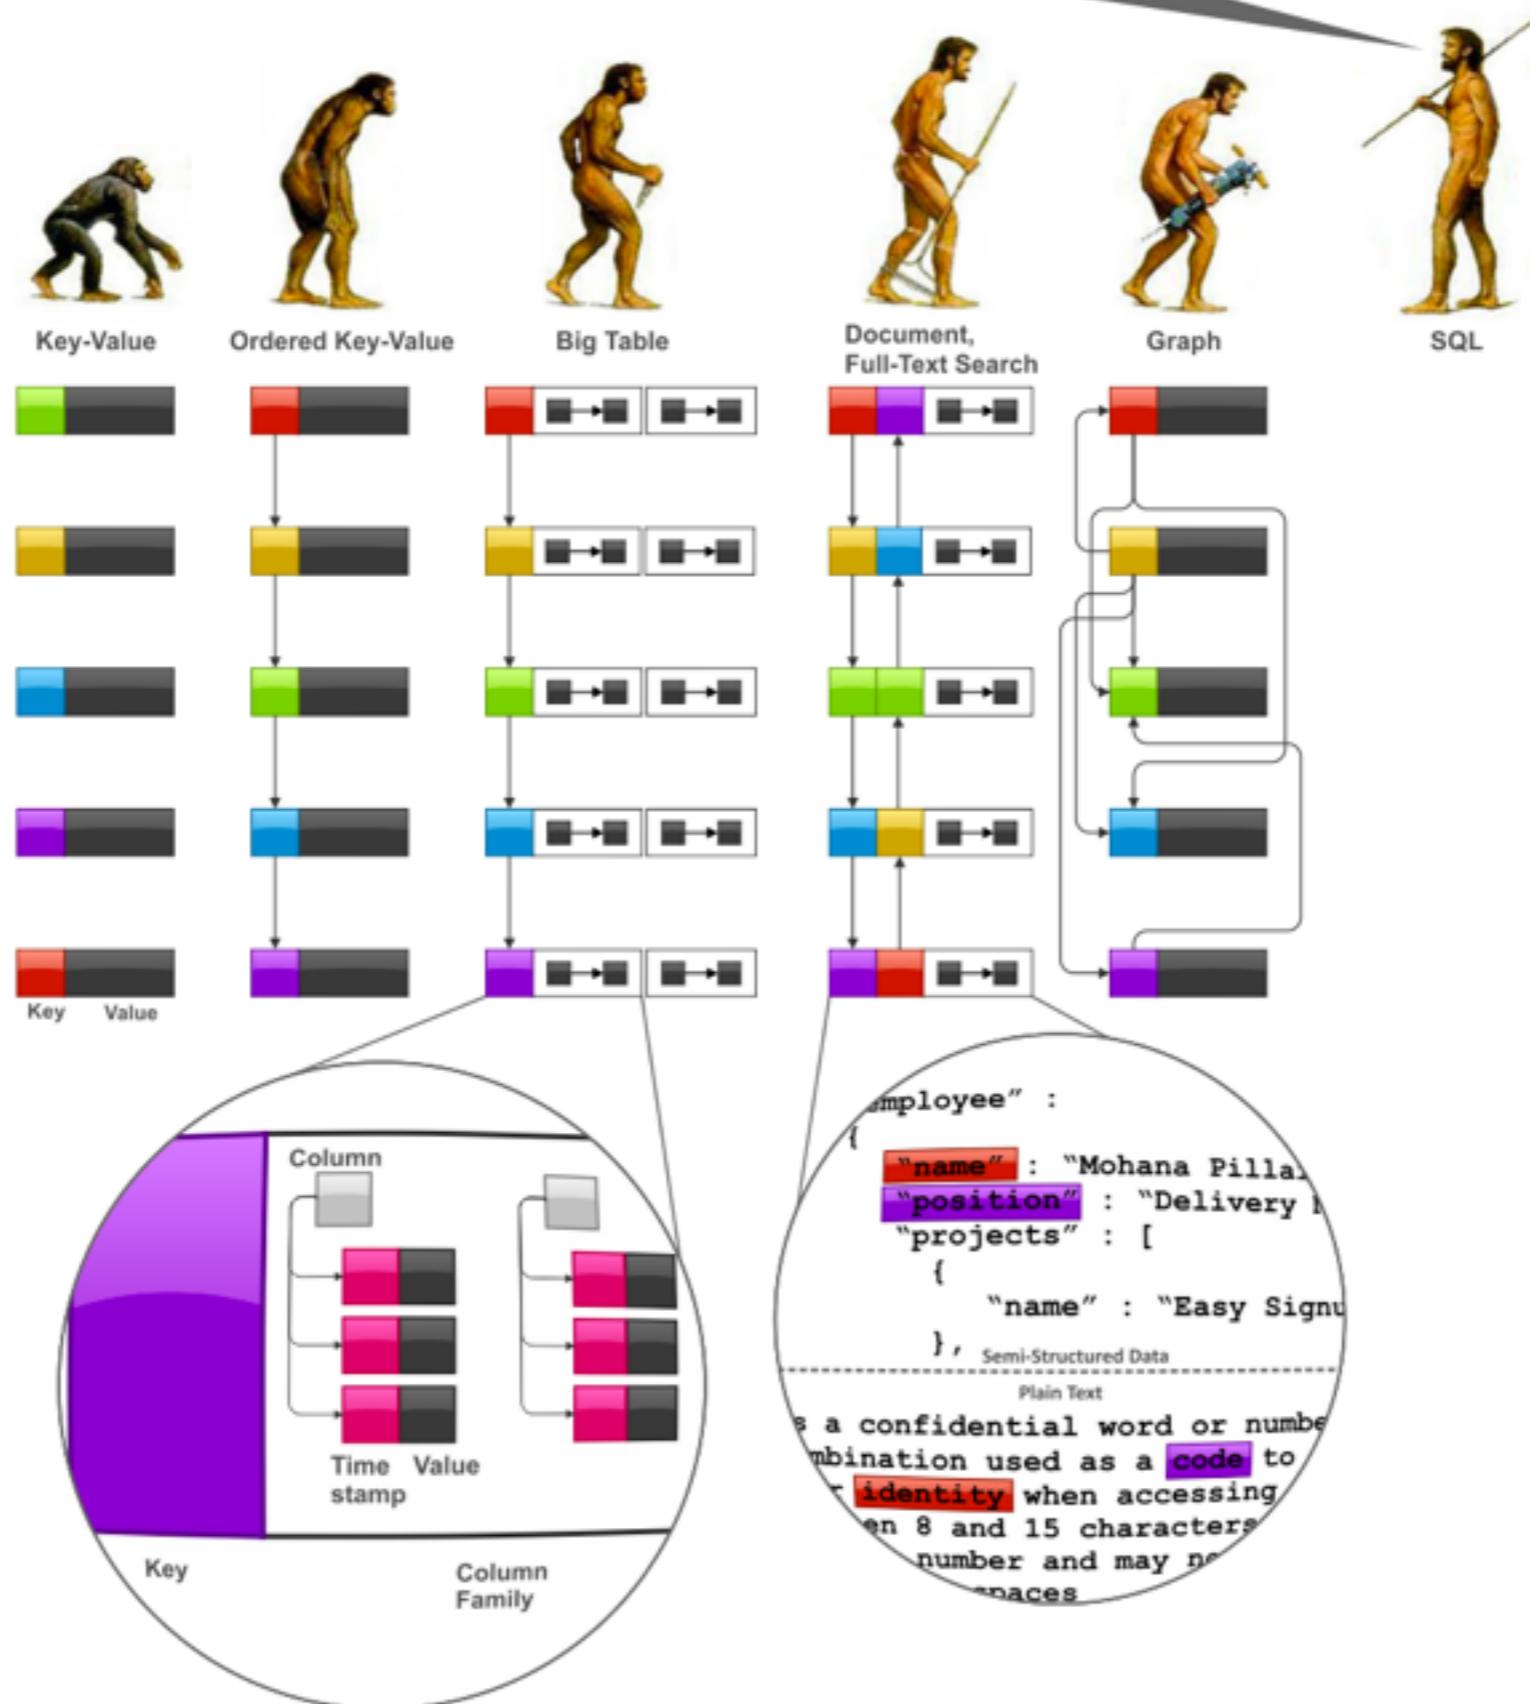
แก้อย่างไรดี ?



Embedding vs Referencing



Stop following me, you fucking freaks!



คำแนะนำสำหรับการอุ่นเครื่อง



Iterate schema design development



Measure performance



Find bottlenecks



Embedding by default (90/10)



Referencing when you need more scaling



Normalize vs Denormalize



Flexible for development and change





**KEEP
CALM**

AND

**use the right tool
for the right job**



Replication



What is Replication ?

Synchronizing data across multiple servers

Redundancy

Increase data availability

No Single Point of Failure (SPoF)



why Replication ?

To keep data safe

High availability of data (24*7)

Disaster recovery

No downtime

Read scaling



Working with MongoDB

Group of mongod instances

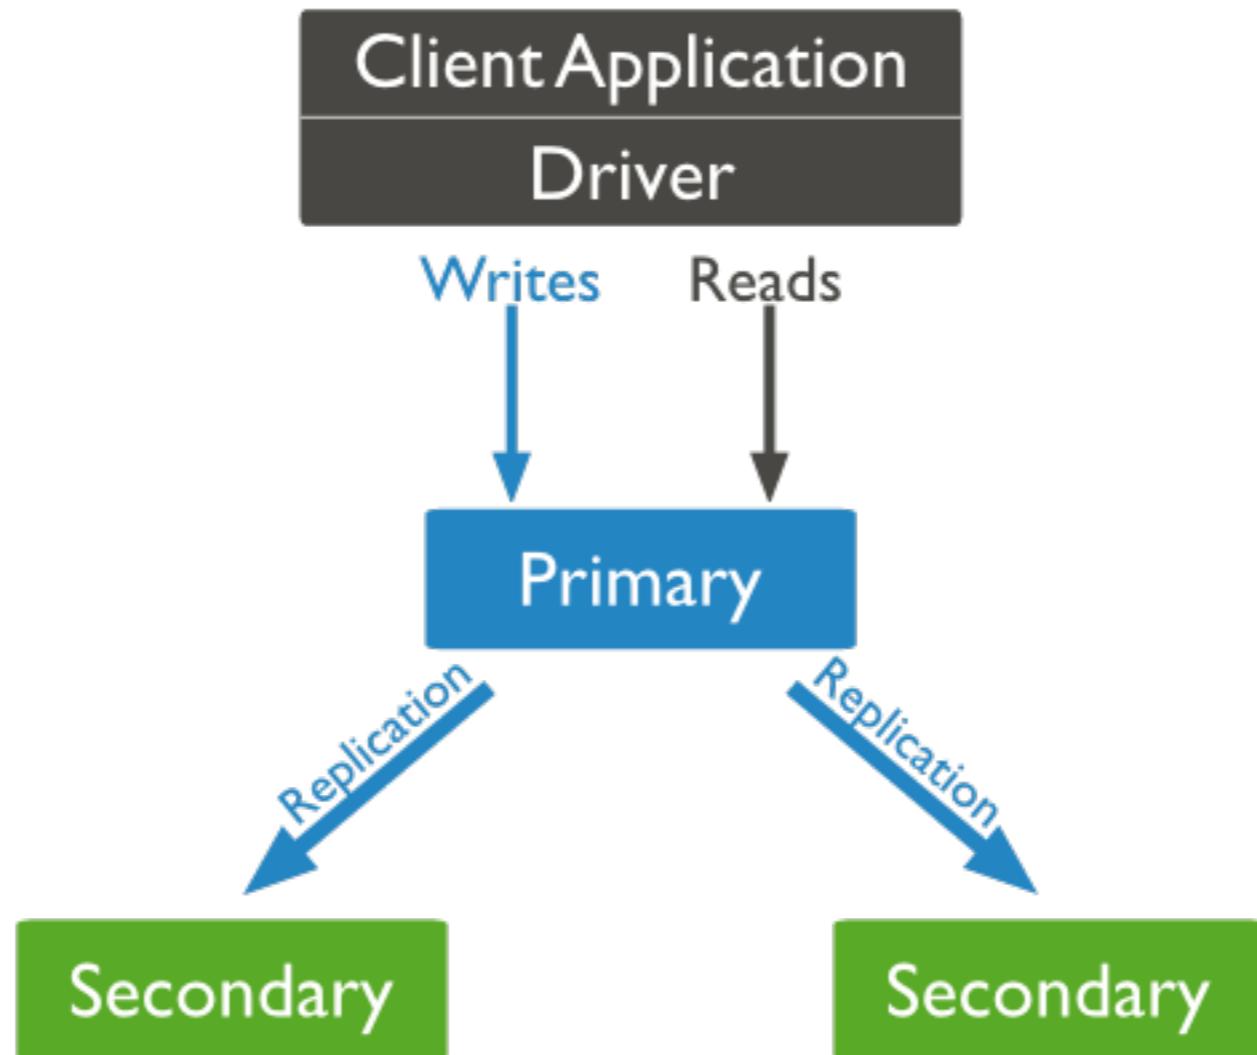
Minimum 3 nodes are required !!

One node is Primary node

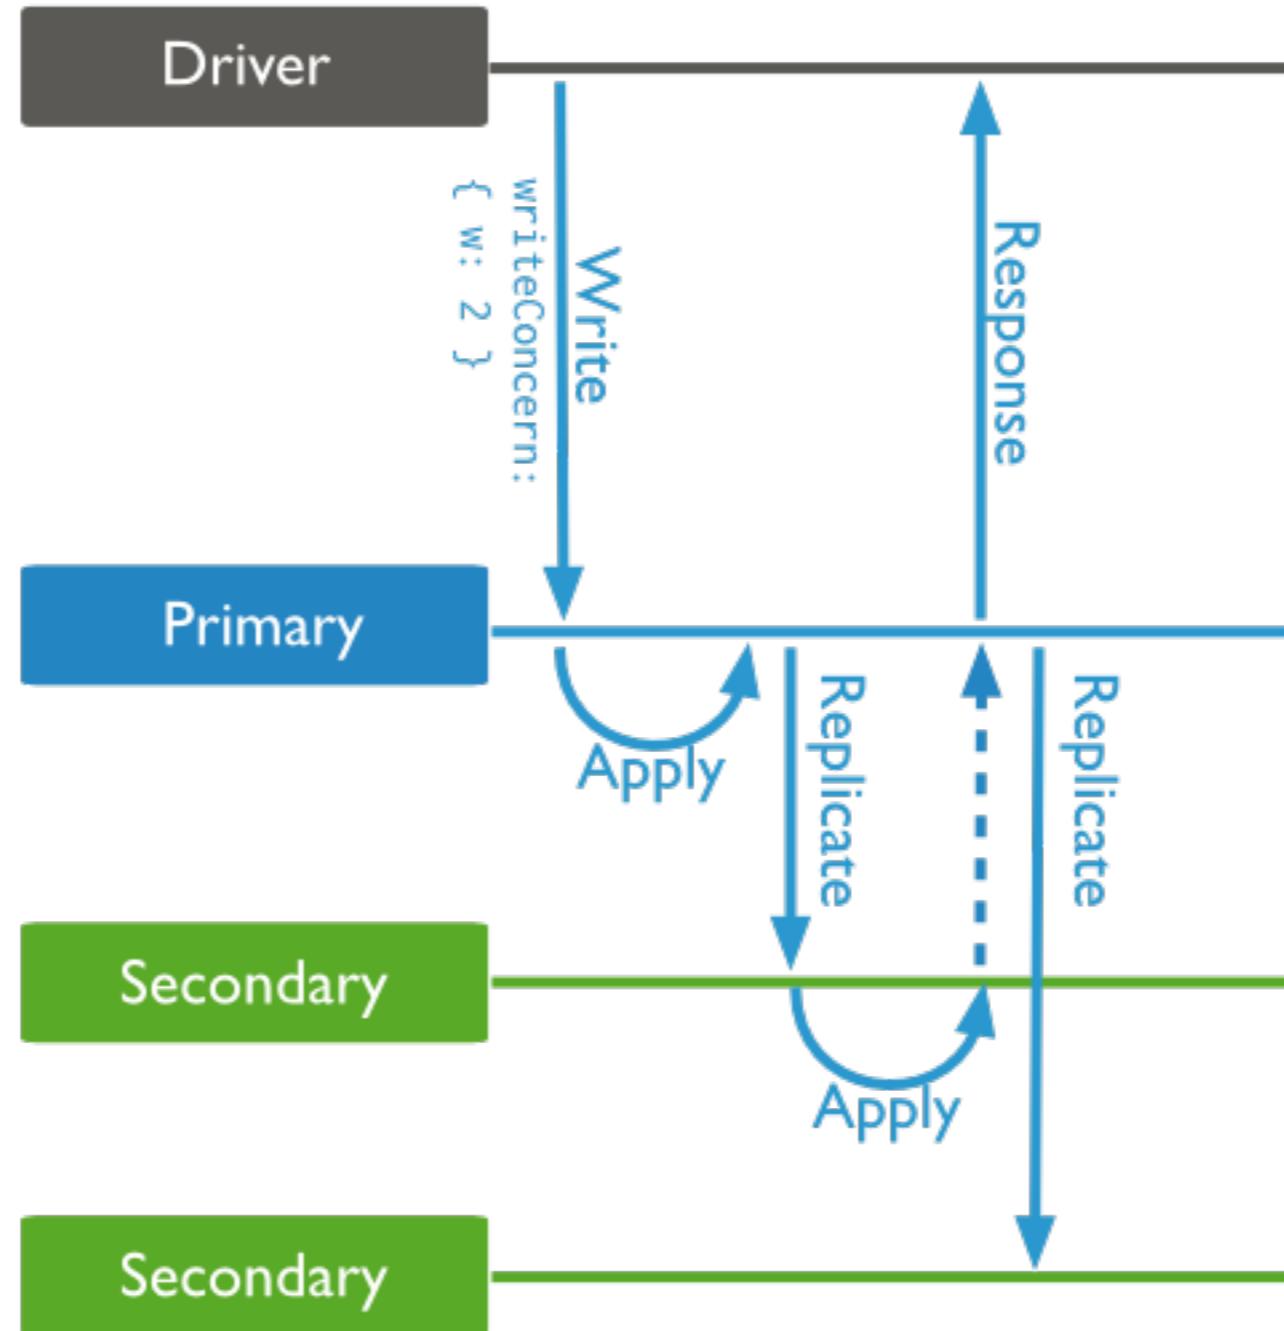
Other nodes is Secondary node



Working with MongoDB



Working with MongoDB



Replica set features

A cluster of N nodes

Anyone node can be primary

All write operations go to primary

Automatic failover

Automatic recovery

Consensus election of primary



Workshop

Setup a replica set

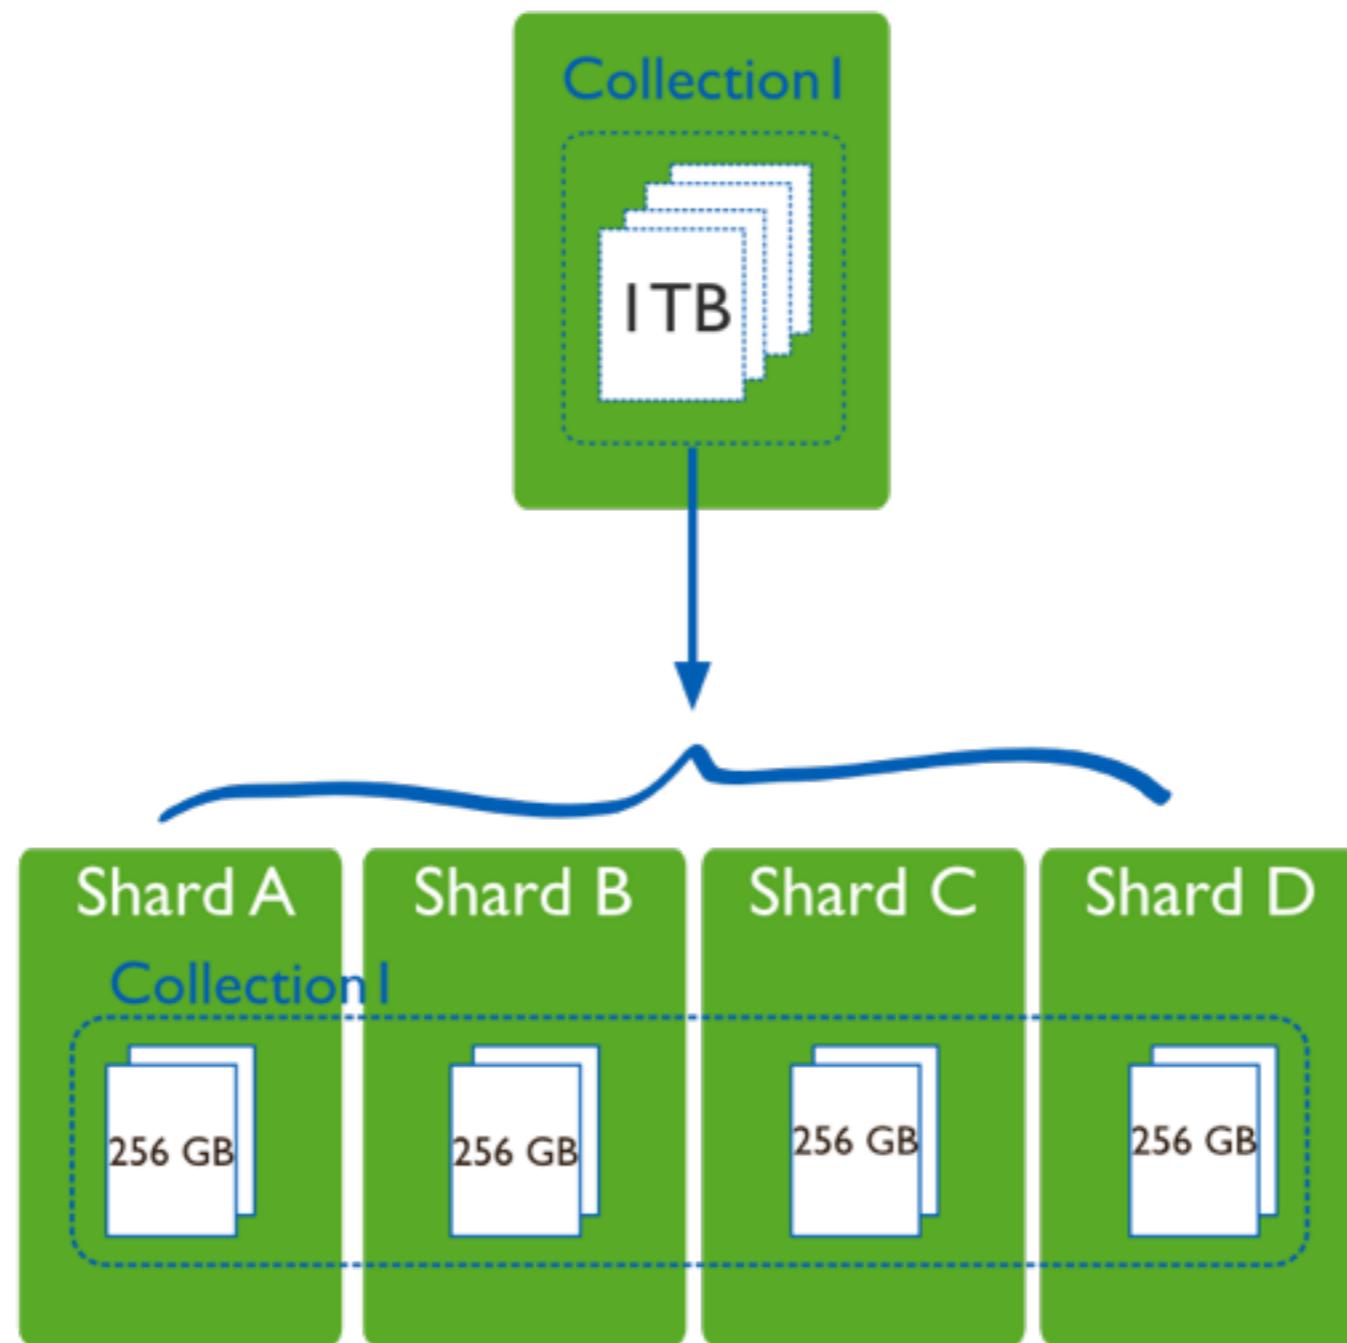
06_replica.txt



Sharding



What is Sharding ?



What is Sharding ?

The process of storing data records
across multiple machine

Approach for data growth



why Sharding ?

In replication all writes go to master

Latency sensitive queries
still go to master

Single replica set has limitation of 12 node



why Sharding ?

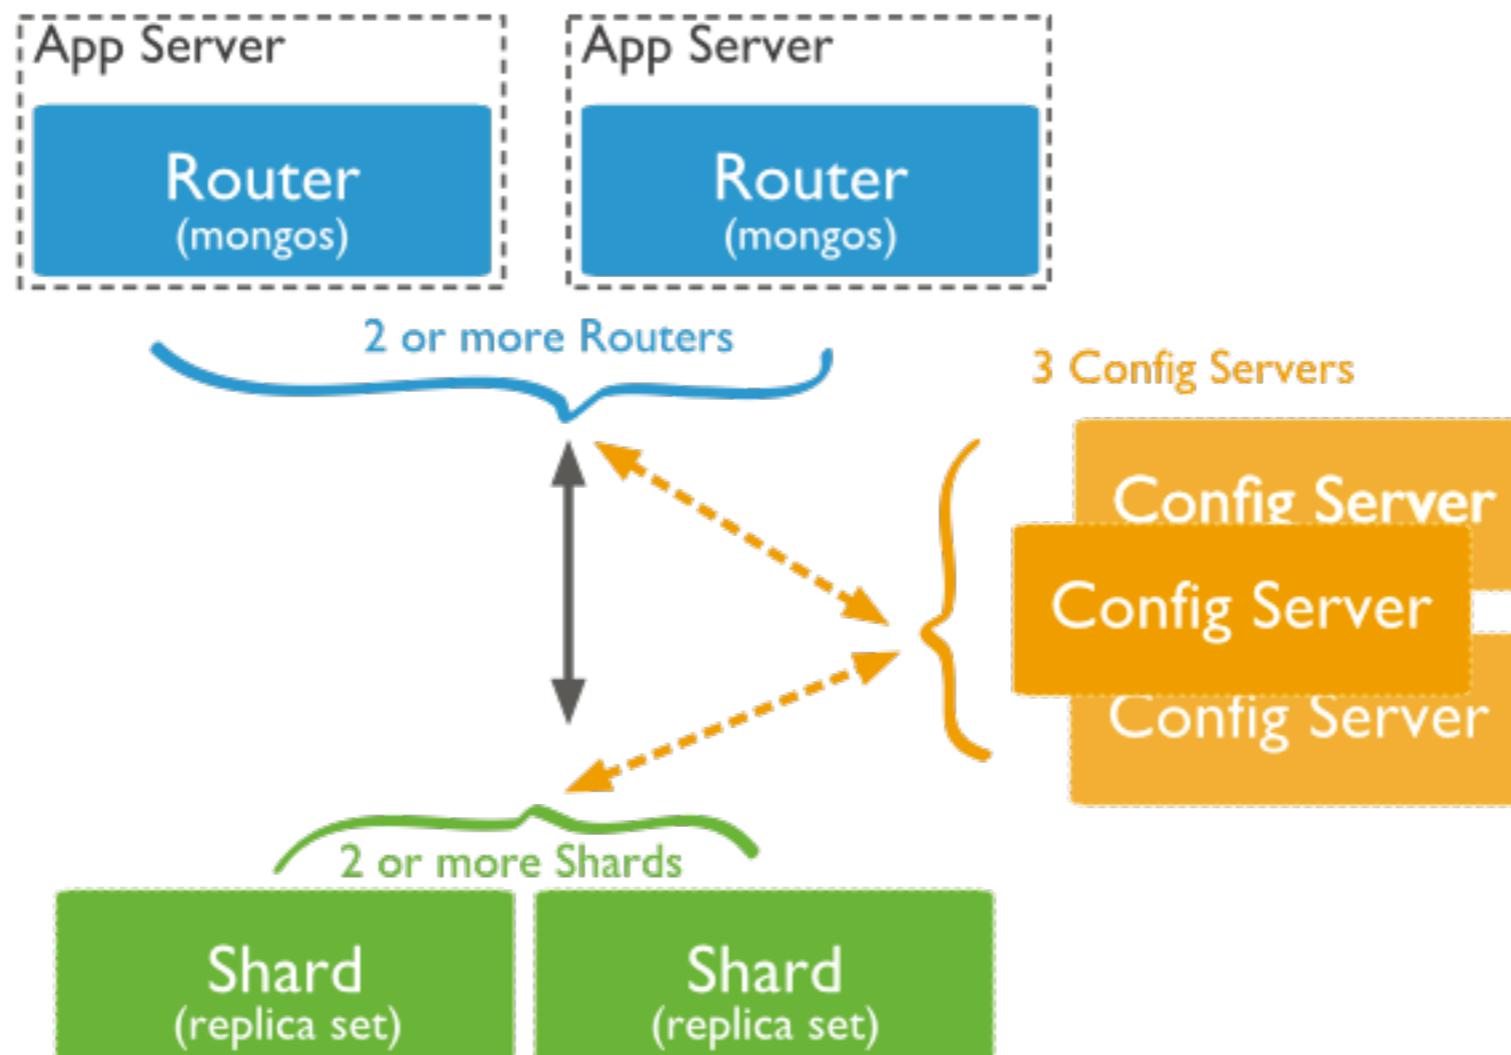
Memory can not be large enough

Local disk is not big enough

Vertical scaling is too expensive



Sharding in MongoDB



Workshop

Working with sharding

07_sharding.txt



Backup and Restore



Create backup

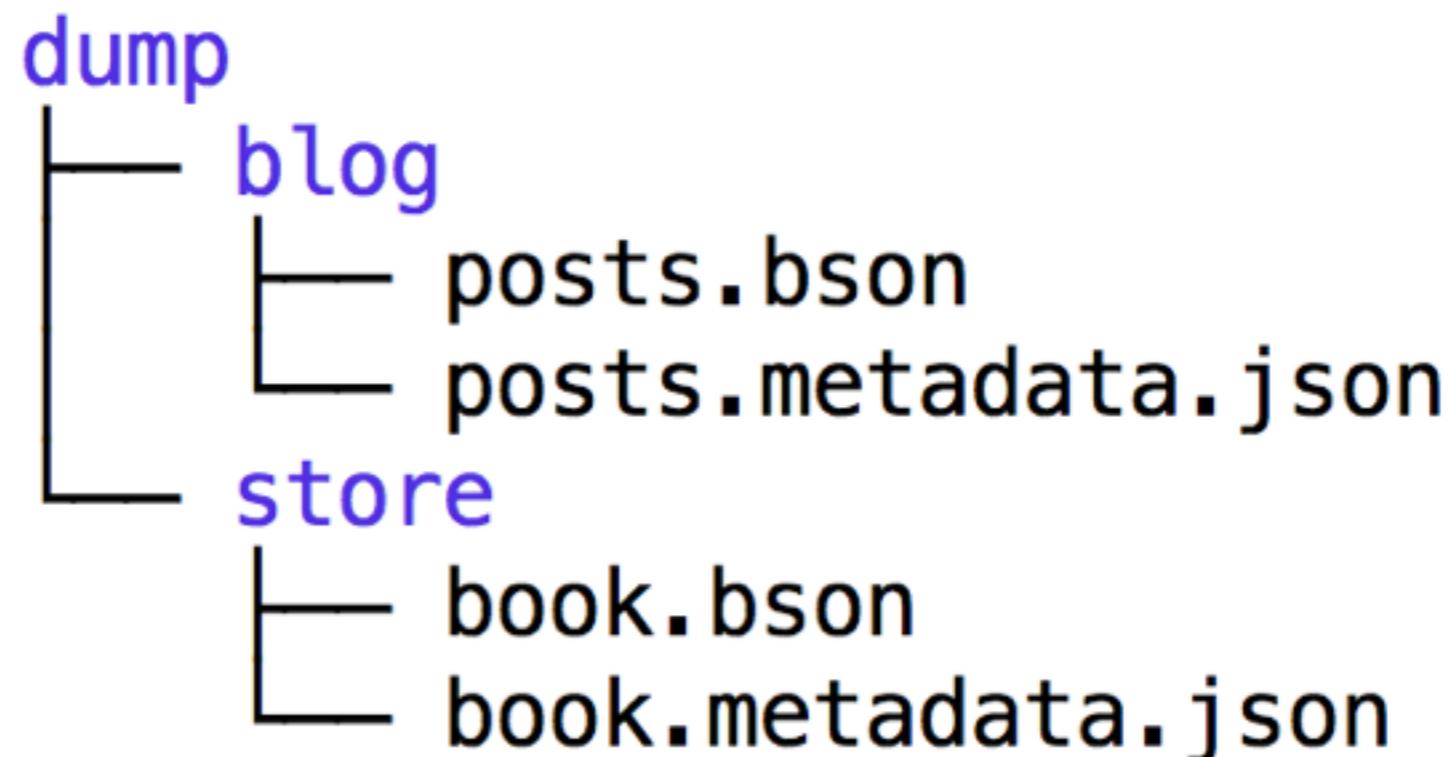
\$mongodump

```
writing blog.posts to
writing store.book to
done dumping store.book (3 documents)
done dumping blog.posts (1 document)
```



Create backup

\$mongodump



Mongodump usage

```
$mongodump --host HOST_NAME –port PORT_NUMBER
```

HOST_NAME = 127.0.0.1

PORT = 27017



Mongodump usage

```
$mongodump --dbpath DB_PATH  
          --out BACKUP_DIRECTORY
```



Mongodump usage

```
$mongodump --collection COLLECTION  
--db DB_NAME
```



Restore data

\$mongorestore



Monitoring tools

mongostat

mongotop

<https://docs.mongodb.com/manual/administration/monitoring/>



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

Monitoring tools

Ganglia
Motop
mtop

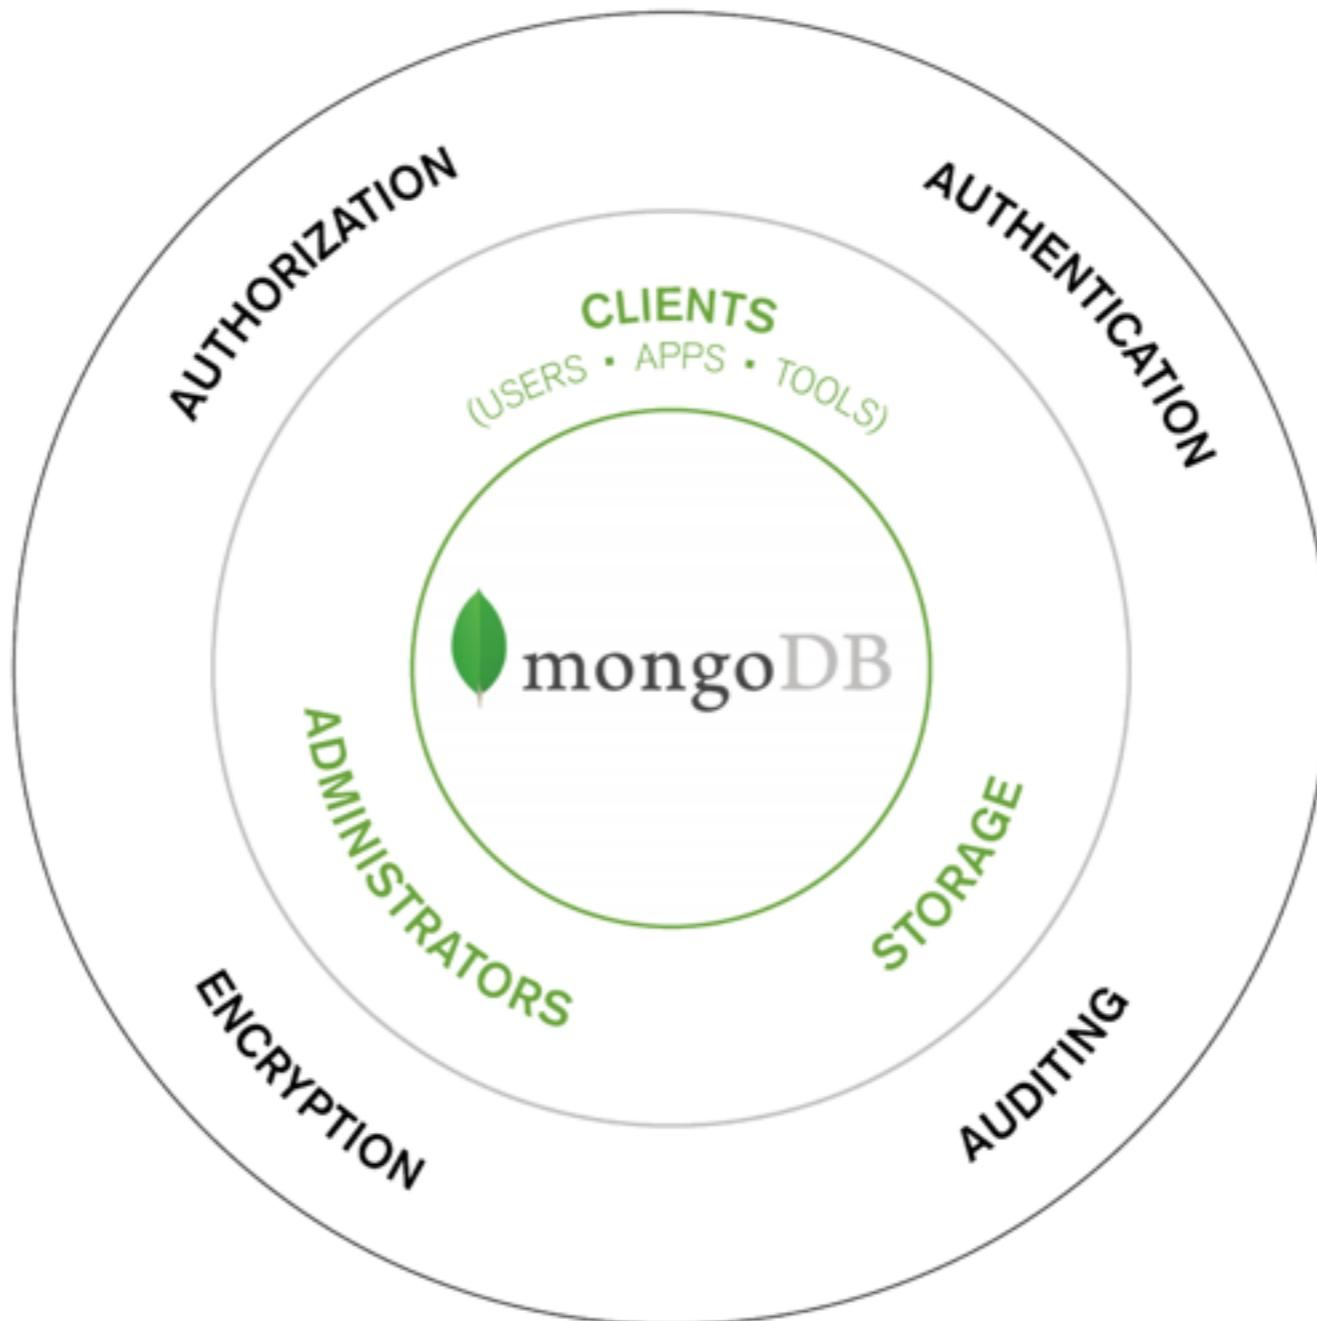
Munin
Nagios

<https://docs.mongodb.com/manual/administration/monitoring/>



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

Security



Security checklist

Enable access control

Enforce authentication

Configure rule-based access control

Encrypt communication

<https://docs.mongodb.com/manual/administration/security-checklist/>



บริษัท สยามช่างนาฏกิจ จำกัด และเพื่อนพ้องน้องพี่

Security checklist

Limit network exposure

Audit system activity

Encrypt and protect data

Run MongoDB with a dedicated user

<https://docs.mongodb.com/manual/administration/security-checklist/>



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

Security checklist

Run MongoDB with secure configuration

Audit system activity

Encrypt and protect data

Run MongoDB with a dedicated user

<https://docs.mongodb.com/manual/administration/security-checklist/>



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

Workshop

Secure MongoDB configuration

08_security.txt



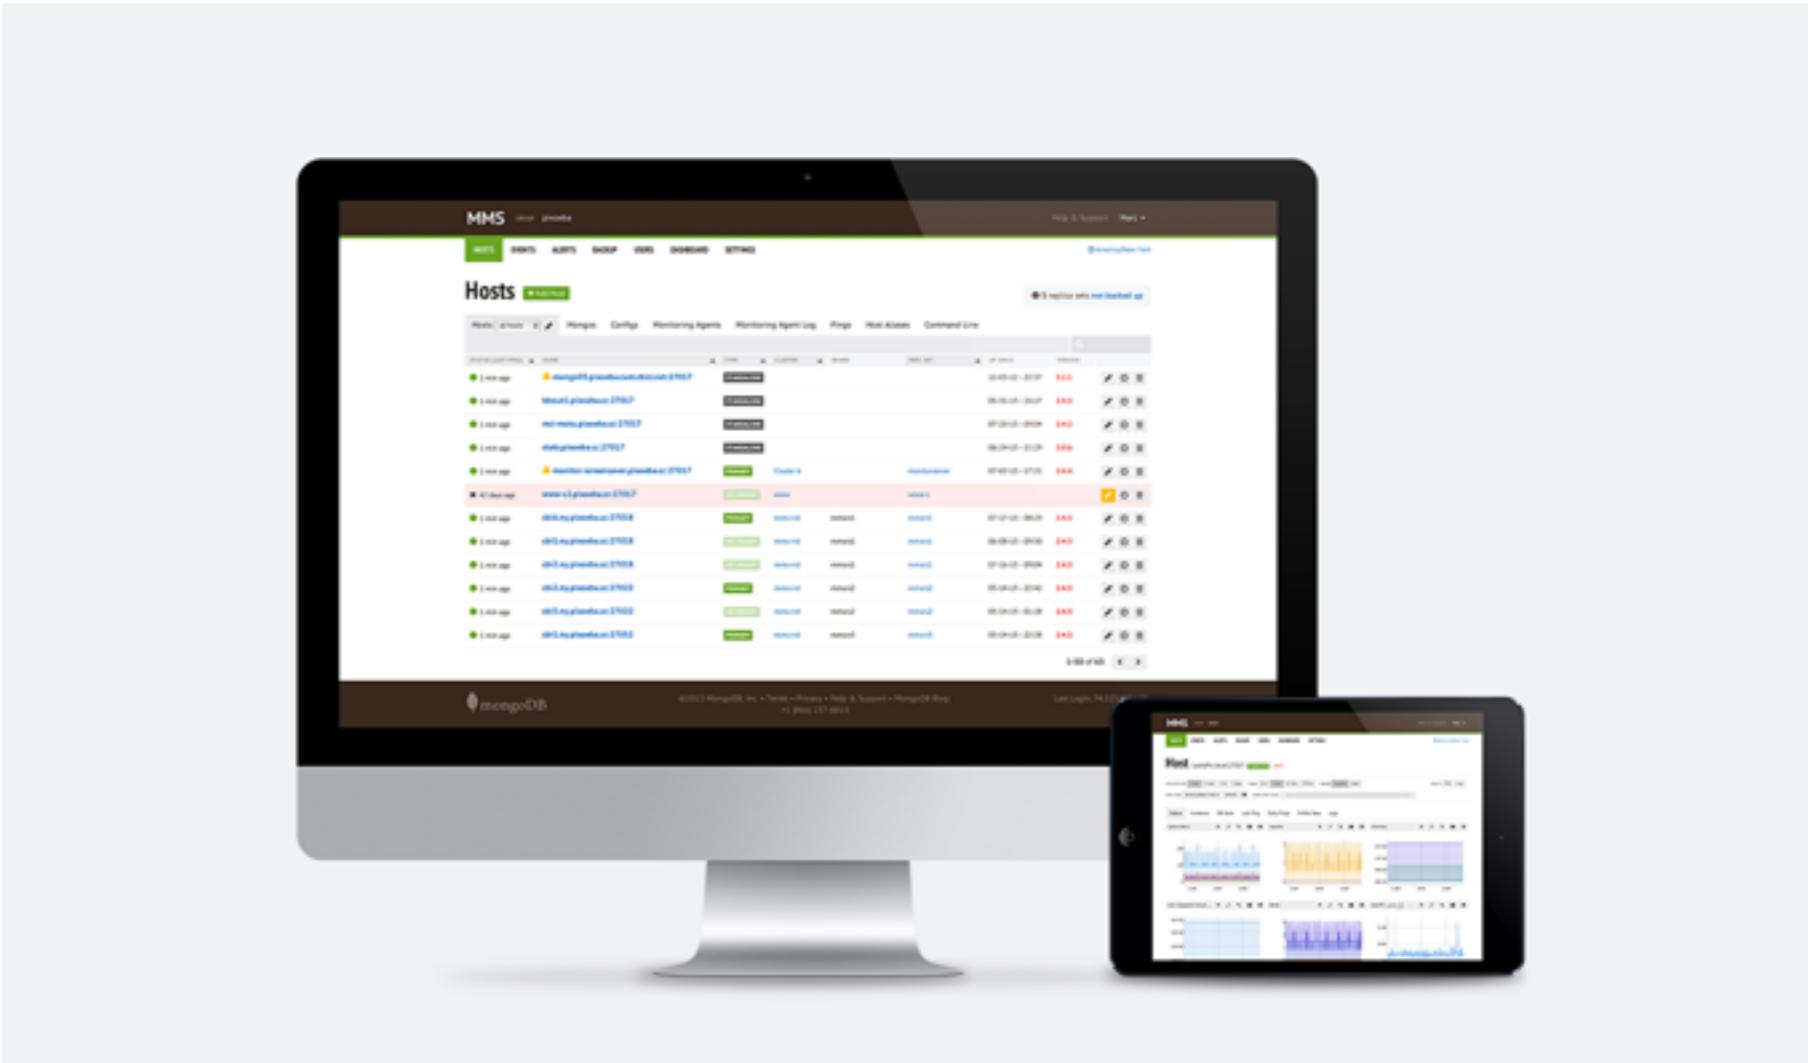
Tools for Administrator

<https://docs.mongodb.com/ecosystem/tools/administration-interfaces/>



บริษัท สยามชานาญกิจ จำกัด และเพื่อนพ้องน้องพี่

MongoDB cloud manager



Mongoclient

objects
10.3.10.54:27017
objectsDB +

Database Stats

Management BETA

- Database Dump/Restore
- File Management
- Admin Queries
- Easy Edit

Collections Add

aytsv01

The server is up for 1179158 seconds

The performances which are shown here are coming from `serverStatus` and `stats` queries. Therefore, before having a look at below graphs consider `mongoclient` is already executing `serverStatus` at regular intervals. See [here](#) for more info

Mongodb version is 3.2.1
Process id is 2678
Process name is mongod

0 Collections 0 Objects
0.00 MBs Data Size 0.00 MBs Storage Size
0 Indexes 0.00 MBs Index Size
0.00 MBs File Size 0 Extents

Memory Usage

Connections , Available: 808

Network

feedback



Mongo-express

Mongo Express Database: test Collection: CollectionB

Viewing Collection: CollectionB

New Document

Simple Advanced

Key Value String Q Find

← First ← Prev Next → Last →

_id	name	age	address	photo	media	moredata	evenMoreData	Delete
568a685def237209ab9d65a0	Bob Jones	12	123 Fake Street					
568a695acee7a79aa5ed879	Oscar Grouch	14	Sesame Street					
568cf29a4df646a3eee55f1c	Data URI Video Example							
568cf5ff03f5b400fce49749	Data URI Audio Example							
568fcfc8e34626a4c1594a5f5	Nested Data Example				{ "key1": "value1", "key2": "some value" }	{ "asdf": "asdfsdf", "doubleNested": { "dN1": 123 } }	{ "Keeeyss": "Valueessss" }	



adminMongo

adminMongo Connections

Database: Blog / Collection - data

Home / Local (connection) / Blog (database) / data (collection)

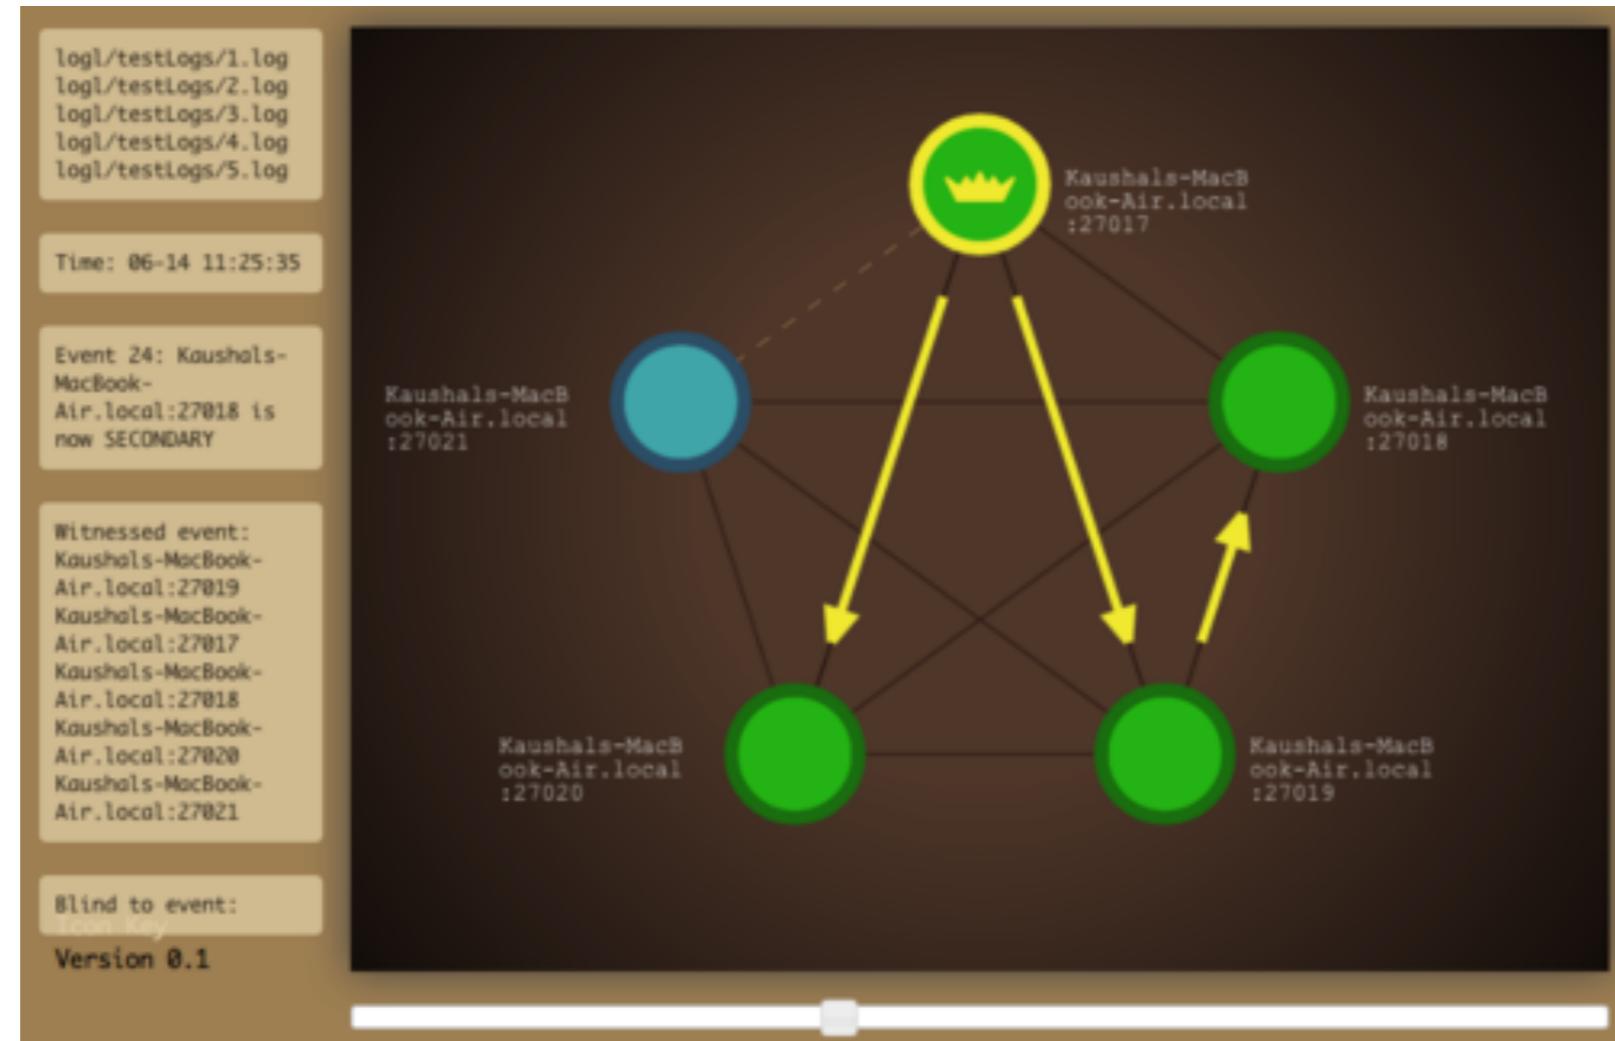
New document Search documents Reset Manage indexes

Document	Actions
{ "_id": "56a97ecdf718fe9a4f59947e", "street": "2 Avenue", "zipcode": "10075", "building": "1480" }	Delete Edit
{ "_id": "56a97eceef718fe9a4f59947f", "street": "2 Avenue", "zipcode": "10075", "building": "1480" }	Delete Edit
{ "_id": "56a97ecff718fe9a4f599480", "street": "2 Avenue", "zipcode": "10075", "building": "1480" }	Delete Edit
{ "_id": "56a97ed1f718fe9a4f599481", "street": "2 Avenue", "zipcode": "10075", "building": "1480" }	Delete Edit
{ "_id": "56a97ed2f718fe9a4f599482", "street": "2 Avenue", "zipcode": "10075", "building": "1480" }	Delete Edit

1 2 3 4



Edda



HumongouS.io

HumongouS.io

Search...

Customers

	Email	Fullscreen	Is active ?	Address	Picture	Gender	Phone
 Alyssa Johnson	mark.b@example.com	Mark Bradley	●	2858 Bollinger Rd			(777) - 115-7961
✓ Store project	ryankelly@example.com	Kelly Ryan	●	3847 Hillcrest Rd			(855) - 575-2507
Dashboard	elssieross@example.com	Elsie Ross	●	6249 Karen Dr			(916) - 585-7247
Customers	j.caleb@example.com	Caleb Jordan	●	Caleb Jordan			(842) - 926-1769
Inventory	nguyen@example.com	Billie Nguyen	●	4206 E Pecan St			(673) - 403-4362
Orders	darlene@example.com	Darlene Richards	●	7505 Robinson Rd			(255) - 752-3492
Orderlines	audreyb@example.com	Audrey Burton	●	4863 Red Saturn Dr			(418) - 301-9724
Products	mrfowler@example.com	Arthur Fowler	●	3413 Adams St			(285) - 600-2453
> Payment Information	ferguson@example.com	Elijah Ferguson	●	3388 Central St			(724) - 316-6447
> Logs	jamie@example.com	Jamie Herrera	●	5228 Kraft Ave			(468) - 857-1837
	dixon.r@example.com	Ramona Dixon	●	8880 Fifth St			(932) - 587-1757
	joe.hall@example.com	Joe Hall	●	2652 Andreas Ave			(934) - 536-7789
	smitdh@example.com	Marsha Schmidt	●	3737 Daisy Dr			(639)-156-7329

1-100 of 2,927   

✓ Store project

Dashboard

Customers

Inventory

Orders

Orderlines

Products

> Payment Information

> Logs

Settings

Logout



Mongo Management Studio

The screenshot shows the interface of Mongo Management Studio version 1.0. On the left, the sidebar lists databases: aamainph, admin, adpsneg, baboon_logs, baboon_rights, blog, demo, zipcode, demo123, do2, do2_dwh, enterprise, icd10, local, lp, lr, mongo_management_studio_rig, test_baboon_logs, test_baboon_rights, test_mms, vi, and wines. The 'demo' database is selected. In the main area, the 'demo.zipcode' collection is displayed. The interface includes a toolbar with 'Run ▶', 'F', 'A', 'C', and other icons, along with 'Add document' and 'Themes' and 'Collection' dropdowns. A search bar at the top of the main area contains the query '1 city: "BARRE"'. Below the search bar is a table showing two documents:

	_id	city	loc	pop	state
<input type="checkbox"/>	"01005"	"BARRE"	Array [2]	4546	"MA"
<input type="checkbox"/>	"05641"	"BARRE"	Array [2]	14994	"VT"

At the bottom of the main area, there are navigation buttons for pages 1 of 1, page 10, and 2 items.



Resources

<https://github.com/up1/course-introduction-mongodb>

